

Backtesting Foundations for Quant Research

Build causal, executable, costed, reconciled, and independently validated simulations.

Library of the Quant

Educational reference manual

Manual Profile

LIBRARY ID	LOTQ-0032
CATEGORY	Backtesting
DIFFICULTY	Beginner
FORMAT	Single-column field-guide manual
USE	Study, lookup, and research-system reference

Educational Use Only

This manual is educational material for quant research study. It is not financial advice, not a trading recommendation, and not a guarantee of correctness or live trading usefulness.

Table of contents

BEGINNER

ORIENTATION

Title	1
How to use this manual	5
Notation and simulation contract	6
Reference backtesting workflow	7

FOUNDATIONS

A backtest is a historical decision simulation	8
The strategy contract: hypothesis to executable rule	9
Information, decision, order, and fill clocks	10
Point-in-time universe and survivorship control	11
Benchmark, cash, and return definition	12

DATA AND STATE I

Adjusted prices and corporate-action integrity	13
Calendars, alignment, and asynchronous markets	14
Missing observations, delistings, and terminal events	15

Table of contents

BEGINNER

DATA AND STATE II

Signal, target, order, fill, and holding state	16
The portfolio ledger and self-financing identity	17

PORTFOLIO AND EXECUTION

Position sizing and capital normalization	18
Rebalance schedules, buffers, and path dependence	19
Transaction costs: explicit and implicit	20
Fill models and slippage assumptions	21
Liquidity, participation, and capacity	22
Short selling, borrow, and financing	23

VALIDATION I

In-sample, validation, and out-of-sample evidence	24
Lookahead bias and information leakage	25
Overlapping holdings and dependent observations	26
Benchmarks, null models, and incremental value	27

Table of contents

BEGINNER

VALIDATION II

Multiple testing and research search accounting	28
Performance metrics and metric contracts	29
Uncertainty, confidence intervals, and resampling	30
Robustness surfaces and sensitivity analysis	31
Regime, subgroup, and conditional performance	32
Attribution: signal, sizing, constraint, and execution	33

SYSTEM DESIGN

Reproducibility, lineage, and immutable artifacts	34
Backtest-to-production parity	35
Backtest QA, invariants, and adversarial tests	36
Promotion, monitoring, and retirement	37

VISUAL LABS AND REVIEW

From historical evidence to governed capital	38
Worked three-day backtest and ledger reconciliation	39
Causal signal, order, fill, and label timeline	40
Signal-to-holding state machine	41
Gross backtest edge to executable net return	42
Equity, drawdown, and path interpretation	43
Lookback and rebalance robustness	44
Backtest behavior across operating states	45
Failure modes and first diagnostics	46
Validation and promotion checklist	47
Final active recall and system index	48

How to use this manual

BEGINNER

1. START WITH THE DECISION CONTRACT

Write the hypothesis, eligible universe, observation and availability clocks, signal rule, sizing, rebalance, constraints, execution, costs, accounting, benchmark, and acceptance criteria before reading performance.

2. FOLLOW STATE, NOT FORMULAS ALONE

Trace signal, target, order, fill, holding, cash, and NAV as separate objects. This reveals whether value came from prediction, portfolio choices, market exposure, or impossible execution.

3. AUDIT CAUSALITY AT EVERY BOUNDARY

Check data vintages, as-of joins, membership, fitted transforms, label overlap, order latency, and first eligible fill. A single leaked boundary invalidates downstream statistics.

4. DEMAND ECONOMIC RECONCILIATION

Every position change must come from a fill or corporate action, every cash movement needs a source, and every return must reconcile to the ledger. Stop on unexplained residuals.

5. EVALUATE THE RESEARCH PROCEDURE

Record all trials, use chronological selection, preserve sealed evidence, account for dependence, and compare with simple benchmarks. The chosen backtest inherits the search that produced it.

6. STRESS IMPLEMENTATION AND SCALE

Vary latency, spread, impact, participation, borrow, missing events, liquidity, and exit conditions. Recompute holdings and risk at each capital level rather than scaling P&L mechanically.

DECISION STANDARD

Promote only when point-in-time data, causal state transitions, accounting, independent net evidence, robustness, capacity, parity, monitoring, rollback, and retirement are defensible as one system.

Notation and simulation contract

BEGINNER

$I_{t,T}$ AND $U_{t,T}$

Information legally available and the eligible point-in-time universe at decision t . Availability, membership, identity, quality, liquidity, and borrow policies are part of these objects.

$S_{i,T}$ AND $MU_{i,T}$

Signal and expected return for asset i . Keep raw evidence, transformed score, calibration, horizon, and forecast uncertainty separately versioned.

$W_{t,T}^0$ AND $W_{t,T}^*$

Unconstrained desired weights and final feasible targets after risk, liquidity, turnover, borrow, and compliance constraints.

$O_{i,T}$ AND $F_{i,T}$

Submitted order and realized fill. Quantity, side, type, timestamps, prices, fees, status, and parent decision define the transition from intent to fact.

$Q_{i,T}$ AND $P_{i,T}$

Actual held quantity and lawful valuation price. Positions change only through fills and corporate actions; stale marks remain flagged state.

NAV_t AND R_t

Net asset value from cash, positions, receivables, and liabilities, and return after external-flow adjustment. Metric denominator and currency are explicit.

ΔW_t AND $TURNOVER_t$

Target trade and one-way portfolio turnover, commonly 0.5 times the absolute weight change for normalized books. Actual traded turnover comes from fills.

$C(O, STATE)$

Trade-level cost function including spread, explicit fees, impact, delay, financing, borrow, and opportunity loss conditional on market state.

MINIMUM RUN IDENTITY

Bind hypothesis, code, parameters, data vintages, calendars, environment, random seed, simulator, metrics, and approvals to one immutable artifact graph.

Reference backtesting workflow

BEGINNER

1

Register the strategy contract

Mechanism, universe, clocks, signal, sizing, execution, costs, metrics, falsifiers, and limits.

2

Freeze point-in-time evidence

Identity, vintages, membership, actions, calendars, prices, borrow, and quality states.

3

Build causal features and labels

As-of joins, fitted transforms, decision frontier, executable entry, horizon, purge, and embargo.

4

Simulate portfolio state

Signal, target, constraints, open orders, fills, holdings, cash, actions, and valuation.

5

Reconcile economics

Ledger, NAV, returns, benchmark, explicit costs, financing, exposures, and attribution residual.

6

Validate the procedure

Temporal selection, baselines, search accounting, dependence, uncertainty, robustness, and regimes.

7

Stress implementation and scale

Latency, slippage, participation, borrow, outages, capacity, common exit, and tail loss.

8

Promote governed state

Immutable artifact, parity, shadow replay, monitoring, owners, rollback, incidents, and retirement.

HARD-STOP RULE

Future knowledge, survivor-only rosters, unreconciled NAV, same-bar fantasy fills, hidden search, unsupported costs, unavailable shorts, or missing rollback blocks promotion regardless of historical performance.

A backtest is a historical decision simulation

BEGINNER

ONE-SENTENCE MEANING

A backtest replays a fully specified trading procedure through historical information, decisions, orders, fills, holdings, cash flows, and marks. Its purpose is to estimate how the procedure might have behaved, not to certify future profit.

MEMORY JOGGER

Simulate what the system could know and do at each instant. If the historical run cannot be reconstructed as a sequence of admissible decisions, it is only a retrospective calculation.

WHY IT MATTERS IN QUANT RESEARCH

This framing forces data timing, portfolio state, execution, and accounting into the research claim. It prevents a predictive statistic from being mistaken for a tradable strategy and makes failure diagnosis possible.

FORMULA INTUITION

State evolves as $S_t = F(S_{t-1}, I_t, \text{rule, constraints, fills})$, where I_t contains only information available by the decision cutoff. Portfolio return is produced by realized holdings and cash flows, not by the signal alone.

SIMPLE MARKET EXAMPLE

A close-to-close momentum signal computed after Monday's close cannot receive Monday's closing price. The simulator forms the decision after the cutoff, generates Tuesday orders, records attainable fills, and carries the resulting positions forward.

PYTHON / SYSTEM IMPLEMENTATION

Represent each event with event_time, availability_time, decision_time, order_time, fill_time, and version IDs. Use a deterministic event loop or vectorized equivalent whose outputs reconcile to an order and position ledger.

CONFUSIONS TO AVOID

A backtest is not a chart of an indicator shifted against future returns. High correlation can coexist with impossible timing, unstable position mapping, unaffordable turnover, or accounting errors.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The simulator joins point-in-time data, signal code, portfolio construction, execution assumptions, accounting, validation, and artifact lineage. Production should implement the same state transitions with live events.

RELATED CONCEPTS AND QUICK RECALL

Related: event-driven simulation, state machine, point-in-time data, execution model. Recall task: Describe the exact historical event that turns one signal value into a held position.

The strategy contract: hypothesis to executable rule

BEGINNER

ONE-SENTENCE MEANING

A strategy contract converts an economic hypothesis into explicit inputs, clocks, universe rules, signal transformations, position logic, execution assumptions, risk limits, and evaluation criteria. It defines what is being tested before results are viewed.

MEMORY JOGGER

Write the rule so another researcher can replay it without judgment calls. Every discretionary repair after seeing performance creates a new candidate and must enter the search record.

WHY IT MATTERS IN QUANT RESEARCH

Clear contracts separate economic reasoning from implementation choices and expose where apparent alpha can be manufactured. They also let validation compare like with like across code versions and datasets.

FORMULA INTUITION

A minimal contract is Procedure = {U_t, I_t, signal, target, constraints, rebalance, execution, costs, accounting, metrics, stop rules}. Each member has a version and an effective time.

SIMPLE MARKET EXAMPLE

A claim that value stocks outperform is incomplete. The contract must specify the valuation field and vintage, eligible securities, ranking date, holding horizon, weighting, delisting treatment, costs, and executable entry.

PYTHON / SYSTEM IMPLEMENTATION

Store a machine-readable manifest beside every result and hash its dependencies. Reject runs with unresolved parameters, mutable external queries, missing calendar versions, or implicit defaults.

CONFUSIONS TO AVOID

A strategy name is not a specification, and a notebook cell order is not governance. Changing one missing-value rule or rebalance timestamp can materially alter the strategy even when the label stays unchanged.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The contract is the boundary between hypothesis registry, feature store, backtest engine, validation suite, model registry, and deployment approval. Every live order should trace back to one approved contract.

RELATED CONCEPTS AND QUICK RECALL

Related: research preregistration, manifests, decision rules, change control. Recall task: List ten fields required to reproduce a strategy without asking its original author.

ONE-SENTENCE MEANING

A causal backtest distinguishes when an observation occurred, when it became available, when the strategy decided, when an order was released, and when a fill could occur. These clocks determine whether the simulated trade was feasible.

MEMORY JOGGER

The market timestamp is not automatically the knowledge timestamp. Add ingestion, computation, approval, and routing latency before selecting a fill window.

WHY IT MATTERS IN QUANT RESEARCH

Most severe backtest errors are timing errors disguised as data convenience. Correct clocks prevent same-bar lookahead and reveal whether short-lived edge survives the path to execution.

FORMULA INTUITION

Require $\text{availability_time} \leq \text{decision_time} < \text{order_time} \leq \text{fill_time}$. A forward label begins at or after the first eligible fill and ends at the declared horizon.

SIMPLE MARKET EXAMPLE

An earnings release stamped 16:05 may contain a fiscal period ending months earlier. A strategy that trades the 16:00 close uses the report before publication even though the underlying quarter is historical.

PYTHON / SYSTEM IMPLEMENTATION

Keep all clock fields in UTC plus venue calendars, and test ordering invariants on every event. Model realistic latency distributions rather than subtracting an arbitrary bar from the signal.

CONFUSIONS TO AVOID

Shifting a feature by one row does not prove causality when rows are irregular or assets trade on different calendars. End-of-day files may also be revised after the date printed in the record.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Clock semantics connect source ingestion, as-of joins, feature computation, scheduling, execution simulation, and label creation. The same timing contract should drive live scheduling and replay.

RELATED CONCEPTS AND QUICK RECALL

Related: as-of joins, latency, same-bar bias, market calendars. Recall task: Draw the five clocks for a signal formed from a post-close filing and identify the earliest lawful fill.

Point-in-time universe and survivorship control

BEGINNER

ONE-SENTENCE MEANING

The point-in-time universe contains securities that were identifiable, eligible, and tradable under rules known at each historical decision. It retains entrants, exits, delistings, and temporary ineligibility instead of projecting today's survivors backward.

MEMORY JOGGER

Reconstruct who could have been chosen before ranking them. Universe membership changes both the opportunity set and every cross-sectional transform derived from it.

WHY IT MATTERS IN QUANT RESEARCH

Survivorship bias removes failed securities and inflates historical quality, liquidity, and returns. Correct rosters are also necessary for realistic turnover, capacity, benchmark comparison, and missing-data behavior.

FORMULA INTUITION

Let $U_t = \{i: \text{listing}_i(t)=\text{active and eligibility}_i(t)=\text{true using information available at } t\}$. All ranks, portfolio constraints, and breadth metrics are conditional on U_t .

SIMPLE MARKET EXAMPLE

Using current index constituents for a 15-year test excludes bankrupt firms and includes later winners before they joined. A historical constituent file restores the actual decision set and captures deletion returns.

PYTHON / SYSTEM IMPLEMENTATION

Use effective-dated security masters, listing histories, identifier maps, exchange status, eligibility reasons, and immutable universe snapshots. Save the exact roster hash with each rebalance.

CONFUSIONS TO AVOID

A vendor's clean continuous history can still be survivor-only. Filtering assets because their future history is short or incomplete uses outcome information even if the price columns themselves are correctly lagged.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Universe construction precedes feature calculation and feeds the simulator, benchmark, borrow model, and capacity analysis. Monitoring should separate signal changes from roster changes.

RELATED CONCEPTS AND QUICK RECALL

Related: survivorship bias, security master, index membership, delisting returns. Recall task: Explain why deleting names with incomplete future data changes the historical selection problem.

Benchmark, cash, and return definition

BEGINNER

ONE-SENTENCE MEANING

Backtest performance depends on a declared return definition covering price, income, financing, costs, and benchmark. The numerator and capital base must match the actual portfolio state.

MEMORY JOGGER

Know what one reported percent means. Distinguish asset return, strategy P&L, return on net asset value, return on gross exposure, and active return versus a benchmark.

WHY IT MATTERS IN QUANT RESEARCH

Ambiguous returns make strategies incomparable and can hide leverage, idle cash, short financing, or uninvestable benchmark drift. Correct definitions allow risk and attribution to reconcile to economic cash flows.

FORMULA INTUITION

Portfolio return $r_t = (\text{NAV}_t - \text{NAV}_{t-1} - \text{external_flow}_t) / \text{NAV}_{t-1}$. Active return is $r_t - r_{b,t}$, while long-short spread return may use gross capital and must state that denominator.

SIMPLE MARKET EXAMPLE

A market-neutral book earns 1 percent of gross exposure with 200 percent gross. The same P&L is 2 percent on NAV before financing, so quoting only the spread return doubles interpretation risk.

PYTHON / SYSTEM IMPLEMENTATION

Maintain a double-entry or fully reconcilable ledger for cash, positions, receivables, fees, borrow, and corporate actions. Publish metric metadata including currency, compounding, denominator, benchmark, and annualization.

CONFUSIONS TO AVOID

Total-return prices may already include dividends, so booking separate dividends double counts them. Dollar-neutral does not mean zero financing or zero NAV risk, and excess return is not alpha without an exposure model.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Return accounting sits after fills and valuation but before performance metrics, attribution, risk, and model comparison. Its definitions should be shared with live portfolio reporting.

RELATED CONCEPTS AND QUICK RECALL

Related: NAV, total return, active return, financing, compounding. Recall task: Reconcile one day's holdings, cash flows, fees, and marks into NAV return and active return.

Adjusted prices and corporate-action integrity

BEGINNER

ONE-SENTENCE MEANING

Corporate-action handling transforms raw market events into economically consistent prices, shares, cash flows, and identifiers without leaking later adjustment knowledge. Splits, dividends, mergers, spinoffs, rights, and delistings require event-specific treatment.

MEMORY JOGGER

Adjust either the history or the holdings ledger coherently, never both. Preserve raw observations and action records so every transformed value remains auditable.

WHY IT MATTERS IN QUANT RESEARCH

Bad adjustments create artificial returns, false signals, and broken position quantities that can dominate a backtest. Corporate actions also affect tradability and cash timing, not just chart continuity.

FORMULA INTUITION

For a split factor a , adjusted historical price may be $P^*_s = P_s/a$ and volume $V^*_s = aV_s$ before the event. A cash dividend belongs in total return or the cash ledger according to the chosen data model.

SIMPLE MARKET EXAMPLE

A 2-for-1 split halves the raw price and doubles shares. Treating the price drop as a loss while leaving position quantity unchanged creates both a false signal and a 50 percent NAV error.

PYTHON / SYSTEM IMPLEMENTATION

Ingest effective-dated action events, test price-share-value conservation, and reconcile transformed returns to vendor total-return series. Model merger consideration and delisting proceeds explicitly.

CONFUSIONS TO AVOID

A backward-adjusted series can embed actions learned later and may not reproduce the data available to the strategy. Splicing identifiers across acquisitions can fabricate a continuous asset that never existed.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The canonical market-data layer supplies raw and adjusted views to features while the accounting engine consumes event cash flows and quantity changes. Both views must reference the same action version.

RELATED CONCEPTS AND QUICK RECALL

Related: split adjustment, dividends, delistings, identifier continuity. Recall task: Show how price, shares, and portfolio value change through a split without creating profit or loss.

Calendars, alignment, and asynchronous markets

BEGINNER

ONE-SENTENCE MEANING

Calendar alignment maps observations, decisions, holdings, and returns onto the actual sessions of each asset and venue. It handles holidays, half-days, time zones, auctions, stale closes, and instruments that trade on different clocks.

MEMORY JOGGER

Align by economic availability, not by matching date labels. A Friday close in one market may precede or follow a Monday observation elsewhere depending on time zone and holiday structure.

WHY IT MATTERS IN QUANT RESEARCH

Misalignment produces artificial lead-lag relationships and impossible portfolio marks. Correct calendars are essential for cross-asset signals, international portfolios, and any rebalance spanning multiple venues.

FORMULA INTUITION

An as-of join selects the latest record with `availability_time <= decision_time`, subject to a maximum age. Portfolio return uses each instrument's lawful mark and records stale-price exposure separately.

SIMPLE MARKET EXAMPLE

A US signal generated after New York closes cannot use the same calendar date's Asian close if that close occurred earlier but cannot fill Asia until its next session. The holding period differs across legs.

PYTHON / SYSTEM IMPLEMENTATION

Use exchange-specific calendar libraries, UTC timestamps, explicit session IDs, and maximum-age rules. Unit tests should cover daylight-saving changes, half-days, holiday gaps, and missing auctions.

CONFUSIONS TO AVOID

Forward-filling prices across closed markets creates zero returns and can suppress measured volatility. Inner-joining calendars discards difficult sessions and changes the tested sample.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Calendar services support data alignment, feature windows, scheduling, execution windows, valuation, risk aggregation, and live replay. Calendar versions belong in the run manifest.

RELATED CONCEPTS AND QUICK RECALL

Related: session calendars, as-of joins, stale marks, cross-asset timing. Recall task: Explain how to value a portfolio when one market is closed and another has moved sharply.

Missing observations, delistings, and terminal events

BEGINNER

ONE-SENTENCE MEANING

Missing data is a state with a cause, while a delisting or terminal event is an economic outcome that must remain in the simulation. Policies should distinguish closure, outage, illiquidity, suspension, inapplicability, and permanent exit.

MEMORY JOGGER

Absence is not a flat return. Determine whether the strategy could hold, trade, value, or receive proceeds before choosing a fallback.

WHY IT MATTERS IN QUANT RESEARCH

Dropping missing or delisted names selectively removes distress and improves backtests. Explicit states preserve tail losses, operational uncertainty, and the costs of forced liquidation.

FORMULA INTUITION

$\text{Coverage}_t = \text{observed eligible values} / \text{eligible assets}$, but return treatment depends on event state. A terminal return includes final cash consideration or a conservative recovery assumption with sensitivity analysis.

SIMPLE MARKET EXAMPLE

A suspended stock has no print for ten sessions and later delists at a 70 percent loss. Forward-filling through the gap is acceptable only as a stale mark with risk flags, not as evidence of realizable zero-volatility value.

PYTHON / SYSTEM IMPLEMENTATION

Persist `missing_reason`, `last_valid_time`, `tradability_state`, `valuation_source`, `terminal_proceeds`, and `confidence`. Run optimistic, base, and conservative recoveries when exact outcomes are unavailable.

CONFUSIONS TO AVOID

Complete-case filtering changes the universe after outcomes are known. Filling missing returns with zero treats uncertainty as safety and understates both drawdown and liquidity risk.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Data-quality states feed eligibility, valuation, execution, risk, and incident handling. The backtest report should attribute results to missing and terminal-event policies.

RELATED CONCEPTS AND QUICK RECALL

Related: stale prices, suspensions, delisting bias, recovery assumptions. Recall task: Choose treatments for a market holiday, feed outage, trading halt, and permanent delisting.

Signal, target, order, fill, and holding state

BEGINNER

ONE-SENTENCE MEANING

A robust backtest keeps predictive signals, desired targets, submitted orders, realized fills, and actual holdings as separate state objects. Each transition can change exposure and deserves independent attribution.

MEMORY JOGGER

Signal is belief; target is intent; order is instruction; fill is fact; holding is current exposure. Collapsing them assumes perfect execution and hides why results changed.

WHY IT MATTERS IN QUANT RESEARCH

State separation makes transaction costs, partial fills, stale orders, constraints, and live reconciliation measurable. It also lets researchers test the signal without silently granting it impossible portfolio access.

FORMULA INTUITION

Compute $\text{target_qty}_t = f(\text{signal}_t, \text{risk}_t, \text{constraints}_t)$. Then $\text{order_qty}_t = \text{target_qty}_t - \text{current_qty}_t - \text{open_order_qty}_t$; holdings update only from confirmed fills and actions.

SIMPLE MARKET EXAMPLE

A signal raises the target from 100 to 300 shares, but a liquidity cap permits an order for 120 and only 70 fill. The next decision begins from 170 held shares plus any still-open order, not the 300-share target.

PYTHON / SYSTEM IMPLEMENTATION

Use immutable event records and derived state snapshots with unique IDs. Reconcile $\text{sum}(\text{fills})$ plus actions to position changes, and prevent duplicate fill ingestion through idempotency keys.

CONFUSIONS TO AVOID

Applying target weights directly to returns assumes instantaneous complete fills. Subtracting every new target from the previous target ignores actual holdings and double-orders when fills are partial.

WHERE THIS FITS IN THE RESEARCH SYSTEM

These objects connect model scoring, portfolio construction, order management, execution, accounting, and attribution. Shared decision and event IDs support production parity tests.

RELATED CONCEPTS AND QUICK RECALL

Related: portfolio state, OMS, partial fills, idempotency, reconciliation. Recall task: Trace one target change through constrained order quantity, partial fill, and next-period holdings.

The portfolio ledger and self-financing identity

BEGINNER

ONE-SENTENCE MEANING

The portfolio ledger records every economic event that changes cash, positions, liabilities, receivables, fees, or net asset value. A self-financing strategy changes wealth through market P&L and internal cash flows rather than unexplained capital creation.

MEMORY JOGGER

If ending NAV cannot be reconstructed from beginning NAV plus valued events, stop. Performance metrics are meaningless until the ledger balances.

WHY IT MATTERS IN QUANT RESEARCH

Ledger discipline catches double-counted dividends, missing fees, incorrect short proceeds, and price-quantity errors. It provides the audit trail needed to compare research, paper trading, and live books.

FORMULA INTUITION

$NAV_t = cash_t + \sum_i q_{i,t} P_{i,t} + receivables_t - liabilities_t$. The change reconciles to mark-to-market P&L, realized P&L, income, financing, fees, and external flows.

SIMPLE MARKET EXAMPLE

Selling short raises cash but creates a negative position liability; it does not create free equity. Borrow fees, dividends owed, margin interest, and closeout proceeds move distinct ledger accounts.

PYTHON / SYSTEM IMPLEMENTATION

Implement balanced journal entries or equivalent invariants for fills and actions. Produce daily reconciliation tables and fail the run when unexplained residual exceeds a tight tolerance.

CONFUSIONS TO AVOID

Summing weighted returns can be an approximation when weights change intraday, leverage varies, or cash flows occur. Ignoring the short liability overstates capital and understates risk.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The ledger is the authoritative source for NAV, exposures, returns, costs, attribution, and broker reconciliation. Model diagnostics should never bypass it for headline performance.

RELATED CONCEPTS AND QUICK RECALL

Related: double-entry accounting, self-financing portfolios, realized P&L, NAV. Recall task: Write the journal effects of buying stock, receiving a dividend, and covering a short.

Position sizing and capital normalization

BEGINNER

ONE-SENTENCE MEANING

Position sizing maps forecasts into desired exposures under a declared capital base, risk target, leverage policy, and name or asset limits. The same signal can produce radically different outcomes under different sizing rules.

MEMORY JOGGER

Backtest the capital allocator, not just the ranking. Record whether weights target dollars, volatility, beta, risk contribution, or another quantity and when each input was observable.

WHY IT MATTERS IN QUANT RESEARCH

Sizing controls concentration, drawdown, turnover, and capacity, often more than small changes in predictive accuracy. It also determines whether comparisons across assets and periods are economically fair.

FORMULA INTUITION

A volatility-scaled raw weight may be w_i proportional to s_i / σ_i , then normalized to a gross or risk target and projected into constraints. Estimated sigma must be lagged and capped to avoid explosive leverage.

SIMPLE MARKET EXAMPLE

Two futures have equal forecasts but annualized volatilities of 10 and 30 percent. Equal contracts triple the second instrument's risk, while inverse-volatility sizing equalizes only a first-order standalone estimate.

PYTHON / SYSTEM IMPLEMENTATION

Version the estimator window, floors, caps, gross target, leverage, contract multiplier, currency conversion, and rounding. Save pre-constraint and final targets for attribution.

CONFUSIONS TO AVOID

Volatility scaling is not risk parity when correlations matter. Normalizing weights after constraints can violate limits, and using future realized volatility leaks information into historical sizing.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Sizing belongs between forecast generation and order creation, consuming risk estimates and account state. Its output feeds constraint checks, turnover control, execution, and exposure reporting.

RELATED CONCEPTS AND QUICK RECALL

Related: volatility targeting, leverage, risk contribution, weight normalization. Recall task: Compare equal weight, inverse-volatility weight, and covariance-aware sizing for the same three forecasts.

Rebalance schedules, buffers, and path dependence

BEGINNER

ONE-SENTENCE MEANING

A rebalance rule determines when the portfolio may change and how existing holdings interact with new targets. Frequency, threshold, buffer, stagger, and order persistence make realized performance path dependent.

MEMORY JOGGER

A signal update is not automatically a trade. Model the exact state transition from current holdings to an executable target under the chosen schedule.

WHY IT MATTERS IN QUANT RESEARCH

Rebalance design governs turnover, delay, capacity, and sensitivity to small rank changes. It can preserve most forecast value while materially altering costs and drawdown behavior.

FORMULA INTUITION

Trade only when $\text{benefit}_t > \text{estimated_cost}_t + \text{uncertainty_buffer}$, or when a holding exits a wider buffer around its entry threshold. Turnover is $0.5 \times \sum_i |w_{i,t} - w_{i,t-1}|$ for a fully normalized book.

SIMPLE MARKET EXAMPLE

A stock enters the top quintile at rank 80.2 and falls to 79.8 next day. Immediate bucket rebalancing churns the position; a 75-85 hysteresis band avoids the round trip until evidence changes materially.

PYTHON / SYSTEM IMPLEMENTATION

Drive schedules from venue calendars and model open orders, cancellations, partial fills, and overlapping batches. Persist a trade reason and expected benefit for every proposed order.

CONFUSIONS TO AVOID

Comparing daily and monthly backtests without holding signal age and execution constant confounds frequency with information. Lower turnover can also be harmful when risk changes quickly or the signal decays fast.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The rebalance controller joins fresh forecasts, holdings, open orders, risk, costs, and calendars. The same controller should be replayable in simulation and production.

RELATED CONCEPTS AND QUICK RECALL

Related: turnover, hysteresis, no-trade bands, staggered rebalancing. Recall task: Design a rebalance rule for a slow signal with noisy ranks and explain its exit condition.

Transaction costs: explicit and implicit

BEGINNER

ONE-SENTENCE MEANING

Transaction costs include commissions, fees, taxes, spread, market impact, delay, borrow, financing, and opportunity loss from unfilled or constrained trades. They depend on order size, liquidity, urgency, venue, and market state.

MEMORY JOGGER

Cost is a function, not one constant basis-point deduction. Estimate it at the trade level and aggregate from simulated orders and fills.

WHY IT MATTERS IN QUANT RESEARCH

Many high-turnover strategies disappear after realistic costs, while cost-aware rebalancing can change which trades should occur. Net evidence is required for capacity and capital allocation.

FORMULA INTUITION

Expected cost may be $\text{spread}/2$ times crossed quantity plus fees plus impact $k \cdot \sigma \cdot \sqrt{Q/\text{ADV}}$ times notional, with state and participation adjustments. Exits and short financing need separate terms.

SIMPLE MARKET EXAMPLE

A 10 bp gross forecast on a thin stock is not attractive if crossing costs 8 bp and expected impact is 12 bp. A larger forecast in a liquid future may remain positive after 2 bp total burden.

PYTHON / SYSTEM IMPLEMENTATION

Calibrate on realized order data where available, segment by asset, size, participation, volatility, and session, and validate prediction intervals. Preserve a simple conservative fallback for unsupported states.

CONFUSIONS TO AVOID

Subtracting a fixed 5 bp from portfolio returns ignores turnover and direction. Using today's liquidity for old dates leaks market structure, while applying institutional costs to tiny paper capital can be equally misleading.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The cost service informs portfolio optimization, trade suppression, execution simulation, net attribution, capacity, and live monitoring. Predicted versus realized shortfall should update the model under governance.

RELATED CONCEPTS AND QUICK RECALL

Related: bid-ask spread, market impact, implementation shortfall, capacity. Recall task: Decompose one round-trip trade into spread, fees, impact, delay, borrow, and opportunity cost.

Fill models and slippage assumptions

BEGINNER

ONE-SENTENCE MEANING

A fill model estimates whether, when, how much, and at what price an order executes under historical market conditions and the strategy's participation. It translates desired trading into attainable holdings.

MEMORY JOGGER

The bar's high and low do not prove your order filled. Queue position, order type, latency, volume ahead, cancellations, and adverse selection determine the execution path.

WHY IT MATTERS IN QUANT RESEARCH

Optimistic fills can create more false alpha than the signal itself, especially for limit-order or intraday strategies. Conservative but testable models expose liquidity and operational risk.

FORMULA INTUITION

Market-order fill may be arrival mid plus half-spread plus impact; a limit order requires a touch or trade-through rule and a fill probability conditional on queue and subsequent flow. Partial quantity changes later exposure.

SIMPLE MARKET EXAMPLE

A buy limit at 100 touches the bar low of 99.95 but may not fill if little volume traded and orders were ahead. Granting a full fill at 100 and then using the rebound creates free optionality.

PYTHON / SYSTEM IMPLEMENTATION

Use quote, trade, or order-book data appropriate to the horizon; enforce latency and participation caps; and run pessimistic queue scenarios. Track unfilled quantity and stale-order cancellation explicitly.

CONFUSIONS TO AVOID

Filling at the next open can still be optimistic during gaps or auctions. Random slippage without conditioning on order and market state adds noise but not realism.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The execution simulator consumes orders and market microstructure, emits fills to the ledger, and produces shortfall diagnostics. Paper and live fills should be replayed through the same interfaces.

RELATED CONCEPTS AND QUICK RECALL

Related: arrival price, limit orders, queue position, partial fills. Recall task: Explain why a touched limit is not guaranteed and specify a conservative historical fill rule.

Liquidity, participation, and capacity

BEGINNER

ONE-SENTENCE MEANING

Capacity is the capital level at which marginal trading, financing, crowding, and exit costs leave insufficient net value or unacceptable risk. It is strategy-, universe-, horizon-, and state-dependent.

MEMORY JOGGER

Scale orders through the market, not just returns through a larger number. Capacity is determined by the most difficult trades and common exits, not average daily volume alone.

WHY IT MATTERS IN QUANT RESEARCH

A strategy can be statistically valid but economically too small, or attractive at research size but destructive when scaled. Capacity analysis turns paper edge into a capital decision.

FORMULA INTUITION

$\text{Participation}_{i,t} = |Q_{i,t}| / \text{available_volume}_{i,t}$, and impact generally rises nonlinearly with participation. Solve $\text{net_edge}(\text{capital})$ after recomputing trades, impact, unfilled quantity, risk, and holding time.

SIMPLE MARKET EXAMPLE

A small-cap rebalance uses 2 percent of daily volume at \$1 million but 20 percent at \$10 million. The larger book takes longer, moves prices, misses targets, and carries stale risk during execution.

PYTHON / SYSTEM IMPLEMENTATION

Sweep capital levels using historical volume, spreads, volatility, borrow, trading windows, and participation limits. Stress simultaneous deleveraging and declining liquidity rather than extrapolating calm averages.

CONFUSIONS TO AVOID

ADV caps alone ignore intraday concentration, order direction, crowding, and exit urgency. Capacity is not the asset size at which historical gross Sharpe remains numerically unchanged.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Capacity connects forecast decay, portfolio turnover, execution, funding, risk limits, and capital allocation. Live utilization should be compared with the approved capacity envelope.

RELATED CONCEPTS AND QUICK RECALL

Related: ADV, participation rate, nonlinear impact, liquidation horizon. Recall task: Sketch net expected return as capital grows and identify the operational break-even point.

Short selling, borrow, and financing

BEGINNER

ONE-SENTENCE MEANING

Short positions require locates, available borrow, fee payments, dividend compensation, margin, recall handling, and executable closeout. These constraints vary by security and time and can invalidate an abstract short leg.

MEMORY JOGGER

A negative target is not permission to short. Simulate availability and economics before the order, then model recalls and forced covers while the position is open.

WHY IT MATTERS IN QUANT RESEARCH

Ignoring borrow selectively overstates distressed, crowded, and small-cap strategies where paper alpha is often concentrated. Financing also changes leverage and long-short comparability.

FORMULA INTUITION

Net short return equals negative asset return minus borrow fee accrual, dividends owed, spread, impact, and financing adjustments. Forced recall adds a state-dependent closeout cost and exposure jump.

SIMPLE MARKET EXAMPLE

The lowest-ranked stock appears to fall 30 percent, but borrow was unavailable before the decline. A backtest that enters the short anyway converts an untradeable observation into profit.

PYTHON / SYSTEM IMPLEMENTATION

Use historical locate or proxy models with conservative eligibility, fee tiers, utilization, recall probability, and buy-in rules. Persist substitution logic and recompute neutrality after unavailable shorts are removed.

CONFUSIONS TO AVOID

A static annual borrow fee cannot represent hard-to-borrow events, and today's easy-borrow list is not historical evidence. Assuming proceeds can always finance longs may violate broker haircuts and margin rules.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Borrow state feeds eligibility, portfolio constraints, cost, execution, ledger liabilities, and stress testing. Live monitoring needs position-level recall and fee escalation triggers.

RELATED CONCEPTS AND QUICK RECALL

Related: locates, hard-to-borrow, recalls, margin, short financing. Recall task: Recalculate a paper short's net result after borrow, dividends, unavailable days, and a forced cover.

In-sample, validation, and out-of-sample evidence

BEGINNER

ONE-SENTENCE MEANING

Temporal validation separates model development, selection, and final evaluation so the same future outcomes do not both shape and judge the strategy. Out-of-sample evidence tests a frozen procedure on later, unseen decisions.

MEMORY JOGGER

Time has direction. Train and choose using the past, then advance the frontier and evaluate without revising the rules inside the test interval.

WHY IT MATTERS IN QUANT RESEARCH

Honest splits estimate how the research procedure generalizes through changing market conditions. They reduce the optimism created by tuning parameters and narratives on the full history.

FORMULA INTUITION

A walk-forward fold fits on times $\leq T$, optionally validates after T , and tests on a later interval after purge and embargo. Aggregate fold returns chronologically rather than randomly shuffling asset-date rows.

SIMPLE MARKET EXAMPLE

A researcher chooses lookback and threshold from 2010-2019, checks them on 2020-2022, then freezes the process for 2023-2025. Revising after seeing 2023 makes that period development evidence.

PYTHON / SYSTEM IMPLEMENTATION

Encode split boundaries in the dataset, fit every transform inside training, cache selected artifacts, and prevent test labels from reaching search dashboards. Keep a final sealed period for the full research procedure.

CONFUSIONS TO AVOID

Calling a period out-of-sample does not make it so after repeated inspection and revisions. Random cross-validation leaks regimes, overlap, and future distribution information into training.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Validation orchestrates feature fitting, model selection, backtest replay, search logs, artifact registry, and promotion. Production outcomes become new forward evidence only after their labels mature.

RELATED CONCEPTS AND QUICK RECALL

Related: walk-forward testing, temporal folds, purge, embargo, sealed test. Recall task: Design three chronological roles for a dataset and state which decisions are allowed in each.

Lookahead bias and information leakage

BEGINNER

ONE-SENTENCE MEANING

Lookahead bias occurs when a historical decision uses values, revisions, membership, transformations, or outcomes unavailable at that time. Leakage can enter through data construction, fitting, labels, universe rules, and evaluation, not only obvious future columns.

MEMORY JOGGER

Ask whether every byte and every fitted parameter existed at the decision frontier. If not, move it behind the frontier or exclude it.

WHY IT MATTERS IN QUANT RESEARCH

Leakage creates convincing, stable results because future information often cleans the past. Detecting it requires lineage and timing tests rather than visual inspection of the equity curve.

FORMULA INTUITION

For decision t , every feature input and fitted state must have availability $\leq t$. Training examples whose label windows overlap a test boundary require purging even when their feature timestamps are earlier.

SIMPLE MARKET EXAMPLE

Restated fundamentals stamped with original filing dates leak later corrections. Full-sample scaling leaks distribution shifts, while today's index membership leaks future selection.

PYTHON / SYSTEM IMPLEMENTATION

Attach availability time and vintage to source records, then fit transforms inside each fold. Adversarial tests should perturb future data and confirm earlier outputs remain unchanged.

CONFUSIONS TO AVOID

A one-period shift does not repair final-vintage data or cross-sectional leakage. Leakage-free feature code can still be invalid if target construction or model selection touches future outcomes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Leakage controls span ingestion, feature graphs, universe service, fold builder, simulator, and evaluation UI. Violations should fail the run rather than appear as warnings.

RELATED CONCEPTS AND QUICK RECALL

Related: data vintages, causal joins, target leakage, purged validation. Recall task: Name five leakage routes that survive a simple shift(1) fix.

Overlapping holdings and dependent observations

BEGINNER

ONE-SENTENCE MEANING

Overlapping signals, labels, and holding periods share returns and portfolio exposure, so observations are not independent. Dependence changes uncertainty, effective sample size, and the design of temporal splits.

MEMORY JOGGER

Count independent decision cohorts and market episodes, not rows. More assets or daily overlapping labels do not automatically create proportionally more evidence.

WHY IT MATTERS IN QUANT RESEARCH

Naive standard errors and random resampling make weak strategies appear precise. Proper dependence handling calibrates confidence and prevents train-test contamination.

FORMULA INTUITION

With autocorrelation ρ_k , effective sample size is roughly $n / (1 + 2 \sum_k \rho_k)$ under simplifying conditions. Use HAC errors, block bootstrap, or nonoverlapping cohorts appropriate to the process.

SIMPLE MARKET EXAMPLE

A monthly signal held for twelve months produces twelve overlapping portfolios. Treating 240 monthly observations as independent ignores shared returns and persistent positions.

PYTHON / SYSTEM IMPLEMENTATION

Store label start and end times, purge intersecting training events, and bootstrap contiguous time blocks or strategy episodes. Report breadth by date separately from time-series support.

CONFUSIONS TO AVOID

Asset-level rows on one date share market shocks and portfolio budgets. Clustering only by asset misses common-date dependence, while clustering only by date may miss serial overlap.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Dependence metadata feeds fold construction, statistical inference, performance reporting, and model comparison. It should be declared in the strategy contract.

RELATED CONCEPTS AND QUICK RECALL

Related: HAC errors, block bootstrap, effective sample size, overlapping labels. Recall task: Estimate how a twelve-month holding period changes the independent evidence in a monthly backtest.

Benchmarks, null models, and incremental value

BEGINNER

ONE-SENTENCE MEANING

A benchmark represents a simpler or already available decision rule against which the candidate must add value after matching universe, timing, risk, and costs. Null models test whether observed performance exceeds plausible chance or mechanical exposure.

MEMORY JOGGER

Do not ask only whether the equity curve rises. Ask what the strategy improves relative to cash, market exposure, naive allocation, existing signals, and randomized alternatives.

WHY IT MATTERS IN QUANT RESEARCH

Benchmarks reveal when complexity merely repackages beta, momentum, volatility, or rebalancing effects. Incremental evidence helps decide whether a model deserves operational cost and risk budget.

FORMULA INTUITION

Incremental return may be candidate net return minus benchmark net return at matched exposure. Null tests can shuffle scores within date, rotate signals in time, or preserve autocorrelation while breaking the hypothesized link.

SIMPLE MARKET EXAMPLE

A machine-learning model shows Sharpe 1.1, but a lagged linear rank with the same turnover and beta shows 1.0. The small increment may not justify model risk or infrastructure.

PYTHON / SYSTEM IMPLEMENTATION

Build a benchmark ladder before tuning: cash or passive, exposure-matched, naive signal, prior production model, and randomized controls. Use identical accounting and execution assumptions.

CONFUSIONS TO AVOID

Comparing a costed candidate with an uncosted benchmark biases against complexity, while comparing different universes can favor it. A factor-regression intercept is not automatically economically implementable alpha.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Benchmark services support validation, attribution, capital allocation, and retirement. Promotion should state the specific incremental capability purchased by the candidate.

RELATED CONCEPTS AND QUICK RECALL

Related: baseline models, factor exposure, permutation tests, incremental alpha. Recall task: Construct four increasingly demanding benchmarks for a new cross-sectional strategy.

Multiple testing and research search accounting

BEGINNER

ONE-SENTENCE MEANING

Every viewed signal, parameter, universe, horizon, filter, cost model, and stopping decision is part of the research search. Multiple testing raises the chance that the best-looking result is noise even when each individual backtest appears reasonable.

MEMORY JOGGER

Count the paths that could have won, including abandoned notebooks and informal charts. Evaluate the selection procedure, not only its final survivor.

WHY IT MATTERS IN QUANT RESEARCH

Search accounting calibrates statistical confidence and forces economic priors, holdouts, and replication to carry more weight. It also preserves negative knowledge and prevents repeated rediscovery of failed ideas.

FORMULA INTUITION

If m independent null trials each use threshold α , family false-positive probability is $1 - (1 - \alpha)^m$. Real trials are dependent, so use resampling, false-discovery controls, or deflated performance measures with the recorded search family.

SIMPLE MARKET EXAMPLE

Testing 100 lookback-threshold pairs and reporting the best Sharpe as if it were preselected hides the optimizer's luck. A broad stable plateau and sealed replication are stronger than one peak.

PYTHON / SYSTEM IMPLEMENTATION

Log every run automatically with code, data, parameters, parent hypothesis, metrics, and whether results were viewed. Reserve untouched periods and use nested temporal selection for the whole pipeline.

CONFUSIONS TO AVOID

Counting only published models understates the trials. Correlated variants are not free of multiplicity, and adding a new holdout after the old one disappointed does not restore independence.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The experiment tracker, hypothesis registry, validation engine, and promotion review share one search ledger. Governance should block unregistered result shopping.

RELATED CONCEPTS AND QUICK RECALL

Related: data snooping, false discovery rate, nested validation, deflated Sharpe. Recall task: Estimate the search family behind one chosen parameter and name evidence that would reduce selection concern.

Performance metrics and metric contracts

BEGINNER

ONE-SENTENCE MEANING

Performance metrics summarize return, risk, path, exposure, turnover, cost, and capacity under explicit sampling and annualization rules. No single number captures a strategy's quality or economic usefulness.

MEMORY JOGGER

Read the metric bundle as a system: return source, volatility, drawdown, tail, dependence, trading burden, support, and uncertainty. Always state units and denominator.

WHY IT MATTERS IN QUANT RESEARCH

Metric contracts prevent inconsistent annualization and selective reporting. They reveal strategies with attractive averages but unacceptable path, concentration, or implementation burden.

FORMULA INTUITION

Sharpe = $\text{mean}(\text{excess return}) / \text{sd}(\text{excess return})$ times $\sqrt{\text{periods per year}}$ only under a declared return frequency and dependence assumption. Max drawdown is the largest peak-to-trough decline in compounded NAV.

SIMPLE MARKET EXAMPLE

Two strategies have the same annual return and volatility, but one loses gradually while the other suffers a 45 percent crash and rapid recovery. Sharpe alone misses the capital and behavioral consequences.

PYTHON / SYSTEM IMPLEMENTATION

Centralize metric code, test against hand calculations, and publish sample dates, frequency, missing policy, annualizer, benchmark, and confidence interval. Include turnover, cost, exposure, and trade support.

CONFUSIONS TO AVOID

Annualizing intraday Sharpe by square-root scaling can be misleading with autocorrelation and changing exposure. Win rate ignores payoff size, and profit factor can be dominated by a few trades.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Metrics sit on reconciled ledger returns and feed reports, model comparison, risk limits, monitoring, and retirement. Dashboard definitions must match validation definitions exactly.

RELATED CONCEPTS AND QUICK RECALL

Related: Sharpe ratio, drawdown, Sortino, turnover, hit rate. Recall task: Given a strong Sharpe, list six additional diagnostics required before promotion.

Uncertainty, confidence intervals, and resampling

BEGINNER

ONE-SENTENCE MEANING

Backtest estimates are uncertain because market histories are finite, dependent, nonstationary, and selected through research. Confidence procedures should preserve the structure that generated the strategy's returns.

MEMORY JOGGER

Attach an uncertainty range to every headline estimate and explain what sources it excludes. Precision from many correlated rows is usually illusory.

WHY IT MATTERS IN QUANT RESEARCH

Calibrated uncertainty distinguishes a durable economic margin from sampling luck and informs sizing, promotion, and retirement thresholds. It also exposes when evidence is too sparse for a confident decision.

FORMULA INTUITION

A confidence interval for mean return can use HAC variance or a block bootstrap over contiguous periods. Report the sampling unit, block length, number of resamples, and whether model selection was repeated inside each sample.

SIMPLE MARKET EXAMPLE

A Sharpe estimate of 1.0 over three years may have a wide interval including weak or negative values, especially with skewed crisis profits. Ten years spanning multiple regimes can still offer limited independent episodes.

PYTHON / SYSTEM IMPLEMENTATION

Choose blocks from measured signal and return dependence, and resample dates when portfolio return is the estimand. Rerun model selection inside bootstrap loops when estimating procedure uncertainty.

CONFUSIONS TO AVOID

A narrow interval around a biased backtest is not useful. IID bootstraps break serial structure, while resampling trades can ignore overlapping positions and shared market shocks.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Uncertainty services support validation, risk budgeting, scorecards, and live drift decisions. Store random seeds and resample indices so intervals are reproducible.

RELATED CONCEPTS AND QUICK RECALL

Related: block bootstrap, HAC covariance, confidence intervals, procedure uncertainty. Recall task: Choose a resampling unit for a monthly cross-sectional strategy and defend the block length.

Robustness surfaces and sensitivity analysis

BEGINNER

ONE-SENTENCE MEANING

Robustness analysis varies reasonable data, parameters, costs, timing, universes, and implementation assumptions to determine whether the conclusion survives nearby choices. It seeks stable mechanisms, not identical performance everywhere.

MEMORY JOGGER

Prefer a defensible plateau to a sharp optimum. Test adverse but plausible assumptions that could overturn the capital decision.

WHY IT MATTERS IN QUANT RESEARCH

Sensitivity analysis reveals hidden dependencies and estimates model risk. It distinguishes genuine edge from one exact timestamp, vendor field, threshold, or favorable cost setting.

FORMULA INTUITION

For parameter vector θ , examine net performance, drawdown, turnover, and exposure over a predeclared neighborhood. Local sensitivity $\Delta \text{metric} / \Delta \theta$ complements scenario and categorical alternatives.

SIMPLE MARKET EXAMPLE

A strategy works for lookbacks from 9 to 15 months and entry delays from one to three days, but fails when costs rise 25 percent. The edge is temporally robust yet economically fragile.

PYTHON / SYSTEM IMPLEMENTATION

Generate bounded grids inside training, freeze selection, and publish the complete surface with support. Add alternative vendors, calendars, fills, corporate-action policies, and stress costs where material.

CONFUSIONS TO AVOID

Robustness testing after choosing the final result can become another search. Demanding every cell be positive is also unreasonable when hypotheses have defined operating regimes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Robustness artifacts feed model review, operating limits, monitoring thresholds, and change impact analysis. The registry should retain failed neighborhoods as negative evidence.

RELATED CONCEPTS AND QUICK RECALL

Related: parameter stability, scenario analysis, model risk, perturbation tests. Recall task: Identify five assumptions whose perturbation would most threaten a daily equity backtest.

Regime, subgroup, and conditional performance

BEGINNER

ONE-SENTENCE MEANING

Conditional analysis examines how a strategy behaves across market states, assets, sides, geographies, liquidity groups, and time without hiding the unconditional result. Conditions must be observable at the decision time if they influence trading.

MEMORY JOGGER

Slice to understand mechanisms and risk, not to rescue an average. Publish support and define every state causally before using it as a filter.

WHY IT MATTERS IN QUANT RESEARCH

Regime analysis can reveal crowding, short-side crashes, liquidity dependence, or exposure drift that aggregate metrics conceal. It guides risk limits and monitoring even when conditional timing is not reliable.

FORMULA INTUITION

For state z_t known by decision t , compare conditional net return $E[r|z]$, drawdown, turnover, exposure, and support. The weighted conditional results must reconcile to unconditional performance.

SIMPLE MARKET EXAMPLE

A strategy earns most profits in calm expansion but loses sharply during liquidity stress because its short book becomes crowded. This may justify stress limits, not a hindsight filter that disables past losses.

PYTHON / SYSTEM IMPLEMENTATION

Use lagged state definitions, minimum support, hierarchical shrinkage where appropriate, and predeclared slices. Separate descriptive diagnostics from independently validated conditional allocation rules.

CONFUSIONS TO AVOID

Defining regimes with future drawdown or full-sample clusters leaks outcomes. Searching many subgroups and reporting only winners is multiple testing in another form.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Conditional diagnostics connect mechanism, validation, stress testing, portfolio controls, and live monitoring. State labels and definitions belong in the artifact manifest.

RELATED CONCEPTS AND QUICK RECALL

Related: market regimes, subgroup analysis, state dependence, conditional sizing. Recall task: Explain why a useful stress diagnostic may still be unsuitable as a live timing filter.

Attribution: signal, sizing, constraint, and execution

BEGINNER

ONE-SENTENCE MEANING

Backtest attribution decomposes realized results into forecast selection, sizing, systematic exposures, constraints, trading, costs, and residual reconciliation. It explains how the final portfolio differed from the raw idea.

MEMORY JOGGER

Follow the chain from score to fill. A profitable book can win for the wrong reason, and a valid signal can be destroyed downstream.

WHY IT MATTERS IN QUANT RESEARCH

Attribution identifies the source of value and the owner of failures, enabling targeted improvement. It also prevents accidental beta or one event from being marketed as signal skill.

FORMULA INTUITION

One decomposition is realized net = raw forecast portfolio + sizing effect + constraint effect + execution timing effect - explicit costs - financing + residual. Components must sum to ledger P&L.

SIMPLE MARKET EXAMPLE

A stock-selection signal is flat, but the optimizer's persistent low-beta tilt profits during a selloff. Exposure attribution shows the backtest is not evidence for the stated selection mechanism.

PYTHON / SYSTEM IMPLEMENTATION

Persist score, forecast, raw target, constrained target, order, fill, holding, exposure, and cost at each decision. Reconcile daily components to the ledger and investigate unexplained residuals.

CONFUSIONS TO AVOID

Factor regressions alone do not show implementation effects, and arithmetic decompositions can become order-dependent. Attribution is explanatory evidence, not proof of causality.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Attribution joins model outputs, optimizer diagnostics, execution, risk exposures, and accounting. The same component definitions should appear in research and production reporting.

RELATED CONCEPTS AND QUICK RECALL

Related: performance attribution, constraint cost, implementation shortfall, factor exposure. Recall task: Build an additive bridge from a raw signal portfolio to realized net P&L.

Reproducibility, lineage, and immutable artifacts

BEGINNER

ONE-SENTENCE MEANING

A reproducible backtest binds code, configuration, data vintages, calendars, environments, random seeds, and outputs so the exact result can be regenerated and audited. Lineage records how every reported number was produced.

MEMORY JOGGER

Save the evidence and procedure, not just the PDF or equity curve. A result without immutable dependencies cannot support a capital decision.

WHY IT MATTERS IN QUANT RESEARCH

Lineage enables independent replication, incident analysis, vendor-change assessment, and rollback. It also prevents silent drift when databases or library defaults change.

FORMULA INTUITION

Run ID = hash(code commit, manifest, data snapshot IDs, environment lock, seed, engine version). The hash is a locator and integrity check, not a substitute for storing human-readable metadata.

SIMPLE MARKET EXAMPLE

A vendor revises historical fundamentals and the old query now returns different rows. An immutable snapshot or vintage pointer lets the team reproduce the approved run and quantify the revision's effect.

PYTHON / SYSTEM IMPLEMENTATION

Write manifests before execution, use content-addressed artifacts, lock dependencies, capture query and snapshot IDs, and make outputs append-only. Include deterministic seeds and stable sort rules.

CONFUSIONS TO AVOID

A Git commit alone does not version external data, database state, calendar packages, or environment. Re-running a mutable notebook and obtaining a similar Sharpe is not exact replication.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The artifact registry links hypothesis, datasets, features, simulator, validation, approvals, deployment, and incidents. Promotion references one immutable graph of dependencies.

RELATED CONCEPTS AND QUICK RECALL

Related: data lineage, content hashes, environment lockfiles, experiment tracking. Recall task: Name every dependency needed to reproduce a backtest five years from now.

Backtest-to-production parity

BEGINNER

ONE-SENTENCE MEANING

Backtest-to-production parity means research and live systems share semantic contracts for data, features, portfolio state, scheduling, costs, and accounting, with differences isolated and measured. The goal is equivalent decisions under equivalent inputs.

MEMORY JOGGER

Replay live events through the research engine and compare at each boundary. Final P&L differences are too late and too aggregated to diagnose parity.

WHY IT MATTERS IN QUANT RESEARCH

Parity reduces deployment surprises from different timestamps, rounding, universe rules, or state handling. It converts research assumptions into testable production interfaces.

FORMULA INTUITION

For identical input event IDs and prior state, compare feature values, forecasts, targets, orders, fills-as-observed, holdings, and ledger entries within declared tolerances. Hash each state transition.

SIMPLE MARKET EXAMPLE

Research uses adjusted end-of-day files and fractional weights, while production receives raw events and whole-share orders. A shadow replay quantifies differences from adjustment timing, rounding, latency, and fills.

PYTHON / SYSTEM IMPLEMENTATION

Package shared libraries or specifications, run golden-event tests, shadow decisions, and daily reconciliation. Record intentional differences such as live safety caps and actual broker fills.

CONFUSIONS TO AVOID

Identical code does not guarantee identical inputs or state. Forcing research to mimic every production outage can also obscure the economic model; isolate operational scenarios explicitly.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Parity testing spans ingestion, feature service, optimizer, scheduler, OMS, execution, ledger, and monitoring. Deployment gates should require resolved boundary-level differences.

RELATED CONCEPTS AND QUICK RECALL

Related: shadow mode, golden datasets, replay, semantic parity. Recall task: Define a parity report that localizes differences before they reach portfolio P&L.

Backtest QA, invariants, and adversarial tests

BEGINNER

ONE-SENTENCE MEANING

Backtest quality assurance uses deterministic invariants, hand-worked fixtures, property tests, differential engines, and adversarial cases to catch errors before interpreting performance. Tests target economic semantics as well as code correctness.

MEMORY JOGGER

Try to break the simulator with zero signals, impossible prices, splits, halts, partial fills, missing data, and future-data perturbations. A realistic equity curve is not a test result.

WHY IT MATTERS IN QUANT RESEARCH

Simulation bugs often produce plausible outputs and survive ordinary unit tests. Economic invariants and independent calculations expose silent leverage, cash, timing, and adjustment errors.

FORMULA INTUITION

Core invariants require NAV reconciliation and forbid position changes without fills or actions. Features cannot use future availability, constraints must hold, and a flat unchanged portfolio earns zero P&L before costs.

SIMPLE MARKET EXAMPLE

A one-asset fixture buys ten shares at 100, pays a \$1 fee, receives a \$2 dividend, and marks at 103. The expected ledger and return are calculated by hand and compared exactly.

PYTHON / SYSTEM IMPLEMENTATION

Create small golden scenarios, randomized property tests, cross-engine comparisons, and future-data mutation tests. Fail fast on unexplained cash, duplicate events, negative impossible volume, or constraint violations.

CONFUSIONS TO AVOID

Snapshot tests can preserve a wrong answer indefinitely. Testing only helper functions misses errors in event ordering, ledger integration, and changing calendars.

WHERE THIS FITS IN THE RESEARCH SYSTEM

QA gates the simulator library, data adapters, strategy plugins, reports, and production replay. Every bug becomes a permanent regression fixture with documented economic impact.

RELATED CONCEPTS AND QUICK RECALL

Related: property testing, golden fixtures, invariants, differential testing. Recall task: Design one minimal fixture that simultaneously tests timing, a partial fill, fees, and NAV reconciliation.

Promotion, monitoring, and retirement

BEGINNER

ONE-SENTENCE MEANING

Promotion authorizes one immutable procedure and operating envelope. Monitoring tests live state against that envelope, and retirement removes capital when evidence or control quality no longer supports the claim.

MEMORY JOGGER

Approve a governed system, not a historical Sharpe. Define warnings, degraded operation, hard stops, rollback, and retirement before deployment.

WHY IT MATTERS IN QUANT RESEARCH

Governance prevents attractive backtests from bypassing causality, feasibility, tail risk, or operations. It also ensures changes receive validation proportional to their semantic impact.

FORMULA INTUITION

Promotion thresholds cover independent net evidence, drawdown, capacity, stress loss, parameter stability, data parity, and control readiness. Monitoring compares live distributions and realized shortfall with validated ranges.

SIMPLE MARKET EXAMPLE

A spread vendor changes methodology, live turnover jumps, and cost error exceeds tolerance. The system degrades capital, freezes model updates, replays affected decisions, and requires revalidation before restart.

PYTHON / SYSTEM IMPLEMENTATION

Use signed manifests, approval records, canary or shadow runs, health tiers, incident playbooks, rollback artifacts, and retirement triggers. Preserve prior versions and negative evidence.

CONFUSIONS TO AVOID

Automatically refitting after losses can absorb data defects or regime noise. Widening limits to keep a strategy alive converts governance into live curve fitting.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The governance layer connects registry, validation, data, portfolio risk, execution, operations, capital owners, and incident management. Every order inherits the approved version and current health state.

RELATED CONCEPTS AND QUICK RECALL

Related: model risk, deployment gates, drift monitoring, rollback, retirement. Recall task: Name three live conditions that warn, de-risk, and stop a strategy without waiting for long-run P&L.

From historical evidence to governed capital

BEGINNER

CORE LOGIC

A defensible backtest is a chain of causal and economic state transitions. The earliest broken boundary invalidates every attractive result downstream, so validation follows the same direction as the simulator.

POINT-IN-TIME EVIDENCE

identity, vintages, membership, actions, prices, calendars, borrow and quality

CAUSAL DECISION

availability cutoff, features, forecast, state, schedule and constraints

EXECUTABLE ORDER PATH

latency, order type, participation, queue, partial fills and cancellations

ECONOMIC PORTFOLIO

holdings, cash, corporate actions, financing, marks and reconciled NAV

INDEPENDENT EVIDENCE

temporal selection, baselines, search, dependence, uncertainty and stress

GOVERNED OPERATION

parity, lineage, monitoring, incident response, rollback and retirement

DIAGNOSTIC QUESTION

Can every reported return be traced backward to a lawful mark, an actual holding, a simulated fill, an admissible order, a causal decision, and point-in-time evidence?

Worked three-day backtest and ledger reconciliation

BEGINNER

SCENARIO

The strategy starts with \$10,000 cash, buys 50 shares, receives a dividend, and sells part of the position. Fees and marks are applied explicitly so NAV return can be reconstructed from economic events.

Event	Qty	Cash flow	Held	Cash	Mark	NAV
Start	-	-	0	\$10,000	\$100.00	\$10,000
Day 1 buy	+50	-\$5,005	50	\$4,995	\$101.00	\$10,045
Day 2 dividend	0	+\$50	50	\$5,045	\$99.00	\$9,995
Day 3 sell	-20	+\$2,058	30	\$7,103	\$103.00	\$10,193

DAY 1

The buy costs $50 \times \$100$ plus a \$5 fee, leaving \$4,995 cash. Marking 50 shares at \$101 gives $\text{NAV} = \$4,995 + \$5,050 = \$10,045$ and a 0.45 percent return.

DAY 2

The \$1 dividend adds \$50 cash while the price falls to \$99. NAV is $\$5,045 + \$4,950 = \$9,995$, so the dividend offsets part of the price decline without being double counted.

DAY 3

Selling 20 shares at \$103 with a \$2 fee raises cash by \$2,058. The remaining 30 shares are worth \$3,090, producing ending NAV of \$10,193 and total return of 1.93 percent.

RECONCILIATION


Ending NAV minus starting NAV is \$193. The result includes the initial \$5 fee, the \$50 dividend, the \$2 sale fee, price changes, and the residual holding. No position changed without a fill, and no cash appeared without an event.

Causal signal, order, fill, and label timeline

BEGINNER

SCENARIO

An end-of-day strategy uses a closing price and a post-close fundamental update. The decision forms after ingestion and computation, the order reaches the next session, and the evaluation label starts at the simulated fill.



OBSERVATION VERSUS AVAILABILITY

The fiscal quarter may end months earlier, but the report is not available until publication and ingestion. The closing price exists only after the 16:00 market event is complete.

DECISION FRONTIER

Feature computation uses records whose availability is no later than 17:00. Universe, transform state, risk inputs, and open positions are frozen into one decision snapshot.

EXECUTABLE ENTRY

The order cannot fill at the already-known close. Auction rules, overnight news, spread, latency, and participation determine the next lawful execution price.

LABEL CONTRACT

Forward return starts from the simulated fill and includes dividends, actions, delistings, and exit cost. The decision is excluded from training folds whose test periods intersect this label.

AUDIT TEST

Mutating any data that arrives after 17:00 must not change this decision. Changing next-open conditions may change the fill and P&L, but never the archived signal.

Signal-to-holding state machine

BEGINNER

PURPOSE

Separating belief, intent, instruction, fact, and exposure prevents perfect-fill assumptions. Each transition is controlled by different data and can explain a gap between research expectation and realized outcome.



NEXT-DECISION RULE

The next order begins from actual holdings plus still-open orders, not from the previous target. Partial fills, cancellations, corporate actions, price moves, and risk changes all alter the next feasible transition.

Gross backtest edge to executable net return

BEGINNER

SCENARIO

A monthly strategy shows 78 basis points of modeled gross return. Delay, spread, impact, fees, financing, borrow, and an exit reserve reduce the expected net result at proposed capital.

78.0 bp

Gross modeled return

reconciled return before implementation burden

-6.0 bp

Signal-to-order decay

value lost before executable entry

-10.0 bp

Bid-ask spread

crossing on entries, exits, and rebalances

-14.0 bp

Market impact

nonlinear cost at planned participation

-3.0 bp

Fees and taxes

commissions, venue fees, and transaction tax

-8.0 bp

Financing and borrow

leverage, short fee, dividends, and haircuts

-7.0 bp

Unfilled opportunity

capacity limits and partial executions

-12.0 bp

Stress exit reserve

common liquidation under thinner depth

+18.0 bp

Expected net return

remaining margin before model uncertainty

DECISION

The 18 bp margin must survive parameter uncertainty and live cost error. Capital should be reduced or trades suppressed when forecast value does not clear the current burden plus safety buffer.


Equity, drawdown, and path interpretation

BEGINNER

SYNTHETIC PATH

The strategy compounds steadily, suffers a sharp loss, partially recovers, and later reaches a new high. The same final return could arise from many paths with different capital, liquidity, and behavioral consequences.

NAV



READ THE PATH

NAV shows compounded wealth; drawdown measures distance below the running peak. Recovery time, loss speed, underwater duration, exposure, and liquidity during the trough matter alongside the maximum percentage.

DIAGNOSTIC NEXT STEP

Attribute the drawdown by signal leg, market exposure, constraint, cost, fill, and data state. Then replay the same episode with delayed execution, lower depth, higher borrow, and forced deleveraging.

Lookback and rebalance robustness

BEGINNER

SYNTHETIC TRAINING SURFACE

Cells combine net return, drawdown, turnover, and support across signal lookback and rebalance frequency. The goal is a coherent neighborhood, not the single highest cell.



SELECTION

Sixty- to 120-day signals with biweekly or monthly rebalancing form a stable region after costs. Choose a simple representative inside nested training and evaluate the frozen selection once on sealed evidence.

CAUTION

Every cell is a trial and neighboring cells share observations. Publish the full surface with turnover, drawdown, exposure, trade support, cost error, and outer-fold behavior rather than promoting one color.

Backtest behavior across operating states

BEGINNER

SYNTHETIC DIAGNOSTIC

Rows describe observable market states and columns separate forecast, execution, funding, and loss behavior. The unconditional result remains visible; adverse states are not filtered away.

	Signal	Turnover	Execution	Funding	Tail loss
Quiet trend	GOOD	LOW	EASY	EASY	LOW
Range market	MIXED	HIGH	MIXED	EASY	MEDIUM
Volatility shock	MIXED	HIGH	HARD	MIXED	HIGH
Liquidity stress	WEAK	HIGH	SEVERE	HARD	HIGH
Short squeeze	MIXED	HIGH	HARD	SEVERE	SEVERE
Data outage	UNKNOWN	MIXED	BLOCKED	MIXED	UNKNOWN

INTERPRETATION

Calm periods support the base case, but stress damages execution and funding before average signal diagnostics fully react. A data outage is an operating state requiring fallback or stop, not a historical return regime.

EVIDENCE RULE

Define states with lagged information, publish support, and validate any conditional allocation independently. Sparse severe states require constructed scenarios and broker or venue evidence in addition to history.

Failure modes and first diagnostics

BEGINNER

EQUITY IMPROVES AFTER A DATA REFRESH

Historical revisions, survivor rosters, adjusted prices, or classification changes may have leaked into earlier decisions. Compare vintages and rerun using immutable evidence.

SIGNAL IS STRONG BUT NAV IS WEAK

Sizing, constraints, partial fills, turnover, costs, borrow, or unintended exposure may destroy forecast value. Bridge score through target, order, fill, holding, and ledger.

NAV CHANGES WITHOUT AN EVENT

The ledger or valuation path is broken. Reconcile cash, quantity, marks, actions, liabilities, fees, and external flows before computing any metric.

LIMIT STRATEGY LOOKS IMPOSSIBLY SMOOTH

Touched bars may be granted full favorable fills without queue or adverse selection. Replay with latency, trade-through rules, participation, and partial fills.

SHARPE RISES WITH MORE ROWS

Overlapping labels, shared market shocks, or pooled assets can inflate precision. Measure effective time support and use block or clustered uncertainty.

ONLY ONE PARAMETER WORKS

The result is likely selection noise or a sparse event. Publish the full neighborhood, search family, outer folds, and economic mechanism.

SHORT LEG DOMINATES PROFIT

Check historical borrow, fees, recalls, delistings, dividends owed, crowding, and forced exit. Unavailable shorts cannot be counted as paper alpha.

LIVE TURNOVER EXCEEDS RESEARCH

Roster churn, stale state, rounding, open-order handling, data timing, or optimizer differences may be responsible. Compare boundary-level parity hashes.

REPAIR PRINCIPLE

Locate the earliest broken boundary: evidence, clock, feature, target, order, fill, holding, ledger, validation, or governance. Downstream tuning cannot repair an upstream defect.

Validation and promotion checklist

BEGINNER

STRATEGY CONTRACT

Hypothesis, universe, clocks, signal, sizing, rebalance, constraints, execution, costs, benchmark, metrics, falsifiers, and limits are explicit and versioned.

POINT-IN-TIME EVIDENCE

Identity, membership, vintages, prices, corporate actions, calendars, liquidity, borrow, risk inputs, and availability times reproduce every decision.

SIMULATION STATE

Signal, target, open order, fill, holding, cash, action, valuation, liability, and NAV transitions are distinct, deterministic, and auditable.

ECONOMIC RECONCILIATION

Every quantity and cash change has an event; return definitions, financing, fees, benchmark, and attribution sum exactly to the ledger.

SEARCH ACCOUNTING

All parameters, universes, filters, horizons, fills, costs, regimes, and viewed subgroups belong to the declared research family.

INDEPENDENT VALIDATION

Chronological selection, purge, embargo, simple baselines, dependence-aware uncertainty, sealed evidence, and replication support the claim.

NET IMPLEMENTATION

Delay, spread, fees, impact, participation, partial fills, borrow, financing, capacity, and stressed exit leave a positive margin at proposed size.

ROBUSTNESS AND RISK

Parameter neighborhoods, vendor and timing alternatives, drawdowns, exposures, concentration, regimes, liquidity shocks, and common liquidation are understood.

OPERATIONS

Immutable lineage, golden fixtures, batch-live parity, shadow mode, health tiers, owners, incident replay, rollback, change control, and retirement have been tested.

HARD STOP

Future knowledge, survivor-only data, unreconciled NAV, impossible fills, hidden search, unsupported capacity, unavailable shorts, or absent rollback blocks capital.

Final active recall and system index

BEGINNER

DEFINE

Write the complete historical decision contract: evidence, clocks, universe, signal, sizing, rebalance, constraints, execution, costs, accounting, benchmark, metrics, and falsifiers.

RECONSTRUCT

Freeze point-in-time identity, membership, vintages, actions, prices, calendars, liquidity, borrow, and risk inputs. Preserve missing and terminal events.

SIMULATE

Keep signal, target, order, fill, holding, cash, liability, mark, and NAV as separate states. Begin each decision from actual portfolio state.

RECONCILE

Prove every position and cash movement, calculate return from the ledger, and bridge raw forecast value through portfolio and execution into realized net P&L.

VALIDATE

Use temporal selection, baselines, search accounting, purge, dependence-aware uncertainty, complete parameter surfaces, regimes, and sealed forward evidence.

IMPLEMENT

Model latency, spread, fees, impact, participation, queue, partial fills, borrow, financing, capacity, and stressed common exit at proposed capital.

TEST

Use hand-worked fixtures, invariants, property tests, future-data mutation, differential engines, golden events, and permanent regression cases.

OPERATE

Bind immutable artifacts, prove backtest-live parity, shadow decisions, monitor data and execution, assign owners, and rehearse fallback and rollback.

REJECT

Stop when a result depends on future data, current survivors, one sharp parameter, pooled-row precision, fantasy fills, missing ledger events, or discretionary repairs.

RETAIN NEGATIVE KNOWLEDGE

Preserve failed strategies, broken assumptions, search history, incidents, and adverse states. Revisit only when genuinely new evidence supports a registered hypothesis.