

Probability Foundations for Quant Research

Events, distributions, probability updates, finite-data estimates, and probabilistic decisions.

Library of the Quant

Educational reference manual

Manual Profile

LIBRARY ID	LOTQ-0024
CATEGORY	Probability
DIFFICULTY	Beginner
FORMAT	Single-column field-guide manual
USE	Study, lookup, and research-system reference

Educational Use Only

This manual is educational material for quant research study. It is not financial advice, not a trading recommendation, and not a guarantee of correctness or live trading usefulness.

Table of contents

BEGINNER

Title	1
Table of contents I	2
Table of contents II	3
How to use this manual	4
Probability notation	5
Probability-to-decision workflow	6
Probability as a model of uncertainty	7
Experiments, outcomes, and sample spaces	8
Events and set operations	9
Probability axioms	10
Equally likely outcomes and counting	11
Complement rule	12
Addition rule and overlap	13
Conditional probability	14
Multiplication rule	15
Joint, marginal, and conditional views	16
Independence	17
Mutual exclusivity versus independence	18
Pairwise and mutual independence	19
Partitions and the law of total probability	20
Bayes' rule	21
Odds and likelihood ratios	22
Random variables	23
Discrete versus continuous variables	24
Probability mass functions	25
Probability density functions	26

Table of contents

BEGINNER

Cumulative distribution functions	27
Survival functions and tail probabilities	28
Quantiles and percentiles	29
Expectation	30
Linearity of expectation	31
Variance and standard deviation	32
Covariance	33
Correlation	34
Transformations of random variables	35
Indicator variables	36
Mixture distributions	37
Convolution and sums	38
Probability generating and moment summaries	39
Law of large numbers	40
Central limit theorem	41
Monte Carlo probability estimation	42
Sampling error and effective sample size	43
Calibration of probabilistic forecasts	44
Proper scoring rules	45
Base rates and class imbalance	46
Rare events and tail estimation	47
Probability trees and path dependence	48
Probability versus frequency	49
Model probability versus real-world uncertainty	50
Probability-to-decision translation	51
Probability tree visual lab	52
Distribution views visual lab	53
Monte Carlo convergence visual lab	54

How to use this manual

BEGINNER

1. DEFINE THE EVENT FIRST

Write the outcome, horizon, universe, timestamp, and information set before assigning a number. An undefined event cannot have an auditable probability.

2. READ EVERY DENOMINATOR

For conditional probabilities and empirical rates, identify the eligible cohort. Most interpretation errors are denominator or conditioning errors.

3. SEPARATE MODEL FROM EVIDENCE

A probability is a model statement; a frequency is observed evidence. Use the second to estimate and validate the first without confusing them.

4. TRACK SUPPORT AND DEPENDENCE

Record event counts, effective sample size, serial dependence, and regimes. Large row counts can still contain little independent information.

5. VALIDATE PROBABILITY QUALITY

Check normalization, calibration, proper scoring rules, uncertainty, tail support, and stability before probabilities drive sizing or alerts.

6. TRANSLATE INTO CONSEQUENCES

Combine probability with payoff magnitude, transaction costs, risk limits, and abstention. A high probability is not automatically a good trade.

DECISION STANDARD

Use a probabilistic estimate only when its event definition, information cutoff, calibration evidence, uncertainty, and downstream decision rule can be reproduced.

Probability notation and reading discipline

BEGINNER

SYMBOL	TYPE	MEANING
Ω	set	Sample space of modeled outcomes
ω	outcome	One element of the sample space
A, B	events	Subsets of Ω
$P(A)$	probability	Mass assigned to event A
$P(A B)$	conditional	Probability of A within B
X, Y	random variables	Numerical functions of outcomes
$p_X(x)$	PMF	Point mass for discrete X
$f_X(x)$	PDF	Density per unit for continuous X
$F_X(x)$	CDF	Probability X is at most x
$E[X], \text{Var}(X)$	moments	Center and dispersion summaries

CONDITIONING AUDIT

Read $P(A|B)$ as: among outcomes satisfying B , what fraction also satisfy A ? The conditioning event changes the reference population. Record B 's support and never reverse the expression without Bayes' rule.

PROBABILITY VERSUS DENSITY

$P(X=x)$ may be positive for a discrete variable. For a continuous variable it is typically zero; probability comes from integrating f_X over an interval. Density values depend on the measurement unit.

Probability-to-decision research workflow

BEGINNER

1

Define outcome

Universe, event, horizon, information cutoff

2

Build evidence

Point-in-time data, labels, eligible sample

3

Choose model

Support, dependence, distribution, prior

4

Estimate

Parameters, probabilities, uncertainty

5

Validate

Calibration, proper scores, stability, tails

6

Stress

Alternative models, regimes, assumptions

7

Decide

Payoffs, costs, constraints, abstention

8

Monitor

Drift, realized frequency, recalibration

HARD GATE

No probability enters capital allocation until it is normalized, timestamp-correct, out-of-sample validated, uncertainty-aware, and linked to a payoff-and-cost decision rule.

Probability as a model of uncertainty

BEGINNER

ONE-SENTENCE MEANING

Probability is a consistent language for assigning numerical weight to uncertain outcomes under a stated model and information set.

MEMORY JOGGER

A probability is conditional on assumptions, data, time, and what is known - even when that conditioning is not written.

WHY IT MATTERS IN QUANT RESEARCH

Quant research compares uncertain payoffs, estimates risk, combines evidence, and designs decisions whose outcomes cannot be known in advance.

FORMULA INTUITION

$0 \leq P(A) \leq 1$; $P(\Omega) = 1$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A model may assign a 55% chance that tomorrow's close exceeds today's close; this is not a promise that 55 of the next 100 days must rise.

PYTHON / SYSTEM IMPLEMENTATION

Represent each forecast with event definition, horizon, information cutoff, model version, and calibration evidence; store probabilities separately from realized labels.

CONFUSIONS TO AVOID

Treating probability as certainty or as a physical property independent of the model hides estimation error and regime dependence.

WHERE THIS FITS IN THE RESEARCH SYSTEM

This is the contract layer of the research system: it defines what an uncertain forecast means before features, models, sizing, or risk limits use it.

RELATED CONCEPTS AND QUICK RECALL

Related: sample space, event, calibration. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Experiments, outcomes, and sample spaces

BEGINNER

ONE-SENTENCE MEANING

A random experiment is a repeatable uncertainty-producing procedure; an outcome is one possible result; the sample space contains every modeled outcome.

MEMORY JOGGER

Define the universe of possible results before assigning probabilities.

WHY IT MATTERS IN QUANT RESEARCH

Clear sample spaces prevent ambiguous labels, missing states, and inconsistent denominators in market-event studies and simulations.

FORMULA INTUITION

$\Omega = \{\omega_1, \omega_2, \dots\}$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

For a next-day sign experiment, Ω might be {up, flat, down}; excluding flat days silently changes every reported probability.

PYTHON / SYSTEM IMPLEMENTATION

Encode mutually exclusive outcome labels, an explicit unknown state, timestamp rules, and a validation test that every observation maps to exactly one state.

CONFUSIONS TO AVOID

Using future information to decide which outcomes were possible creates survivorship and look-ahead bias in the modeled sample space.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The ontology sits between cleaned market data and statistical estimation; downstream probabilities are valid only for the versioned outcome map.

RELATED CONCEPTS AND QUICK RECALL

Related: event partition, state space, label schema. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Events and set operations

BEGINNER

ONE-SENTENCE MEANING

An event is a subset of the sample space, and unions, intersections, and complements combine market conditions precisely.

MEMORY JOGGER

OR is union, AND is intersection, NOT is complement.

WHY IT MATTERS IN QUANT RESEARCH

Event algebra supports compound filters, trigger definitions, scenario counts, and exact interpretation of backtest conditions.

FORMULA INTUITION

A union B; A intersection B; A^c . Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Let A be a positive return and B be high volume. A intersection B is a positive-return, high-volume day; A union B includes either condition.

PYTHON / SYSTEM IMPLEMENTATION

Build boolean masks from atomic predicates, test precedence with parentheses, and log counts for each component, overlap, union, and complement.

CONFUSIONS TO AVOID

Confusing inclusive OR with exclusive OR undercounts cases where both conditions occur.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Event definitions belong in a reusable feature-and-label registry so research, monitoring, and execution evaluate the same logic.

RELATED CONCEPTS AND QUICK RECALL

Related: Boolean algebra, Venn diagram, compound condition. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability axioms

BEGINNER

ONE-SENTENCE MEANING

The probability axioms require nonnegative probabilities, total mass one, and additive mass for disjoint events.

MEMORY JOGGER

The rules are simple; most useful probability identities follow from them.

WHY IT MATTERS IN QUANT RESEARCH

They provide sanity checks for forecast vectors, transition matrices, scenario weights, and simulation outputs.

FORMULA INTUITION

$P(A) \geq 0$; $P(\Omega) = 1$; $P(\text{union } A_i) = \sum P(A_i)$ for disjoint A_i . Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If up, flat, and down are mutually exclusive and exhaustive, their probabilities cannot be negative and must sum to one.

PYTHON / SYSTEM IMPLEMENTATION

Assert finite values, nonnegativity within tolerance, normalized row sums, and mutually exclusive class mappings before any decision consumes probabilities.

CONFUSIONS TO AVOID

Renormalizing invalid negative scores may hide a modeling bug rather than create legitimate probabilities.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Use these checks as hard data-quality gates at model output, scenario generation, and risk aggregation boundaries.

RELATED CONCEPTS AND QUICK RECALL

Related: normalization, additivity, probability mass. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Equally likely outcomes and counting

BEGINNER

ONE-SENTENCE MEANING

When finite outcomes are genuinely symmetric, event probability is favorable outcomes divided by all possible outcomes.

MEMORY JOGGER

Counting works only after defending equal likelihood.

WHY IT MATTERS IN QUANT RESEARCH

It helps with permutation tests, randomized controls, combinatorial scenarios, and simple benchmark simulations.

FORMULA INTUITION

$P(A) = |A| / |\Omega|$ under equal likelihood. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If a day is randomly selected from 20 labeled days, 7 of which are positive, the empirical selection probability is 7/20.

PYTHON / SYSTEM IMPLEMENTATION

Use exact combinatorial functions or stable log-counts for large spaces; document the randomization mechanism that justifies symmetry.

CONFUSIONS TO AVOID

Market paths are rarely equally likely; counting path shapes without weights creates a misleading probability model.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Counting is useful in the experiment-design service for randomized resampling, not as a default model of price dynamics.

RELATED CONCEPTS AND QUICK RECALL

Related: permutation, combination, symmetry. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Complement rule

BEGINNER

ONE-SENTENCE MEANING

The complement rule obtains an event probability by subtracting the probability of everything outside it from one.

MEMORY JOGGER

Sometimes it is easier to count failure than success.

WHY IT MATTERS IN QUANT RESEARCH

It simplifies at-least-one calculations, breach probabilities, uptime estimates, and quality-control checks.

FORMULA INTUITION

$P(A^c) = 1 - P(A)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If the modeled chance of no limit breach tomorrow is 0.97, the chance of at least one breach is 0.03.

PYTHON / SYSTEM IMPLEMENTATION

Compute the better-conditioned side directly when probabilities are extremely small, using log survival functions when available.

CONFUSIONS TO AVOID

Subtracting a rounded probability near one can lose precision and materially distort rare-event estimates.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The risk-monitoring layer often stores both breach probability and survival probability so alerts can be checked from either direction.

RELATED CONCEPTS AND QUICK RECALL

Related: survival probability, tail event, numerical precision. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Addition rule and overlap

BEGINNER

ONE-SENTENCE MEANING

The addition rule adds two event probabilities and subtracts their overlap so shared outcomes are counted once.

MEMORY JOGGER

Add, then remove the double count.

WHY IT MATTERS IN QUANT RESEARCH

It measures the chance that at least one signal, risk flag, or market condition is active without assuming exclusivity.

FORMULA INTUITION

$P(A \cup B) = P(A) + P(B) - P(A \cap B)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If momentum fires on 30% of days, volatility fires on 25%, and both fire on 10%, either fires on 45%.

PYTHON / SYSTEM IMPLEMENTATION

Materialize A, B, A&B, and A|B masks from the same timestamped panel; reconcile union counts with the identity before reporting.

CONFUSIONS TO AVOID

Simply summing trigger rates can exceed one when signals overlap substantially.

WHERE THIS FITS IN THE RESEARCH SYSTEM

This logic belongs in signal-combination analytics, where overlap explains capacity, diversification, and alert volume.

RELATED CONCEPTS AND QUICK RECALL

Related: intersection, inclusion-exclusion, signal overlap. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Conditional probability

BEGINNER

ONE-SENTENCE MEANING

Conditional probability updates an event's probability after restricting attention to cases where another event occurred.

MEMORY JOGGER

Change the denominator to the information-conditioned subset.

WHY IT MATTERS IN QUANT RESEARCH

It quantifies regime-specific hit rates, post-event risk, transition behavior, and the meaning of predictive features.

FORMULA INTUITION

$P(A|B) = P(A \text{ intersection } B) / P(B)$, $P(B) > 0$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If 40 of 100 days are high-volatility and 26 of those have negative returns, $P(\text{negative}|\text{high vol}) = 26/40 = 65\%$.

PYTHON / SYSTEM IMPLEMENTATION

Compute numerator and denominator from one point-in-time cohort; report support, uncertainty, and the exact conditioning rule beside the estimate.

CONFUSIONS TO AVOID

Reversing $P(A|B)$ and $P(B|A)$ is a base-rate error, not a small notation mistake.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Conditional estimates feed model diagnostics and regime routing, but should be blocked when support falls below a declared threshold.

RELATED CONCEPTS AND QUICK RECALL

Related: Bayes rule, contingency table, regime. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Multiplication rule

BEGINNER

ONE-SENTENCE MEANING

The multiplication rule expresses a joint probability as a conditional probability times the probability of the conditioning event.

MEMORY JOGGER

Joint equals branch probability times probability of reaching the branch.

WHY IT MATTERS IN QUANT RESEARCH

It builds scenario trees, path likelihoods, default-event chains, and multistage execution models.

FORMULA INTUITION

$P(A \text{ intersection } B) = P(A|B)P(B)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If high volatility occurs with probability 0.20 and a loss given high volatility occurs with probability 0.60, their joint probability is 0.12.

PYTHON / SYSTEM IMPLEMENTATION

For event trees, multiply along each path and sum mutually exclusive terminal paths; validate that all terminal probabilities total one.

CONFUSIONS TO AVOID

Multiplying marginal probabilities assumes independence and is wrong when the second event depends on the first.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The scenario engine uses this identity to turn local transition probabilities into full path weights for stress and planning.

RELATED CONCEPTS AND QUICK RECALL

Related: joint probability, chain rule, scenario tree. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Joint, marginal, and conditional views

BEGINNER

ONE-SENTENCE MEANING

A joint distribution describes variables together; marginals summarize one variable; conditionals describe one given another.

MEMORY JOGGER

Joint is the full table, marginal is a row or column sum, conditional is a normalized slice.

WHY IT MATTERS IN QUANT RESEARCH

Moving among these views clarifies dependence, regime behavior, portfolio co-movement, and feature-target relationships.

FORMULA INTUITION

$P(X=x) = \sum_y P(X=x, Y=y)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A return-sign by volatility-state table can yield the overall loss rate, the high-volatility frequency, and loss probability within high volatility.

PYTHON / SYSTEM IMPLEMENTATION

Store tidy contingency tables with counts and weights, then derive marginals and conditionals from the same artifact to prevent denominator drift.

CONFUSIONS TO AVOID

Comparing conditionals built from different universes or time windows produces apparent relationships that may be purely compositional.

WHERE THIS FITS IN THE RESEARCH SYSTEM

This table is a compact diagnostic artifact between exploratory analysis and formal modeling, suitable for registry snapshots and monitoring.

RELATED CONCEPTS AND QUICK RECALL

Related: contingency table, dependence, aggregation. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Independence

BEGINNER

ONE-SENTENCE MEANING

Two events are independent when learning that one occurred does not change the probability of the other under the model.

MEMORY JOGGER

Independent means no probability update, not merely low correlation.

WHY IT MATTERS IN QUANT RESEARCH

Independence assumptions simplify simulations and risk aggregation, but false independence severely understates clustered market risk.

FORMULA INTUITION

$P(A \text{ intersection } B) = P(A)P(B)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If two venue outages are independent at 1% each, simultaneous outage is 0.01%; common infrastructure can make the true rate much higher.

PYTHON / SYSTEM IMPLEMENTATION

Compare joint frequency with the product of marginals across time and regimes; use domain structure and uncertainty, not a single test alone.

CONFUSIONS TO AVOID

Failure to reject dependence is not proof of independence, especially with small samples or nonlinear relationships.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Independence assumptions must be explicit configuration in scenario and portfolio engines, with stress alternatives for common shocks.

RELATED CONCEPTS AND QUICK RECALL

Related: conditional independence, correlation, copula. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Mutual exclusivity versus independence

BEGINNER

ONE-SENTENCE MEANING

Mutually exclusive events cannot occur together; independent events may occur together without changing one another's probabilities.

MEMORY JOGGER

Disjoint means impossible overlap; independent means unchanged odds.

WHY IT MATTERS IN QUANT RESEARCH

The distinction prevents invalid signal logic, incorrect scenario weights, and mistaken diversification claims.

FORMULA INTUITION

Disjoint: $P(A \text{ intersection } B)=0$; independent: $P(A \text{ intersection } B)=P(A)P(B)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Positive and negative return signs for one day are mutually exclusive; positive returns on two different days are not automatically independent.

PYTHON / SYSTEM IMPLEMENTATION

Test label partitions for zero overlap, while assessing independence with joint-versus-product diagnostics and temporal structure.

CONFUSIONS TO AVOID

Nonempty mutually exclusive events cannot be independent because observing one makes the other impossible.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Schema validation handles exclusivity; statistical validation handles independence. They are different controls owned by different pipeline stages.

RELATED CONCEPTS AND QUICK RECALL

Related: partition, dependence, event logic. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Pairwise and mutual independence

BEGINNER

ONE-SENTENCE MEANING

Pairwise independence checks every pair, while mutual independence requires every finite subset to factor jointly.

MEMORY JOGGER

All pairs can look independent while the group still obeys a hidden constraint.

WHY IT MATTERS IN QUANT RESEARCH

Higher-order dependence matters when combining many signals, counterparties, systems, or tail scenarios.

FORMULA INTUITION

$P(\bigcap_i A_i) = \prod_i P(A_i)$ for every relevant subset. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Three binary indicators can be constructed so each pair is independent while the third is determined by the first two.

PYTHON / SYSTEM IMPLEMENTATION

Use multivariate simulations and interaction diagnostics rather than relying solely on a correlation matrix or pairwise tests.

CONFUSIONS TO AVOID

Declaring a feature set independent because all pairwise correlations are near zero ignores nonlinear and higher-order structure.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The dependency model should version pairwise, group, and stress assumptions separately so aggregation behavior is auditable.

RELATED CONCEPTS AND QUICK RECALL

Related: higher-order dependence, interaction, factor model. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Partitions and the law of total probability

BEGINNER

ONE-SENTENCE MEANING

A partition breaks the sample space into disjoint exhaustive states, letting an event probability be assembled across those states.

MEMORY JOGGER

Weight each conditional rate by how often its state occurs.

WHY IT MATTERS IN QUANT RESEARCH

It combines regime-specific forecasts, scenario-conditioned losses, and segment-level model performance into an overall estimate.

FORMULA INTUITION

$P(A) = \sum_i P(A|B_i)P(B_i)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If calm days occur 80% with 45% loss probability and stressed days 20% with 70% loss probability, total loss probability is 50%.

PYTHON / SYSTEM IMPLEMENTATION

Validate that regime states are exhaustive and nonoverlapping, then store state weights, state-conditional estimates, and contribution terms.

CONFUSIONS TO AVOID

A good conditional model can still give a bad aggregate forecast if regime probabilities are misspecified.

WHERE THIS FITS IN THE RESEARCH SYSTEM

This identity links the regime classifier to the forecast aggregator and exposes whether errors come from state weights or within-state models.

RELATED CONCEPTS AND QUICK RECALL

Related: partition, mixture model, regime weighting. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Bayes' rule

BEGINNER

ONE-SENTENCE MEANING

Bayes' rule reverses a conditional probability by combining likelihood with prior probability and normalizing across possible causes.

MEMORY JOGGER

Posterior is prior times evidence, divided by total evidence.

WHY IT MATTERS IN QUANT RESEARCH

It supports signal interpretation, anomaly diagnosis, regime inference, model updating, and base-rate-aware alerting.

FORMULA INTUITION

$P(B|A) = P(A|B)P(B)/P(A)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A crash alert that catches 80% of crash days can still have low precision when crashes are rare and false alerts are common.

PYTHON / SYSTEM IMPLEMENTATION

Implement in log space for small likelihoods, include every competing hypothesis, and monitor posterior calibration and prior drift.

CONFUSIONS TO AVOID

Ignoring the base rate turns a sensitive detector into an apparently reliable predictor when most alerts are false.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Bayesian updates belong in the inference service; the alerting layer should consume calibrated posterior risk, not raw detector scores.

RELATED CONCEPTS AND QUICK RECALL

Related: prior, likelihood, posterior. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Odds and likelihood ratios

BEGINNER

ONE-SENTENCE MEANING

Odds express probability as success-to-failure weight, while a likelihood ratio shows how evidence multiplies prior odds.

MEMORY JOGGER

Evidence is additive in log-odds and multiplicative in odds.

WHY IT MATTERS IN QUANT RESEARCH

Odds provide a stable way to combine diagnostic evidence and interpret classifier updates across different base rates.

FORMULA INTUITION

$\text{odds}(A) = p/(1-p)$; $\text{posterior odds} = \text{prior odds} \times \text{likelihood ratio}$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A prior probability of 20% gives odds 1:4. Evidence with likelihood ratio 3 changes odds to 3:4, or probability 42.9%.

PYTHON / SYSTEM IMPLEMENTATION

Clip probabilities only for numerical safety, compute log-odds with stable functions, and keep likelihood-ratio calibration out of the final test period.

CONFUSIONS TO AVOID

A likelihood ratio is not the posterior probability; the prior remains essential.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model-combination layer can aggregate independent evidence in log-odds, while validation checks whether the resulting probabilities remain calibrated.

RELATED CONCEPTS AND QUICK RECALL

Related: logit, evidence, diagnostic test. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Random variables

BEGINNER

ONE-SENTENCE MEANING

A random variable maps each outcome in the sample space to a numerical value that can be analyzed.

MEMORY JOGGER

It is a function of outcomes, not a variable that changes without structure.

WHY IT MATTERS IN QUANT RESEARCH

Returns, P&L, spread changes, order fills, drawdowns, and realized volatility are modeled as random variables.

FORMULA INTUITION

$X: \Omega \rightarrow \mathbb{R}$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

For each next-day market outcome, X might map the realized close-to-close return to a number such as -0.012 or 0.006.

PYTHON / SYSTEM IMPLEMENTATION

Define value, unit, timestamp, horizon, corporate-action treatment, missing-state policy, and valid domain in the data contract.

CONFUSIONS TO AVOID

Treating prices and returns as interchangeable random variables leads to inconsistent stationarity, scale, and aggregation assumptions.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Random-variable definitions form the semantic layer shared by feature engineering, simulation, risk, and performance attribution.

RELATED CONCEPTS AND QUICK RECALL

Related: outcome map, return, measurable quantity. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Discrete versus continuous variables

BEGINNER

ONE-SENTENCE MEANING

Discrete variables take countable values; continuous models assign density across intervals of possible values.

MEMORY JOGGER

For continuous X , a point has probability zero even though intervals have positive probability.

WHY IT MATTERS IN QUANT RESEARCH

The distinction determines whether to use probability masses, densities, sums, integrals, and appropriate estimators.

FORMULA INTUITION

Discrete: $\sum_x p(x)=1$; continuous: $\int f(x)dx=1$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Trade count is discrete; a modeled return is continuous, although stored prices and tick sizes make observed values granular.

PYTHON / SYSTEM IMPLEMENTATION

Preserve integer types for counts and document discretization or binning when continuous measurements are converted into categories.

CONFUSIONS TO AVOID

Calling a density value a probability can produce values above one and incorrect interval calculations.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The schema registry should label support type because likelihood functions, sampling code, and visual diagnostics depend on it.

RELATED CONCEPTS AND QUICK RECALL

Related: PMF, PDF, support. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability mass functions

BEGINNER

ONE-SENTENCE MEANING

A probability mass function assigns probability to each possible value of a discrete random variable.

MEMORY JOGGER

Mass lives on points and all point masses sum to one.

WHY IT MATTERS IN QUANT RESEARCH

PMFs model counts, states, fills, defaults, signal votes, and finite scenario outcomes.

FORMULA INTUITION

$p_X(x) = P(X=x)$; $\sum_x p_X(x) = 1$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A fill-count model might assign probabilities to 0, 1, 2, or 3 partial fills during an interval.

PYTHON / SYSTEM IMPLEMENTATION

Use explicit support arrays aligned by value, validate normalization, and retain zero-probability states when their absence is semantically meaningful.

CONFUSIONS TO AVOID

Dropping unobserved but possible states makes empirical PMFs overconfident and can break downstream joins.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The discrete distribution object supplies scenario sampling and exact likelihoods to execution and operational-risk models.

RELATED CONCEPTS AND QUICK RECALL

Related: support, categorical distribution, histogram. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability density functions

BEGINNER

ONE-SENTENCE MEANING

A probability density function distributes probability per unit of a continuous variable; interval probability is area under the curve.

MEMORY JOGGER

Height is density; area is probability.

WHY IT MATTERS IN QUANT RESEARCH

Densities support return models, likelihoods, tail calculations, simulation, and comparisons across continuous outcomes.

FORMULA INTUITION

$P(a \leq X \leq b) = \int_a^b f_X(x) dx$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A daily-return density can be highest near zero, while the chance of a return between -1% and 1% is the area across that interval.

PYTHON / SYSTEM IMPLEMENTATION

Use numerically stable distribution libraries, check units after scaling, and verify integration or sampling with synthetic tests.

CONFUSIONS TO AVOID

Comparing densities after changing units without the Jacobian can create apparently different distributions for the same variable.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Density models sit in the probabilistic model registry with parameter estimates, fit window, transformations, and goodness-of-fit evidence.

RELATED CONCEPTS AND QUICK RECALL

Related: CDF, likelihood, Jacobian. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Cumulative distribution functions

BEGINNER

ONE-SENTENCE MEANING

A cumulative distribution function gives the probability that a random variable is less than or equal to a threshold.

MEMORY JOGGER

The CDF accumulates probability from the far left up to x .

WHY IT MATTERS IN QUANT RESEARCH

CDFs answer quantile, breach, percentile, and calibration questions without requiring the reader to interpret density height.

FORMULA INTUITION

$F_X(x) = P(X \leq x)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If $F(-0.02) = 0.04$, the model assigns a 4% chance to a return of -2% or worse.

PYTHON / SYSTEM IMPLEMENTATION

Use empirical CDFs as model-free baselines, compare modeled and observed CDFs, and implement right-continuous threshold conventions consistently.

CONFUSIONS TO AVOID

For losses versus returns, left and right tails reverse; ambiguous sign conventions corrupt risk reporting.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The risk API should expose CDF, survival, and quantile functions under one sign convention to prevent repeated ad hoc tail logic.

RELATED CONCEPTS AND QUICK RECALL

Related: quantile, survival function, percentile. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Survival functions and tail probabilities

BEGINNER

ONE-SENTENCE MEANING

A survival function gives the probability that a variable exceeds a threshold and is often more accurate in the far tail than one minus a rounded CDF.

MEMORY JOGGER

Ask how much probability survives beyond x .

WHY IT MATTERS IN QUANT RESEARCH

It directly represents exceedance risk for losses, latency, spread, slippage, and time-to-event variables.

FORMULA INTUITION

$S_X(x) = P(X > x) = 1 - F_X(x)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If the slippage survival at 12 basis points is 0.01, one modeled execution in 100 exceeds 12 basis points.

PYTHON / SYSTEM IMPLEMENTATION

Call a library survival function instead of subtracting from one for extreme thresholds; validate tail estimates on log scales.

CONFUSIONS TO AVOID

Tail estimates based on a handful of observations carry large uncertainty even when the plotted curve appears smooth.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The limit-monitoring service maps survival probabilities to alert thresholds and must store sample support and extrapolation flags.

RELATED CONCEPTS AND QUICK RECALL

Related: hazard, exceedance, extreme tail. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Quantiles and percentiles

BEGINNER

ONE-SENTENCE MEANING

A quantile is a threshold below which a stated fraction of probability lies, subject to conventions for discrete data and interpolation.

MEMORY JOGGER

Probability goes in; a value threshold comes out.

WHY IT MATTERS IN QUANT RESEARCH

Quantiles define VaR-like limits, robust summaries, execution benchmarks, clipping rules, and forecast intervals.

FORMULA INTUITION

$q_{\alpha} = \inf\{x: F_X(x) \geq \alpha\}$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

The 5th percentile of return is a downside threshold exceeded on the lower side about 5% of the modeled time.

PYTHON / SYSTEM IMPLEMENTATION

Pin the quantile definition and interpolation method, use training-only estimates for preprocessing, and report uncertainty for small tails.

CONFUSIONS TO AVOID

The 95th percentile of loss and the 5th percentile of return may describe similar risk but have opposite signs and labels.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Quantile artifacts are reused by preprocessing, risk limits, dashboards, and backtests, so convention metadata must travel with every value.

RELATED CONCEPTS AND QUICK RECALL

Related: VaR, median, inverse CDF. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Expectation

BEGINNER

ONE-SENTENCE MEANING

Expectation is the probability-weighted long-run average of a random variable under the model.

MEMORY JOGGER

Multiply outcomes by their weights, then add or integrate.

WHY IT MATTERS IN QUANT RESEARCH

Expected return, expected cost, expected loss, and expected utility organize decisions, while not describing the full range of outcomes.

FORMULA INTUITION

$E[X] = \sum x \times p(x)$ or $\int x f(x) dx$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A trade paying +2 with probability 0.6 and -1 with probability 0.4 has expected payoff 0.8 units.

PYTHON / SYSTEM IMPLEMENTATION

Estimate on a declared sample and horizon, keep units explicit, and compare the mean with robust alternatives and uncertainty intervals.

CONFUSIONS TO AVOID

Positive expectation does not imply a high probability of profit or acceptable drawdown; rare large gains can dominate the mean.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Expected values feed ranking and optimization only after the risk layer receives the accompanying distribution and estimation uncertainty.

RELATED CONCEPTS AND QUICK RECALL

Related: mean, expected utility, sample average. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Linearity of expectation

BEGINNER

ONE-SENTENCE MEANING

The expectation of a sum equals the sum of expectations, whether or not the terms are independent when the expectations exist.

MEMORY JOGGER

Means add without an independence assumption.

WHY IT MATTERS IN QUANT RESEARCH

It enables portfolio expected-return aggregation, cost decomposition, and attribution across trades or scenarios.

FORMULA INTUITION

$E[aX+bY]=aE[X]+bE[Y]$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A portfolio's expected return is the weighted sum of asset expected returns even when the assets are correlated.

PYTHON / SYSTEM IMPLEMENTATION

Compute contribution vectors before aggregation and reconcile their sum to the portfolio expectation at machine precision.

CONFUSIONS TO AVOID

Independence becomes important for variance and distribution shape, not for the linearity identity itself.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The attribution engine uses this identity to decompose expected alpha, fees, financing, borrow, and impact into auditable components.

RELATED CONCEPTS AND QUICK RECALL

Related: weighted sum, portfolio return, attribution. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Variance and standard deviation

BEGINNER

ONE-SENTENCE MEANING

Variance is expected squared distance from the mean; standard deviation returns that dispersion to the variable's original units.

MEMORY JOGGER

Variance squares surprises, so large deviations matter disproportionately.

WHY IT MATTERS IN QUANT RESEARCH

They summarize volatility, forecast uncertainty, execution dispersion, and the risk of aggregated positions.

FORMULA INTUITION

$\text{Var}(X) = E[(X - E[X])^2] = E[X^2] - E[X]^2$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Returns of -2%, 0%, and 2% around mean zero have average squared deviations that determine their volatility.

PYTHON / SYSTEM IMPLEMENTATION

Use numerically stable estimators, define population versus sample denominator, and monitor both variance and tail diagnostics.

CONFUSIONS TO AVOID

Low variance does not guarantee low downside risk when distributions are skewed, heavy-tailed, or regime-switching.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Variance estimates flow into risk models and optimizers; their window, frequency, annualization, and missing-data policy must be versioned.

RELATED CONCEPTS AND QUICK RECALL

Related: volatility, second moment, dispersion. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Covariance

BEGINNER

ONE-SENTENCE MEANING

Covariance measures whether two variables tend to deviate from their means in the same or opposite direction.

MEMORY JOGGER

Positive moves together; negative moves oppositely; magnitude depends on units.

WHY IT MATTERS IN QUANT RESEARCH

It determines portfolio variance, hedge behavior, factor exposure aggregation, and multivariate simulation.

FORMULA INTUITION

$\text{Cov}(X, Y) = E[(X - E[X])(Y - E[Y])]$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A stock and market index with positive covariance tend to have same-signed deviations, increasing combined risk when both positions are long.

PYTHON / SYSTEM IMPLEMENTATION

Align timestamps and calendars, use point-in-time returns, inspect missing pairs, and estimate a positive-semidefinite covariance matrix.

CONFUSIONS TO AVOID

Near-zero covariance does not establish independence and can conceal nonlinear or tail dependence.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The centralized risk service should own covariance estimation so research, backtests, and production sizing use the same snapshot.

RELATED CONCEPTS AND QUICK RECALL

Related: correlation, portfolio variance, dependence. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Correlation

BEGINNER

ONE-SENTENCE MEANING

Correlation standardizes covariance by both variables' standard deviations, producing a unitless linear-dependence measure between -1 and 1.

MEMORY JOGGER

Correlation is covariance after scale is removed.

WHY IT MATTERS IN QUANT RESEARCH

It compares co-movement across assets and features, supports diversification analysis, and diagnoses redundant signals.

FORMULA INTUITION

$\rho_{XY} = \text{Cov}(X, Y) / (\sigma_X \sigma_Y)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Two assets can have identical correlation but very different volatilities, so equal weights can still contribute very different risk.

PYTHON / SYSTEM IMPLEMENTATION

Estimate on aligned returns, report window and confidence, stress regimes, and avoid computing when either variance is effectively zero.

CONFUSIONS TO AVOID

Correlation is not causation, independence, tail dependence, or a stable constant across market regimes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Correlation belongs in exploratory and monitoring views; portfolio calculations should retain the full covariance scale.

RELATED CONCEPTS AND QUICK RECALL

Related: covariance, scatterplot, tail dependence. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Transformations of random variables

BEGINNER

ONE-SENTENCE MEANING

A transformation creates a new random variable by applying a deterministic function to an existing one.

MEMORY JOGGER

Changing the value scale changes the distribution, and sometimes its support or density.

WHY IT MATTERS IN QUANT RESEARCH

Returns from prices, losses from P&L, volatility from squared returns, and signals from ranks are all transformations.

FORMULA INTUITION

$Y=g(X)$; for monotone g , $F_Y(y)=F_X(g^{-1}(y))$ with direction care. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If $L=-R$, a 5% lower-tail return event becomes a 95% upper-tail loss event under a sign reversal.

PYTHON / SYSTEM IMPLEMENTATION

Treat every transformation as a versioned pipeline step with domain checks, units, invertibility notes, and train-only fitted parameters.

CONFUSIONS TO AVOID

Applying log to nonpositive values or mixing simple and log returns creates undefined values and inconsistent aggregation.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The transformation graph is stored with feature lineage so simulations and inverse mappings reproduce exactly what models saw.

RELATED CONCEPTS AND QUICK RECALL

Related: change of variables, log return, standardization. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Indicator variables

BEGINNER

ONE-SENTENCE MEANING

An indicator variable equals one when an event occurs and zero otherwise, converting event logic into arithmetic.

MEMORY JOGGER

The mean of a 0/1 indicator is the event frequency.

WHY IT MATTERS IN QUANT RESEARCH

Indicators make hit rates, counts, conditional summaries, labels, and event-study regressions easy to express.

FORMULA INTUITION

$I_A \in \{0,1\}$; $E[I_A] = P(A)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A daily breach indicator averages 0.03 over a stable sample when breaches occur on 3% of observed days.

PYTHON / SYSTEM IMPLEMENTATION

Store boolean and integer representations consistently, retain missing as unknown rather than zero, and verify event timestamps.

CONFUSIONS TO AVOID

Filling unknown outcomes with zero biases event probabilities downward and falsely improves apparent safety.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Indicators connect event definitions to estimators, dashboards, alert-rate monitoring, and supervised-learning labels.

RELATED CONCEPTS AND QUICK RECALL

Related: Bernoulli variable, event rate, binary label. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Mixture distributions

BEGINNER

ONE-SENTENCE MEANING

A mixture distribution combines several component distributions using probabilities that sum to one.

MEMORY JOGGER

Choose a state, then draw from that state's distribution.

WHY IT MATTERS IN QUANT RESEARCH

Mixtures represent calm and stressed markets, multiple liquidity states, heterogeneous assets, and latent regimes.

FORMULA INTUITION

$f(x) = \sum_k \pi_k f_k(x)$, $\pi_k \geq 0$, $\sum \pi_k = 1$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Daily returns may be modeled as an 85% calm component with low volatility and a 15% stressed component with high volatility and negative mean.

PYTHON / SYSTEM IMPLEMENTATION

Estimate with identifiability controls, multiple initializations, and time-aware validation; store component labels and weight ordering rules.

CONFUSIONS TO AVOID

Components are modeling devices, not necessarily real economic regimes, and labels can switch across fits.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The scenario generator draws a regime first and an outcome second, while monitoring separately checks weights and within-regime fit.

RELATED CONCEPTS AND QUICK RECALL

Related: law of total probability, latent state, regime model. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Convolution and sums

BEGINNER

ONE-SENTENCE MEANING

The distribution of a sum combines the component distributions through convolution, with dependence assumptions determining whether simple formulas apply.

MEMORY JOGGER

Many small risks aggregate by adding outcomes, not by averaging their probabilities.

WHY IT MATTERS IN QUANT RESEARCH

Convolution models cumulative P&L, total demand, aggregate latency, and multi-period returns under specified dependence.

FORMULA INTUITION

$f_{X+Y}(z) = \int f_X(x)f_Y(z-x)dx$ for independent X, Y . Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Independent trade costs with discrete PMFs can be convolved to obtain the distribution of total cost across two orders.

PYTHON / SYSTEM IMPLEMENTATION

Use FFT methods for large discrete supports, direct simulation for complex dependence, and verify mean and support against analytic identities.

CONFUSIONS TO AVOID

Convolving marginals as if independent can severely understate joint tail risk during crowded or stressed conditions.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The aggregation engine selects analytic, FFT, or simulation methods based on support type and dependency metadata.

RELATED CONCEPTS AND QUICK RECALL

Related: sum distribution, independence, FFT. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability generating and moment summaries

BEGINNER

ONE-SENTENCE MEANING

Generating functions and moments summarize distributions through weighted powers or transforms when those quantities exist.

MEMORY JOGGER

A compact transform can encode many probabilities or moments at once.

WHY IT MATTERS IN QUANT RESEARCH

They help derive count aggregates, expected totals, variance identities, and approximations for operational events.

FORMULA INTUITION

$M_X(t) = E[e^{tX}]$; $G_X(s) = E[s^X]$ for nonnegative integer X . Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

For independent counts, probability-generating functions multiply, making the total-count distribution easier to derive.

PYTHON / SYSTEM IMPLEMENTATION

Use symbolic derivations only with domain checks; in numerical work prefer stable log transforms and verify moments by simulation.

CONFUSIONS TO AVOID

Heavy-tailed distributions may not have finite high moments or moment-generating functions near zero.

WHERE THIS FITS IN THE RESEARCH SYSTEM

These tools belong in model validation and analytic test fixtures, not as opaque production features without domain guarantees.

RELATED CONCEPTS AND QUICK RECALL

Related: moment, transform, cumulant. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Law of large numbers

BEGINNER

ONE-SENTENCE MEANING

The law of large numbers states that under suitable conditions, sample averages converge toward the expected value as observations accumulate.

MEMORY JOGGER

More representative observations stabilize averages, but dependence and drift can slow or break the intuition.

WHY IT MATTERS IN QUANT RESEARCH

It underpins empirical frequencies, mean estimates, simulation averages, and performance measurement.

FORMULA INTUITION

$(1/n) \sum_i X_i \rightarrow E[X]$ in probability under stated conditions. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A simulated strategy's average payoff becomes less erratic as independent scenario count grows, although tail metrics may converge much more slowly.

PYTHON / SYSTEM IMPLEMENTATION

Track estimates against effective sample size, run convergence plots, and use block methods when observations are serially dependent.

CONFUSIONS TO AVOID

More data from the same biased, nonstationary, or leaked process does not guarantee convergence to the desired quantity.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The simulation service uses convergence diagnostics and stopping rules rather than a fixed scenario count chosen without error targets.

RELATED CONCEPTS AND QUICK RECALL

Related: sample mean, consistency, effective sample size. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Central limit theorem

BEGINNER

ONE-SENTENCE MEANING

The central limit theorem explains why suitably normalized sums can approach a normal distribution under specific regularity conditions.

MEMORY JOGGER

It is about sums and scaling, not a claim that raw market returns are normal.

WHY IT MATTERS IN QUANT RESEARCH

It motivates approximate standard errors, confidence calculations, and Monte Carlo error estimates.

FORMULA INTUITION

$\sqrt{n}(\bar{X} - \mu)/\sigma \Rightarrow N(0,1)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

The average of many independent finite-variance simulated costs may be approximately normal even when individual costs are skewed.

PYTHON / SYSTEM IMPLEMENTATION

Check dependence, tail behavior, sample size, and convergence with simulation; use robust or bootstrap alternatives when conditions are doubtful.

CONFUSIONS TO AVOID

Heavy tails, strong dependence, changing regimes, or selection can make normal approximations poor where risk decisions matter most.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The analytics layer may use CLT-based error bars for stable averages while the tail-risk layer retains empirical and stress methods.

RELATED CONCEPTS AND QUICK RECALL

Related: normal approximation, standard error, asymptotics. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Monte Carlo probability estimation

BEGINNER

ONE-SENTENCE MEANING

Monte Carlo estimates probabilities and expectations by drawing scenarios from a defined model and averaging their outcomes.

MEMORY JOGGER

Simulate, evaluate, average, and quantify simulation error.

WHY IT MATTERS IN QUANT RESEARCH

It handles path-dependent payoffs, nonlinear rules, complex portfolios, and uncertainty that resists closed-form calculation.

FORMULA INTUITION

$P(A)$ approximately $(1/N)\sum_i I_A(\omega_i)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

To estimate barrier-breach probability, simulate price paths, record whether each path crosses the barrier, and average the indicators.

PYTHON / SYSTEM IMPLEMENTATION

Use seeded reproducible streams, vectorized batches, convergence diagnostics, independent validation, and recorded generator/model versions.

CONFUSIONS TO AVOID

A million simulations from a misspecified model estimate the wrong probability very precisely.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The scenario platform separates model error from Monte Carlo error and registers seeds, draws, transformations, and payoff code.

RELATED CONCEPTS AND QUICK RECALL

Related: simulation error, random seed, variance reduction. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Sampling error and effective sample size

BEGINNER

ONE-SENTENCE MEANING

Sampling error is the difference between an estimate and its target caused by observing a finite, dependent sample.

MEMORY JOGGER

The row count is not always the information count.

WHY IT MATTERS IN QUANT RESEARCH

It calibrates confidence in hit rates, tail probabilities, correlations, and backtest summaries.

FORMULA INTUITION

$SE(\hat{p})$ approximately $\sqrt{p(1-p)/n}$ for independent Bernoulli trials. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A 60% hit rate from 20 trades is much less precise than 60% from 2,000 independent trades, and clustered trades reduce effective support further.

PYTHON / SYSTEM IMPLEMENTATION

Report raw count, event count, time span, cluster structure, and effective sample size; use block bootstrap when serial dependence is material.

CONFUSIONS TO AVOID

Presenting estimates without support makes unstable small-sample results look comparable to well-supported ones.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Every metric artifact should carry numerator, denominator, effective support, and uncertainty fields into dashboards and model reviews.

RELATED CONCEPTS AND QUICK RECALL

Related: standard error, block bootstrap, confidence interval. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Calibration of probabilistic forecasts

BEGINNER

ONE-SENTENCE MEANING

A probabilistic forecast is calibrated when events assigned probability p occur about p of the time among comparable forecasts.

MEMORY JOGGER

Among 70% forecasts, roughly 70% should realize over a stable evaluation set.

WHY IT MATTERS IN QUANT RESEARCH

Calibration turns model scores into decision-ready probabilities for sizing, expected value, and risk limits.

FORMULA INTUITION

$E[Y \mid \hat{p}=p]$ approximately p . Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If only 50% of events occur in the bucket forecast at 70%, the model is overconfident there.

PYTHON / SYSTEM IMPLEMENTATION

Use time-split reliability curves, Brier score decomposition, sample counts, and calibration models fitted only on training or validation data.

CONFUSIONS TO AVOID

Good ranking does not imply good calibration, and aggregate calibration can hide regime-specific failure.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The probability calibration service sits after predictive scoring and before expected-value, sizing, or alert-threshold consumers.

RELATED CONCEPTS AND QUICK RECALL

Related: reliability curve, Brier score, discrimination. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Proper scoring rules

BEGINNER

ONE-SENTENCE MEANING

A proper scoring rule rewards honest probability forecasts by making the true belief distribution optimal in expectation.

MEMORY JOGGER

Score the full probability, not just whether the class guess was right.

WHY IT MATTERS IN QUANT RESEARCH

Proper scores compare forecast quality, penalize overconfidence, and support calibration-aware model selection.

FORMULA INTUITION

Brier = $(p - y)^2$; log loss = $-(y \log p + (1 - y) \log(1 - p))$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A 99% bullish forecast is penalized much more by log loss when the market falls than a cautious 60% forecast.

PYTHON / SYSTEM IMPLEMENTATION

Clip only at documented numerical bounds, compute scores out of sample, retain per-observation contributions, and segment by regime.

CONFUSIONS TO AVOID

Accuracy discards confidence information; choosing thresholds before evaluating probabilities can conceal poor calibration.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The experiment registry should rank probabilistic models using proper scores plus economic utility, not hit rate alone.

RELATED CONCEPTS AND QUICK RECALL

Related: log loss, Brier score, entropy. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Base rates and class imbalance

BEGINNER

ONE-SENTENCE MEANING

A base rate is the unconditional prevalence of an event in the relevant population and time period.

MEMORY JOGGER

Rare events remain rare even when a detector is sensitive.

WHY IT MATTERS IN QUANT RESEARCH

Base rates determine alert precision, posterior probability, benchmark accuracy, and the value of a trading classifier.

FORMULA INTUITION

$P(A) = \text{event count} / \text{eligible count}$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

If severe losses occur on 1% of days, a detector with a 5% false-positive rate can produce many more false than true alerts.

PYTHON / SYSTEM IMPLEMENTATION

Freeze the eligible population, monitor prevalence drift, use precision-recall diagnostics, and compare against a constant-base-rate forecast.

CONFUSIONS TO AVOID

Balancing classes during training changes the effective prior; raw model outputs may need correction before they are probabilities.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The monitoring layer tracks production prevalence and triggers recalibration or threshold review when it departs from training assumptions.

RELATED CONCEPTS AND QUICK RECALL

Related: prior probability, precision, prevalence drift. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Rare events and tail estimation

BEGINNER

ONE-SENTENCE MEANING

Rare-event probability estimation focuses on outcomes with little direct support, where model assumptions and uncertainty dominate.

MEMORY JOGGER

The farther into the tail, the more the answer depends on extrapolation.

WHY IT MATTERS IN QUANT RESEARCH

It is central to crash risk, extreme slippage, liquidity failure, operational outages, and stress capital.

FORMULA INTUITION

For k observed events in n trials, $\hat{p}=k/n$, but uncertainty is large when k is small. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

Seeing zero exchange outages in 100 days does not prove zero outage probability; it only provides limited evidence at that horizon.

PYTHON / SYSTEM IMPLEMENTATION

Combine empirical counts, confidence bounds, stress scenarios, domain constraints, and importance sampling where appropriate; mark extrapolated regions.

CONFUSIONS TO AVOID

Reporting a precise tail percentile from only a few effective observations creates false confidence.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The tail-risk service should expose estimate, interval, support, model family, threshold, and stress override as separate fields.

RELATED CONCEPTS AND QUICK RECALL

Related: extreme value, importance sampling, confidence bound. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability trees and path dependence

BEGINNER

ONE-SENTENCE MEANING

A probability tree represents sequential events as branches whose path probabilities multiply and whose terminal alternatives sum.

MEMORY JOGGER

Multiply down a path; add across disjoint paths.

WHY IT MATTERS IN QUANT RESEARCH

Trees clarify multistage trades, regime transitions, execution outcomes, and contingent risk controls.

FORMULA INTUITION

$P(\text{path}) = \prod_t P(\text{branch}_t \mid \text{history}_t)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

An order may first be filled or not filled, then experience favorable or adverse price movement; payoff depends on the entire branch history.

PYTHON / SYSTEM IMPLEMENTATION

Store nodes with state, timestamp, conditional probability, action, payoff, and parent identifier; validate sibling probabilities at every node.

CONFUSIONS TO AVOID

Merging branches that reach the same current state can be invalid when future behavior depends on how that state was reached.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The decision engine can evaluate policies on the tree while the audit log retains the branch assumptions used at each decision point.

RELATED CONCEPTS AND QUICK RECALL

Related: chain rule, transition, scenario tree. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability versus frequency

BEGINNER

ONE-SENTENCE MEANING

Probability is a model-based uncertainty statement, while observed frequency is a sample statistic used to estimate or test that statement.

MEMORY JOGGER

Forecast probability comes before the outcome; frequency summarizes outcomes afterward.

WHY IT MATTERS IN QUANT RESEARCH

Keeping them distinct prevents in-sample frequencies from being presented as certain future odds.

FORMULA INTUITION

p is a model parameter or forecast; $\hat{p}=k/n$ is an estimate. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A historical 52% up-day frequency is evidence about a probability, not a guarantee that the next day has exactly 52% upside chance.

PYTHON / SYSTEM IMPLEMENTATION

Store forecasts before realization, maintain immutable outcome joins, and compare forecast buckets with later frequencies using uncertainty bands.

CONFUSIONS TO AVOID

Using realized labels to revise same-timestamp forecasts is leakage disguised as calibration.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The forecast ledger separates prediction-time artifacts from evaluation-time outcomes and enables honest longitudinal calibration.

RELATED CONCEPTS AND QUICK RECALL

Related: estimation, backtest, forecast ledger. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Model probability versus real-world uncertainty

BEGINNER

ONE-SENTENCE MEANING

A probability distribution describes uncertainty inside a chosen model; model uncertainty concerns whether that model and its parameters are adequate.

MEMORY JOGGER

A precise answer can be conditional on a fragile model.

WHY IT MATTERS IN QUANT RESEARCH

Quant decisions need sensitivity to alternative distributions, parameters, regimes, and data-generating assumptions.

FORMULA INTUITION

Total uncertainty = outcome variation + parameter uncertainty + model uncertainty (conceptually). Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A normal return model may report a tiny crash probability, while a heavy-tailed or regime-mixture model reports materially more risk.

PYTHON / SYSTEM IMPLEMENTATION

Run challenger models, parameter posterior or bootstrap draws, regime stresses, and decision-regret comparisons instead of trusting one fitted distribution.

CONFUSIONS TO AVOID

Narrow confidence intervals conditional on one model do not measure uncertainty about model misspecification.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The governance layer registers champion and challenger distributions, assumptions, validation failures, and fallback stress overlays.

RELATED CONCEPTS AND QUICK RECALL

Related: epistemic uncertainty, robustness, model risk. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability-to-decision translation

BEGINNER

ONE-SENTENCE MEANING

A probability forecast becomes useful only after it is combined with payoffs, costs, constraints, and the consequences of being wrong.

MEMORY JOGGER

Probability alone does not choose the trade; expected consequences do.

WHY IT MATTERS IN QUANT RESEARCH

The same event probability can justify buying, selling, hedging, or doing nothing under different payoff asymmetries and costs.

FORMULA INTUITION

Expected utility(action)= $\sum_s P(s|information) U(action,s)$. Check the event definition, conditioning information, support, and units before calculating.

SIMPLE MARKET EXAMPLE

A 55% win probability is unattractive if wins earn 1 unit, losses cost 2 units, and trading costs consume 0.1 unit.

PYTHON / SYSTEM IMPLEMENTATION

Create a decision table that joins calibrated scenario probabilities to payoff distributions, costs, exposure limits, and abstention alternatives.

CONFUSIONS TO AVOID

Optimizing classification accuracy without modeling payoff magnitude and selection can reduce economic performance.

WHERE THIS FITS IN THE RESEARCH SYSTEM

This is the handoff from the probability layer to portfolio construction, execution, and risk approval, with assumptions preserved end to end.

RELATED CONCEPTS AND QUICK RECALL


Related: expected value, utility, decision threshold. Define the concept in plain English, reproduce its formula, and name one data or modeling failure that would invalidate its use.

Probability tree: volatility and return sign

BEGINNER

HOW TO READ THE TREE

Branch probabilities are conditional on the node already reached. Multiply down a path to obtain a joint probability; add terminal paths only when they are mutually exclusive.



WORKED PATH PROBABILITIES

Calm and gain: $0.80 \times 0.55 = 0.44$. Calm and loss: $0.80 \times 0.45 = 0.36$. Stress and gain: $0.20 \times 0.30 = 0.06$. Stress and loss: $0.20 \times 0.70 = 0.14$. Terminal mass totals 1.00.

AGGREGATE RESULT

Overall gain probability is $0.44 + 0.06 = 0.50$. Overall loss probability is $0.36 + 0.14 = 0.50$. The equal aggregate rates hide a much worse conditional outcome in stress.

SYSTEM LESSON


Store local branch estimates and path contributions. When the overall forecast changes, attribution can separate changing regime weights from changing within-regime probabilities.

PMF, density, CDF, and survival views

BEGINNER

ONE DISTRIBUTION, DIFFERENT QUESTIONS

A PMF or density describes local mass; a CDF answers probability below a threshold; a survival function answers probability above it. Use the view that matches the decision.



TAIL INTERPRETATION

The shaded area is $P(X > \text{threshold})$, not the curve height at the threshold. The CDF gives the complementary left area. In extreme tails, call the survival function directly for numerical accuracy.

UNIT WARNING

If returns change from decimals to percent, the horizontal scale multiplies by 100 and density height divides by 100 so total area remains one. Probability itself is invariant to the display unit.

EMPIRICAL BASELINE

Before fitting a parametric curve, compare it with an empirical CDF and observed exceedance counts. Inspect the center and each tail separately; a smooth central fit can conceal poor loss-tail behavior.

RECALL


Explain why $f(x)$ can exceed one without violating probability axioms. Then state which function directly answers a limit-breach question and which function returns a percentile threshold.

Monte Carlo convergence and sampling error

BEGINNER

WHAT CONVERGENCE SHOULD LOOK LIKE

A running estimate should stabilize around its target while its typical Monte Carlo error shrinks roughly with $1/\sqrt{N}$. Rare events require far more scenarios than central averages.



ERROR ARITHMETIC

For a Bernoulli event with $p=0.20$ and $N=1,000$ independent draws, standard error is $\sqrt{0.20 \times 0.80 / 1,000} = 1.26$ percentage points. A rough 95% Monte Carlo band is about ± 2.48 points.

DEPENDENCE PENALTY

If paths share random shocks, are duplicated, or observations are serially correlated, N overstates effective support. Estimate effective sample size or redesign the simulation streams before trusting the narrow band.

MODEL ERROR IS SEPARATE

The blue line can converge perfectly to the probability implied by a wrong model. Stress distributions, parameters, dependence, and regimes in addition to reducing Monte Carlo noise.

STOPPING RULE

Choose scenario count from a decision-relevant error tolerance and confidence target. Record seed, generator, model version, batch estimates, and convergence evidence for reproducibility.

Formula sheet and worked calculations

BEGINNER

EVENT IDENTITIES

Complement: $P(A^c) = 1 - P(A)$. Addition: $P(A \cup B) = P(A) + P(B) - P(A \cap B)$. Conditional: $P(A|B) = P(A \cap B) / P(B)$. Multiplication: $P(A \cap B) = P(A|B)P(B)$.

BAYES WORKED EXAMPLE

Suppose crash prevalence is 1%, detector sensitivity is 80%, and false-positive rate is 5%.
 $P(\text{alert}) = 0.80(0.01) + 0.05(0.99) = 0.0575$. Therefore $P(\text{crash}|\text{alert}) = 0.008 / 0.0575 = 13.9\%$, not 80%.

TOTAL PROBABILITY WORKED EXAMPLE

Calm probability 0.75 with 40% loss chance; stress probability 0.25 with 68% loss chance. Overall loss probability is $0.75(0.40) + 0.25(0.68) = 0.47$. Stress contributes 0.17 of total probability mass.

EXPECTATION AND VARIANCE

For X taking -1 with probability 0.4 and +2 with probability 0.6: $E[X] = 0.8$. $E[X^2] = 2.8$. $\text{Var}(X) = 2.8 - 0.64 = 2.16$, so standard deviation is 1.470.

PORTFOLIO AGGREGATION

Expected portfolio return is $w^T \mu$. Variance is $w^T \Sigma w$, which includes covariance cross-terms. Linearity of expectation requires no independence; variance aggregation does.

RECALL

Derive the addition rule from disjoint pieces. Explain every denominator in Bayes' example, then distinguish outcome variability, parameter uncertainty, and model uncertainty.

Failure modes and diagnostic table

BEGINNER

PROBABILITIES DO NOT SUM TO ONE

Check missing states, duplicate labels, negative transformed scores, row alignment, and rounding. Do not blindly renormalize until the source defect is understood.

HIGH ACCURACY, POOR DECISIONS

Inspect class imbalance, base-rate benchmark, probability calibration, payoff asymmetry, costs, and selection. Accuracy alone discards confidence and economics.

EXCELLENT BACKTEST CALIBRATION, WEAK LIVE CALIBRATION

Audit leakage, population drift, regime mix, label delay, recalibration timing, and whether production scores use the registered transform.

TAIL PROBABILITY CHANGES DRAMATICALLY

Report effective tail count, threshold, model family, parameter uncertainty, dependence, and extrapolation. Use stress bounds instead of a false point estimate.

CONDITIONAL RATE APPEARS IMPRESSIVE

Check denominator size, multiple comparisons, cohort selection, timestamp eligibility, and the reverse conditional. Compare with unconditional base rate.

SIMULATION ESTIMATE IS STABLE BUT WRONG

Convergence controls Monte Carlo error only. Validate model assumptions, random-number transformations, dependence, payoff code, and analytic toy cases.

Implementation and model-registry checklist

BEGINNER

EVENT CONTRACT

Event name, plain definition, universe, horizon, observation clock, cutoff, label delay, tie handling, unknown state, and owner.

EVIDENCE SNAPSHOT

Source tables, as-of timestamp, eligible count, event count, exclusions, missingness, dependence structure, effective sample size, and immutable hash.

PROBABILITY MODEL

Support, distribution or algorithm, prior, parameters, conditioning features, transformations, fit window, seed, code version, and assumptions.

VALIDATION EVIDENCE

Time splits, calibration curve, proper scores, uncertainty, subgroup and regime checks, tail support, challenger models, and economic utility.

RUNTIME CHECKS

Finite probabilities, range, normalization, schema, feature freshness, support drift, prevalence drift, calibration lag, latency, and fallback.

DECISION HANDOFF

Scenario probabilities, payoffs, costs, constraints, threshold or utility rule, abstention behavior, approval, monitoring, and rollback artifact.

Final active recall index

BEGINNER

FOUNDATIONS

Define experiment, outcome, sample space, event, probability axioms, complement, union, intersection, conditional probability, and independence using one market example.

DISTRIBUTION VIEWS

Contrast PMF, PDF, CDF, survival function, and quantile. State what each consumes, what it returns, and how units affect it.

UPDATING

Build a probability tree, apply total probability, reverse a conditional with Bayes' rule, and explain the role of base rates and likelihood ratios.

MOMENTS AND AGGREGATION

Derive expectation, variance, covariance, and correlation. Explain why means add without independence but portfolio variance depends on covariance.

ESTIMATION AND VALIDATION

Explain LLN, CLT, Monte Carlo error, effective sample size, calibration, proper scoring rules, rare-event uncertainty, and model risk.

DECISION TEST

Before using a probability, answer: what precisely is the event, what information is available, how much support exists, is the forecast calibrated, what is the payoff, and what breaks the model?