

Stationarity, Trends, and Regime Changes

Stationarity assumptions, trends, unit roots, breaks, regimes, and state-dependent validation.

Library of the Quant

Educational reference manual

Manual Profile

LIBRARY ID	LOTQ-0028
CATEGORY	Time Series
DIFFICULTY	Intermediate
FORMAT	Single-column field-guide manual
USE	Study, lookup, and research-system reference

Educational Use Only

This manual is educational material for quant research study. It is not financial advice, not a trading recommendation, and not a guarantee of correctness or live trading usefulness.

Table of contents

INTERMEDIATE

ORIENTATION

Title	1
How to use this manual	5
Notation and state contract	6
Stationarity and regime research workflow	7

STATIONARITY FOUNDATIONS

Stationarity is a modeling contract	8
Strict and weak stationarity	9
Mean stationarity	10
Variance stationarity	11
Lag-only autocovariance	12
Ergodicity and learning from one path	13
Local stationarity	14
Structural stability versus predictive stability	15

Table of contents

INTERMEDIATE

TRENDS AND TRANSFORMS

Deterministic trend	16
Stochastic trend	17
Unit roots and near-unit roots	18
First differencing	19
Log transforms and growth rates	20
Detrending and residual analysis	21
Seasonality and calendar structure	22
Cointegration and stable combinations	23
Fractional integration and long memory	24

TESTS AND DIAGNOSTICS

Augmented Dickey-Fuller test	25
Phillips-Perron test	26
KPSS test	27
Rolling estimates as visual diagnostics	28
Chow and known-break tests	29
Unknown changepoint detection	30
CUSUM monitoring	31
Breakpoint uncertainty	32

Table of contents

INTERMEDIATE

REGIME MODELS

A regime is a decision-relevant state	33
Volatility regimes	34
Markov-switching models	35
Hidden Markov models	36
Threshold autoregression	37
Feature-space clustering	38
Regime labels and validation	39
Regime uncertainty and soft decisions	40

RESEARCH AND SYSTEM DESIGN

Point-in-time regime features	41
Walk-forward validation under drift	42
Strategy robustness across regimes	43
Regime monitoring and response	44
Regime model governance	45

VISUAL LABS AND REVIEW

One path, four stationarity views	46
Choosing a stability treatment	47
ADF, PP, and KPSS evidence matrix	48
Break event, evidence, and action	49
Regime stability across choices	50
Method selection by failure mode	51
Failure modes and diagnostic actions	52
Formula sheet and interpretation map	53
Worked example: break to legal action	54
Final active recall and system index	55

How to use this manual

INTERMEDIATE

1. STATE THE INVARIANT

Name the mean, scale, dependence, distribution, or payoff mapping that a method requires to remain stable. Tie the statement to a horizon, universe, and decision rather than asking whether the chart is simply stationary.

2. SEPARATE PATH MECHANISMS

Distinguish deterministic trend, stochastic trend, seasonality, slow local drift, abrupt break, and persistent latent state. Similar plots can imply very different forecasts after a shock.

3. USE OPPOSING EVIDENCE

Combine plots, economic events, ADF or PP unit-root tests, KPSS stationarity-null tests, break diagnostics, and later forecast performance. Conflicting results describe uncertainty that the decision must absorb.

4. PRESERVE CAUSAL TIMING

Compute every trend, scale, detector, and regime probability from information available at the forecast origin. Keep retrospective smoothed states for explanation only, never as historical live features.

5. VALIDATE EPISODES

Count independent regimes and transitions, not only rows. Measure state uncertainty, detection lag, costs, boundary sensitivity, and the outcome of wrong-state decisions.

6. CONNECT EVIDENCE TO ACTION

Specify whether instability changes a transform, window, position size, model, risk limit, or fallback. A regime label without an approved action and owner is not an operational control.

DECISION STANDARD

Promote a state-dependent system only when its invariance target, data clock, transformations, tests, episode evidence, causal probability, action map, monitor, fallback, and rollback are reproducible.

Notation and state contract

INTERMEDIATE

SYMBOL	OBJECT	MEANING
X_t	observed process	Value available on the declared observation grid
μ_t	conditional center	Location estimate that may vary through time
σ_t	conditional scale	Point-in-time volatility or dispersion state
γ_k	lag covariance	Covariance at separation k under a stable model
B	lag operator	$B X_t = X_{(t-1)}$; differences use $1-B$
d	integration order	Number or fraction of differences used for stability
τ	break time	Estimated event boundary, distinct from detection time
S_t	latent regime	Unobserved state indexing model parameters
p_{ij}	transition probability	Chance of moving from state i to state j
F_t	information set	All values legally available at decision cutoff t

TWO CLOCKS FOR EVERY BREAK

The event time is when the market mechanism changed; the detection time is when accumulated evidence crossed the live threshold. Historical simulations can act only after detection, computation, decision, and execution latency.

MINIMUM STATE METADATA

Store feature cutoff, fit cutoff, transformation version, regime probability, label-alignment rule, detector state, event estimate, detection timestamp, confidence, authorized action, and fallback. A color-coded state alone cannot be replayed.

Stationarity and regime research workflow

INTERMEDIATE

1

Define the decision

Target, horizon, execution, loss, and risk consequence

2

Declare invariances

Mean, scale, dependence, tails, or payoff mapping

3

Validate the clock

Grid, calendars, availability, revisions, and quality

4

Diagnose the path

Trend, differences, seasonality, ACF, and scale

5

Challenge stability

Opposing nulls, events, breaks, and boundary ranges

6

Build causal state

Filtered probability, latency, confidence, and fallback

7

Walk forward

Refit, purge, score transitions, costs, and wrong states

8

Govern operation

Version, monitor, respond, roll back, and retire

HARD GATE

A regime result is not live-ready until retrospective labels have been replaced by causal state estimates, action timing is executable, episode-level validation is complete, and a conservative fallback has been tested.

Stationarity is a modeling contract

INTERMEDIATE

ONE-SENTENCE MEANING

Stationarity means the probabilistic behavior required by a model is stable under shifts in time, so estimates from one period can inform another. The exact requirement depends on whether the method uses means, variances, covariances, quantiles, or a fuller distribution.

MEMORY JOGGER

Do not ask whether a chart looks stationary in the abstract. Ask which feature of the process must remain invariant, over which horizon, and for which decision.

WHY IT MATTERS IN QUANT RESEARCH

Forecast coefficients, standard errors, risk estimates, and validation splits all borrow strength across observations. That borrowing becomes misleading when the generating mechanism drifts faster than the estimator can adapt.

FORMULA INTUITION

For weak stationarity, $E[X_t] = \mu$ and $\text{Cov}(X_t, X_{t-k}) = \gamma_k$ do not depend on calendar time. The statement concerns a process, while a realized sample offers only noisy evidence about it.

SIMPLE MARKET EXAMPLE

Daily equity returns may have roughly stable mean over a year while their volatility doubles during a crisis. A mean model can remain usable even when a constant-variance likelihood becomes badly calibrated.

PYTHON / SYSTEM IMPLEMENTATION

Define a stationarity contract beside every feature and model: target property, observation grid, window, tolerances, tests, and response. Store diagnostics by forecast origin rather than computing one full-history verdict.

CONFUSIONS TO AVOID

Stationarity is not synonymous with flat prices, normality, independence, or absence of seasonality. Passing one test never proves every property needed by a downstream model.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The data and feature registry declares which invariances each consumer assumes. Model validation then checks those properties on training folds, later folds, and operational monitoring windows.

RELATED CONCEPTS AND QUICK RECALL

Related: weak stationarity, invariance, forecast origin, model assumption. Recall task: Name the exact invariant required by a rolling z-score model and explain why a single full-sample test is insufficient.

Strict and weak stationarity

INTERMEDIATE

ONE-SENTENCE MEANING

Strict stationarity keeps every finite-dimensional joint distribution unchanged by a time shift. Weak stationarity keeps only the mean, variance, and lag covariance invariant.

MEMORY JOGGER

Strict means the whole probabilistic movie looks the same after shifting the clock. Weak means its center, scale, and linear dependence pattern stay the same.

WHY IT MATTERS IN QUANT RESEARCH

Linear forecasting, spectral analysis, and covariance estimation often rely on second-order structure rather than full distributions. Tail-risk and nonlinear decision rules can fail even when weak stationarity approximately holds.

FORMULA INTUITION

Strict: $(X_{t_1}, \dots, X_{t_m})$ has the same distribution as $(X_{(t_1+h)}, \dots, X_{(t_m+h)})$ for all valid m , times, and shifts h . Weak conditions alone do not constrain skewness, kurtosis, or tail dependence.

SIMPLE MARKET EXAMPLE

A return process can keep zero mean, constant variance, and zero autocorrelation while alternating between symmetric and negatively skewed conditional distributions. It is weakly stable in selected moments but not strictly stationary.

PYTHON / SYSTEM IMPLEMENTATION

Document which order of moments each estimator uses and compare conditional distribution diagnostics across eras. Use quantile, tail, and calibration monitors when decisions depend on more than covariance.

CONFUSIONS TO AVOID

For a Gaussian process, weak stationarity implies strict stationarity because mean and covariance determine the joint law. Extending that shortcut to heavy-tailed or mixture data is unjustified.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The assumption catalog maps model components to required distributional order. A mean forecaster, volatility model, and expected-shortfall engine can therefore have different health states on the same feed.

RELATED CONCEPTS AND QUICK RECALL

Related: joint distribution, Gaussian process, second-order structure, tail stability. Recall task: Construct a process with stable mean and variance but changing skewness, then identify which models would notice.

Mean stationarity

INTERMEDIATE

ONE-SENTENCE MEANING

Mean stationarity says the process has a time-invariant unconditional expectation over the population being modeled. In practice, researchers look for economically material drift relative to sampling uncertainty and the forecast horizon.

MEMORY JOGGER

A stable center does not mean every rolling average is flat. Sample means wander, so the question is whether their movement is plausible under one stable process or signals a changed center.

WHY IT MATTERS IN QUANT RESEARCH

Signals expressed as deviations from a historical mean depend directly on this assumption. A stale center can turn a normal repricing into a false contrarian trade or hide a persistent new premium.

FORMULA INTUITION

$E[X_t] = \mu$ for all t under the model. A rolling estimate $m_t = (1/L) \sum X_i$ is a noisy state estimate whose variance depends on L and serial dependence.

SIMPLE MARKET EXAMPLE

A value spread historically centered at zero shifts to 1.5% after a funding rule change. A strategy that continues buying at +1% treats the new equilibrium as a temporary deviation.

PYTHON / SYSTEM IMPLEMENTATION

Track rolling and expanding means with uncertainty bands, segment by known events, and compare forecast loss under frozen versus adaptive centers. Require a minimum effect size before changing production state.

CONFUSIONS TO AVOID

A smooth rolling mean can be created by overlapping windows even when the underlying process changes abruptly. Conversely, a short noisy excursion is not automatically a permanent mean break.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature service owns centering parameters and effective dates. Change-control records whether an update was scheduled, event-triggered, or a temporary degraded-mode adjustment.

RELATED CONCEPTS AND QUICK RECALL

Related: location shift, rolling mean, equilibrium level, adaptive baseline. Recall task: Explain how window length affects the delay and false-alarm rate of a mean-shift detector.

Variance stationarity

INTERMEDIATE

ONE-SENTENCE MEANING

Variance stationarity requires the unconditional dispersion around the mean to remain constant through time. Financial returns often violate it through volatility clustering, jumps, leverage effects, and liquidity episodes.

MEMORY JOGGER

A zero-centered series can still become much more dangerous when its scale changes. Separate direction from uncertainty before deciding that a process is stable.

WHY IT MATTERS IN QUANT RESEARCH

Risk limits, z-scores, likelihood weights, confidence intervals, and position sizes depend on scale. Underestimated variance creates excessive leverage and narrow prediction intervals exactly when losses accelerate.

FORMULA INTUITION

$\text{Var}(X_t) = \sigma^2$ under a constant-variance model. Conditional heteroskedastic models instead allow $\text{Var}(X_t | F_{t-1}) = \sigma_t^2$ while seeking stable dynamics for the variance state.

SIMPLE MARKET EXAMPLE

A strategy calibrated to 1% daily volatility enters a 3% regime without resizing. The same raw forecast produces three times the intended risk and a ninefold increase in variance contribution.

PYTHON / SYSTEM IMPLEMENTATION

Monitor robust rolling scale, squared-return dependence, interval coverage, and forecast-to-realized volatility ratios. Distinguish feed staleness from genuine calm by checking quote age and event counts.

CONFUSIONS TO AVOID

Changing variance does not necessarily imply changing expected return. Standardizing by a contemporaneous estimate can leak information if the estimate includes the outcome being predicted.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The risk-state service publishes point-in-time scale estimates and confidence flags to sizing, validation, and execution. Fallback logic applies conservative risk when the scale model becomes uncalibrated.

RELATED CONCEPTS AND QUICK RECALL

Related: heteroskedasticity, volatility clustering, scale break, conditional variance. Recall task: Calculate the risk error when true volatility triples but position size remains fixed, then name two live diagnostics.

Lag-only autocovariance

INTERMEDIATE

ONE-SENTENCE MEANING

A weakly stationary process has covariance between two observations that depends on their separation, not on their absolute dates. This makes lag patterns portable across forecast origins under the stated model.

MEMORY JOGGER

The same distance on the clock should imply the same linear relationship. If lag-one behavior differs before and after a market redesign, one global autocovariance is an average of incompatible states.

WHY IT MATTERS IN QUANT RESEARCH

AR models, standard errors, effective sample size, and spectral methods depend on stable lag structure. Shifting dependence can create coefficients that look precise but forecast no individual regime well.

FORMULA INTUITION

$\text{Cov}(X_t, X_{t-k}) = \gamma_k$ for every valid t . The autocorrelation $\rho_k = \gamma_k / \gamma_0$ standardizes the relation but inherits instability in both numerator and variance.

SIMPLE MARKET EXAMPLE

A spread shows positive lag-one persistence in calm markets and negative lag-one reversal during liquidation events. The full-sample correlation near zero hides both economically distinct behaviors.

PYTHON / SYSTEM IMPLEMENTATION

Plot rolling lag estimates with pair counts, confidence intervals, and regime annotations. Recompute every window using only observations available by its endpoint and compare with state-conditional estimates.

CONFUSIONS TO AVOID

Overlapping returns, moving averages, and stale prices mechanically create autocorrelation. A stable-looking correlogram can describe the transformation pipeline rather than the economic process.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The dependence monitor tracks material lags required by each model, while the validation store retains window definitions and data-quality context. Alerts route to feature, data, or model owners based on cause.

RELATED CONCEPTS AND QUICK RECALL

Related: autocorrelation, serial dependence, overlap, spectral density. Recall task: Give one economic and one mechanical reason for changing lag-one correlation, and design a test that separates them.

Ergodicity and learning from one path

INTERMEDIATE

ONE-SENTENCE MEANING

Ergodicity allows time averages from a sufficiently long realization to converge to population quantities. Stationarity describes invariance, while ergodicity determines whether one observed market history can reveal that invariant structure.

MEMORY JOGGER

We usually observe one path through market history, not thousands of parallel universes. Ergodicity is the bridge from that path to an estimated population.

WHY IT MATTERS IN QUANT RESEARCH

Sample means, covariances, and model fits are useful only if accumulated observations provide new information. Persistent latent states and structural breaks can make a long dataset behave like a small effective sample.

FORMULA INTUITION

For an ergodic mean, $(1/T)\sum_{t=1}^T X_t$ converges to $E[X_t]$ as T grows. Strong serial dependence slows convergence and widens uncertainty relative to an IID calculation.

SIMPLE MARKET EXAMPLE

Ten years of daily returns may contain only a few independent volatility episodes. Counting 2,500 rows as 2,500 equivalent regime observations creates false confidence about crisis behavior.

PYTHON / SYSTEM IMPLEMENTATION

Report effective sample size, block-bootstrap uncertainty, regime counts, and event diversity. Prefer evaluation across independent instruments or later episodes when the economic mechanism permits pooling.

CONFUSIONS TO AVOID

A stationary process need not be ergodic, and a visually long history need not contain broad state coverage. More high-frequency rows inside one event do not replace new independent regimes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The evidence registry stores row count, effective count, regime coverage, and dependence-adjusted uncertainty. Promotion rules can require minimum event diversity rather than only a date span.

RELATED CONCEPTS AND QUICK RECALL

Related: effective sample size, mixing, convergence, block bootstrap. Recall task: Explain why 2,500 daily observations may provide far fewer than 2,500 independent pieces of evidence.

Local stationarity

INTERMEDIATE

ONE-SENTENCE MEANING

Local stationarity treats a process as approximately stable inside a limited neighborhood even though its behavior evolves over longer horizons. It is a pragmatic approximation, not permission to ignore drift.

MEMORY JOGGER

Zoom in until the process changes slowly relative to the estimation window. The window must still contain enough information to estimate the state with acceptable noise.

WHY IT MATTERS IN QUANT RESEARCH

Many financial models are useful only locally because market structure, participation, and volatility evolve. The central design problem becomes balancing adaptation speed against estimation variance.

FORMULA INTUITION

A time-varying representation can write $X_u(t, T)$ with parameters $\theta_u(t/T)$ that change smoothly in rescaled time. Within a short block around u , θ_u is approximated as constant.

SIMPLE MARKET EXAMPLE

A 60-day volatility model adapts to a new regime faster than a five-year estimate but becomes noisy during quiet periods. A 250-day window is steadier yet can remain stale for months after a break.

PYTHON / SYSTEM IMPLEMENTATION

Evaluate several predeclared windows by state-estimation error, decision turnover, and out-of-sample loss. Use shrinkage toward long-run parameters rather than choosing a window from the final Sharpe ratio.

CONFUSIONS TO AVOID

Local stationarity does not solve abrupt breaks or sparse tails. Continuously retuning the window on recent outcomes can chase noise and consume the holdout period.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The state-estimation layer versions window policy, shrinkage target, refresh cadence, and warm-up behavior. Monitoring compares fast and slow estimates to detect meaningful divergence.

RELATED CONCEPTS AND QUICK RECALL

Related: adaptive window, time-varying parameter, bias-variance tradeoff, forgetting factor. Recall task: Choose a window policy for a fast volatility shift and state how you would prevent adaptation from becoming overfitting.

Structural stability versus predictive stability

INTERMEDIATE

ONE-SENTENCE MEANING

Structural stability concerns whether the data-generating relationships remain consistent, while predictive stability concerns whether a model continues to deliver calibrated out-of-sample decisions. One can fail before the other becomes visible.

MEMORY JOGGER

A model may keep scoring well through offsetting shifts, or its coefficients may move while forecasts remain robust. Monitor mechanism and outcome separately.

WHY IT MATTERS IN QUANT RESEARCH

Quant systems need early warnings without triggering unnecessary retraining. Input drift, parameter drift, residual drift, and decision degradation provide different evidence and require different responses.

FORMULA INTUITION

Parameter stability asks $\theta_t = \theta$; predictive stability asks whether loss $L(y_t, f_t)$ and calibration remain within approved bounds. Neither equality is observed without estimation error.

SIMPLE MARKET EXAMPLE

A trend feature weakens while a volatility feature strengthens, leaving portfolio Sharpe unchanged. Aggregate performance hides a changed dependency that may fail when volatility normalizes.

PYTHON / SYSTEM IMPLEMENTATION

Track feature distributions, coefficients, residuals, forecast calibration, turnover, and exposure attribution on aligned windows. Define alert tiers based on materiality and persistence rather than any single p-value.

CONFUSIONS TO AVOID

Stable PnL does not prove stable structure, and changing coefficients do not automatically justify intervention. Repeated testing of monitoring windows creates its own false-alarm problem.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The control plane separates observation, warning, degraded mode, and hard stop. Each state has an owner, evidence threshold, conservative fallback, and rollback path.

RELATED CONCEPTS AND QUICK RECALL

Related: concept drift, calibration, parameter drift, model health. Recall task: Describe a case where forecast performance is stable but structural risk is rising, then specify an appropriate response.

Deterministic trend

INTERMEDIATE

ONE-SENTENCE MEANING

A deterministic trend is a predictable function of time around which deviations may be stationary. Once the trend specification is known or estimated, future trend values are mechanically extrapolated rather than accumulated from random shocks.

MEMORY JOGGER

The path has a moving center, but the movement follows a stated clock pattern. Remove the clock component and study whether the remainder behaves stably.

WHY IT MATTERS IN QUANT RESEARCH

Trend regressions, capacity growth, inflation adjustment, and slowly changing market participation can require deterministic terms. An omitted trend creates serial residuals and exaggerated evidence of persistence.

FORMULA INTUITION

$X_t = a + b \cdot t + u_t$ with stationary u_t is linear trend-stationary. A polynomial or spline replaces $b \cdot t$ when theory supports curvature, but added flexibility increases extrapolation risk.

SIMPLE MARKET EXAMPLE

Average daily volume rises steadily as an ETF gains adoption, while deviations around the adoption path revert. Differencing removes the slope but may add noise compared with a defensible trend model.

PYTHON / SYSTEM IMPLEMENTATION

Fit trend terms inside each training fold, archive the time origin and basis, and transform validation data without refitting. Compare residual ACF and forecast loss against naive differencing.

CONFUSIONS TO AVOID

A fitted line can absorb a one-time level shift and falsely suggest gradual growth. High in-sample R-squared against time does not establish a stable future trend.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature pipeline stores detrending parameters with their fit cutoff. Forecast services receive both trend forecast and residual component so reconstruction remains auditable.

RELATED CONCEPTS AND QUICK RECALL

Related: trend-stationary process, time regression, residual stationarity, extrapolation. Recall task: Contrast the forecast after a temporary shock for a deterministic-trend model and a random-walk model.

Stochastic trend

INTERMEDIATE

ONE-SENTENCE MEANING

A stochastic trend is a persistent path generated by accumulated shocks, so innovations have lasting effects on the level. It is typically removed by differencing rather than by subtracting a fixed function of time.

MEMORY JOGGER

The trend itself wanders because each surprise becomes part of tomorrow's baseline. There is no predetermined line to which the level must return.

WHY IT MATTERS IN QUANT RESEARCH

Asset prices, some macro levels, and cumulative flows often behave more like integrated processes than stable deviations around a line. Modeling them as mean-reverting levels can produce hazardous contrarian positions.

FORMULA INTUITION

$X_t = X_{(t-1)} + c + \epsilon_t$ implies $\Delta X_t = c + \epsilon_t$. The variance of X_t grows with horizon because shocks accumulate instead of decaying.

SIMPLE MARKET EXAMPLE

A stock falls after permanent bad news and then fluctuates around the lower level. A deterministic-trend residual model keeps forecasting recovery toward a line that no longer represents value.

PYTHON / SYSTEM IMPLEMENTATION

Compare level and first-difference diagnostics, use random-walk baselines, and evaluate long-horizon forecast variance. Preserve the integration choice in the model artifact and reintegrate forecasts carefully.

CONFUSIONS TO AVOID

A stochastic trend is not a smooth trend, and differencing does not prove economic efficiency. Some level combinations can remain stable through cointegration even when individual levels are integrated.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The target registry declares whether models consume levels, differences, or equilibrium residuals. Backtests reconcile transformed forecasts to economically executable level decisions.

RELATED CONCEPTS AND QUICK RECALL

Related: integrated process, permanent shock, random walk, differencing. Recall task: Explain why the forecast variance of a random-walk level expands with horizon and why a trend line is misleading.

Unit roots and near-unit roots

INTERMEDIATE

ONE-SENTENCE MEANING

A unit root occurs when persistence is strong enough that shocks do not decay, while a near-unit root decays so slowly that finite samples can look integrated. Distinguishing them is statistically difficult and economically consequential.

MEMORY JOGGER

When ϕ is one, the process remembers a shock forever; when ϕ is 0.99, it forgets only after a long time. Short histories struggle to tell those stories apart.

WHY IT MATTERS IN QUANT RESEARCH

Mean-reversion speed, forecast intervals, spread trading, and long-horizon risk depend on the persistence classification. Overconfident rejection of a unit root can lead to oversized convergence bets.

FORMULA INTUITION

In $X_t = c + \phi X_{t-1} + \epsilon_t$, $\phi=1$ is a unit root. For $|\phi| < 1$ the half-life is $\log(0.5)/\log(|\phi|)$, which becomes extremely long as ϕ approaches one.

SIMPLE MARKET EXAMPLE

A spread estimated with $\phi=0.98$ has a half-life near 34 periods, while $\phi=0.995$ implies about 138 periods. Transaction costs and structural change can dominate before convergence arrives.

PYTHON / SYSTEM IMPLEMENTATION

Report coefficient uncertainty, implied half-life intervals, test sensitivity, and out-of-sample convergence. Simulate the fitted near-unit process to understand the detector's finite-sample behavior.

CONFUSIONS TO AVOID

Failing to reject a unit root is not proof of one, and rejecting it does not guarantee fast or tradable mean reversion. Standard critical values fail when deterministic terms are chosen incorrectly.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The signal-review service translates persistence into holding period, capital usage, stop policy, and regime risk. Promotion limits exposure when half-life uncertainty spans impractical horizons.

RELATED CONCEPTS AND QUICK RECALL

Related: AR(1), half-life, local-to-unity, persistence uncertainty. Recall task: Compute two half-lives for nearby ϕ values and explain why statistical stationarity may still be operationally unusable.

First differencing

INTERMEDIATE

ONE-SENTENCE MEANING

First differencing replaces a level with the change between adjacent observations and removes one unit root under common integrated-process models. It changes both the statistical object and the economic question.

MEMORY JOGGER

Ask what changed rather than where the level stands. Differencing turns accumulated shocks into individual increments but discards long-run level information.

WHY IT MATTERS IN QUANT RESEARCH

Returns, changes in yields, and changes in macro series are often more stable than their levels. Correct differencing can prevent spurious regressions and make baselines realistic.

FORMULA INTUITION

$\Delta X_t = X_t - X_{t-1}$. If $(1-B)X_t$ is stationary, X_t is integrated of order one, written $I(1)$, with B denoting the lag operator.

SIMPLE MARKET EXAMPLE

A yield moves from 4.00% to 4.08%; its difference is +8 basis points. Modeling the change addresses next-period movement, whereas modeling the level addresses the yield state itself.

PYTHON / SYSTEM IMPLEMENTATION

Compute differences after calendar and corporate-action validation, retain the original level, and unit-test interval alignment. Fit scalars and imputers to differenced training data only.

CONFUSIONS TO AVOID

Over-differencing introduces unnecessary negative autocorrelation and amplifies measurement noise. A missing observation makes the next valid difference span a longer interval unless the grid is explicitly restored.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The transformation registry records order, lag, units, gap policy, and inverse mapping. Downstream consumers declare whether they require change forecasts or reconstructed levels.

RELATED CONCEPTS AND QUICK RECALL

Related: lag operator, integration order, over-differencing, inverse transform. Recall task: Show how an irregular time gap changes the meaning of a first difference and state how the pipeline should record it.

Log transforms and growth rates

INTERMEDIATE

ONE-SENTENCE MEANING

A log transform converts multiplicative scale into additive differences and can stabilize proportional variation when values are positive. It does not by itself remove a stochastic trend or guarantee constant variance.

MEMORY JOGGER

Equal percentage moves become equal vertical distances in log space. That helps with growth interpretation, but the clock behavior still needs separate diagnosis.

WHY IT MATTERS IN QUANT RESEARCH

Prices, volumes, sizes, and economic levels often vary in percentage rather than absolute units. Logs can reduce right skew and make trend coefficients interpretable as approximate growth rates.

FORMULA INTUITION

$Y_t = \log(X_t)$, and $\Delta Y_t = \log(X_t/X_{t-1})$. For small changes, the log difference approximates the simple growth rate but diverges for large moves.

SIMPLE MARKET EXAMPLE

Volume doubling from one to two million and from five to ten million gives the same log change, even though the absolute changes differ by four million shares.

PYTHON / SYSTEM IMPLEMENTATION

Reject or separately model nonpositive values, record the offset if one is economically justified, and compare residual scale across level bands. Reconcile cumulative log changes to the final level ratio.

CONFUSIONS TO AVOID

Adding an arbitrary constant before logging changes interpretation and can dominate small values. A stable log variance can coexist with breaks in mean, dependence, or tail shape.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature contract labels raw, log-level, and log-difference columns distinctly. Monitoring tracks domain violations and reconstruction error rather than silently clipping inputs.

RELATED CONCEPTS AND QUICK RECALL

Related: Box-Cox transform, proportional change, variance stabilization, log return. Recall task: State when logging is economically natural and why log-level stationarity still must be tested.

Detrending and residual analysis

INTERMEDIATE

ONE-SENTENCE MEANING

Detrending subtracts an estimated systematic path so the remaining component can be modeled separately. Its validity depends on the trend family, fit window, and whether the trend is stable out of sample.

MEMORY JOGGER

Separate the moving baseline from deviations around it. Then test the deviations rather than declaring the original level stable.

WHY IT MATTERS IN QUANT RESEARCH

Residual signals, relative-value spreads, and anomaly thresholds often rely on detrended inputs. A poor trend fit can manufacture apparent cycles or erase real slow-moving alpha.

FORMULA INTUITION

$\hat{u}_t = X_t - \hat{m}_t$, where $\hat{m}_t$ may be linear, spline, filter, or economically specified. Uncertainty in $\hat{m}_t$ carries into residual statistics and forecasts.

SIMPLE MARKET EXAMPLE

A smooth moving average fitted symmetrically around each date creates attractive residual reversals because future observations help define the historical baseline. The live one-sided filter behaves differently.

PYTHON / SYSTEM IMPLEMENTATION

Fit trend operators inside chronological folds, use causal one-sided filters for live features, and save boundary behavior. Compare residual distribution and ACF across forecast origins.

CONFUSIONS TO AVOID

Two-sided filters leak future data, while overly flexible splines can treat temporary shocks as trend. Detrended residual stationarity does not prove the trend forecast remains valid.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The preprocessing artifact stores trend basis, parameters, timestamps, and causal mode. The forecast ledger records trend and residual contributions to every final prediction.

RELATED CONCEPTS AND QUICK RECALL

Related: trend filter, residual, one-sided smoother, leakage. Recall task: Explain why a centered moving average is unsuitable for a historical live signal and propose a causal alternative.

Seasonality and calendar structure

INTERMEDIATE

ONE-SENTENCE MEANING

Seasonality is a repeatable pattern tied to a known calendar or event phase, such as time of day, weekday, month, contract cycle, or announcement schedule. A stationary residual can remain after the seasonal mean and scale are modeled.

MEMORY JOGGER

The process may be stable at the same phase, not at every timestamp. Compare like with like before interpreting apparent drift.

WHY IT MATTERS IN QUANT RESEARCH

Intraday volume, spread, volatility, and returns often have strong clock patterns. Ignoring them creates false anomalies at the open and understated risk around scheduled events.

FORMULA INTUITION

$X_{t-s}(\text{phase}(t)) + u_t$, or both mean and scale may vary by phase. Seasonal differencing uses $X_t - X_{t-s}$ when an appropriate fixed period s exists.

SIMPLE MARKET EXAMPLE

A bid-ask spread of eight basis points is unusual at midday but normal during the opening minute. One unconditional threshold misclassifies both phases.

PYTHON / SYSTEM IMPLEMENTATION

Estimate seasonal profiles from training history with robust pooling, require phase coverage, and update on a fixed schedule. Validate daylight-saving, holiday, half-day, and venue-calendar behavior.

CONFUSIONS TO AVOID

Seasonality is not a deterministic guarantee, and calendar bins can drift after market-design changes. Seasonal differencing can remove economically meaningful event effects or create gaps around irregular calendars.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The calendar service provides phase keys and session metadata to feature and monitor layers. Anomaly scores use phase-conditional references with explicit fallback for rare bins.

RELATED CONCEPTS AND QUICK RECALL

Related: seasonal adjustment, intraday pattern, calendar effect, phase conditioning. Recall task: Design a phase-conditional spread anomaly and specify how holidays and daylight-saving transitions are handled.

Cointegration and stable combinations

INTERMEDIATE

ONE-SENTENCE MEANING

Cointegration occurs when individually nonstationary series have a linear combination that is stationary, representing a persistent long-run equilibrium. The stable object is the spread, not either level alone.

MEMORY JOGGER

Two wandering paths can share a tether. Estimate the tether carefully, then model deviations only if the relation survives later data and trading frictions.

WHY IT MATTERS IN QUANT RESEARCH

Pairs, curve, basis, and cross-market relative-value strategies often depend on equilibrium relationships. Cointegration provides a more coherent foundation than high level correlation alone.

FORMULA INTUITION

If X_t and Y_t are $I(1)$ but $z_t = Y_t - \beta X_t$ is $I(0)$, they are cointegrated with vector $(1, -\beta)$. An error-correction model links short-run changes to lagged disequilibrium.

SIMPLE MARKET EXAMPLE

Two share classes track the same issuer with different liquidity. Their levels wander, yet a hedged price combination can revert until borrow, conversion, or market access changes.

PYTHON / SYSTEM IMPLEMENTATION

Estimate hedge ratios only on training data, test residual stability, include deterministic terms deliberately, and validate half-life, turnover, borrow, and break risk out of sample.

CONFUSIONS TO AVOID

Correlation of trending levels can be high without cointegration. Re-estimating beta on the full sample makes the spread look artificially stable and leaks future structural information.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The spread-model artifact binds universe eligibility, hedge coefficients, fit window, break monitor, and exit policy. Execution consumes actual leg prices and liquidity rather than only normalized residuals.

RELATED CONCEPTS AND QUICK RECALL

Related: error correction, equilibrium residual, spurious regression, pairs trading. Recall task: Explain why high correlation is insufficient for a pairs trade and how the hedge ratio must be frozen in validation.

Fractional integration and long memory

INTERMEDIATE

ONE-SENTENCE MEANING

Fractional integration describes persistence between ordinary stationarity and a unit root by allowing a non-integer differencing order. Shocks decay slowly, creating long memory without permanent accumulation.

MEMORY JOGGER

Dependence fades like a long tail rather than disappearing quickly or lasting forever. A short window can mistake that slow decay for a changing regime.

WHY IT MATTERS IN QUANT RESEARCH

Volatility, order flow, and some spreads show persistent autocorrelation across many lags. Long-memory models can improve horizon scaling but are vulnerable to structural breaks that mimic persistence.

FORMULA INTUITION

$(1-B)^d X_t = \epsilon_t$ with fractional d . Roughly, $0 < d < 0.5$ can be covariance-stationary with long memory, while larger d changes stationarity properties.

SIMPLE MARKET EXAMPLE

Absolute returns remain autocorrelated for months after volatility spikes. A fractional model explains gradual decay, whereas a two-regime model explains persistence through long state durations.

PYTHON / SYSTEM IMPLEMENTATION

Compare fractional, break, and regime-switching challengers on later periods. Estimate d with uncertainty, inspect low-frequency sensitivity, and avoid using a single log-periodogram slope as proof.

CONFUSIONS TO AVOID

Structural shifts, seasonality, aggregation, and smooth trends can all imitate long memory. Fractional differencing may preserve more level information but complicates warm-up and live state management.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model registry stores differencing coefficients, truncation length, initialization history, and reconstruction logic. Monitoring checks whether estimated memory remains stable across expanding windows.

RELATED CONCEPTS AND QUICK RECALL

Related: long-range dependence, Hurst behavior, fractional differencing, structural break. Recall task: Describe how a sequence of volatility breaks can masquerade as long memory and design a challenger comparison.

Augmented Dickey-Fuller test

INTERMEDIATE

ONE-SENTENCE MEANING

The augmented Dickey-Fuller test asks whether an autoregressive root is consistent with one after adding lagged differences to absorb short-run serial correlation. Its null is a unit root, not general nonstationarity.

MEMORY JOGGER

Evidence must be strong enough to reject permanent persistence. The deterministic terms and augmentation lags define the question before the statistic is read.

WHY IT MATTERS IN QUANT RESEARCH

ADF is a common screen before level regression, mean-reversion research, or differencing. Used responsibly, it narrows a modeling choice rather than issuing a universal stationary-or-not label.

FORMULA INTUITION

$\Delta X_t = a + b t + \Delta X_{(t-1)} + \sum \phi_i \Delta X_{(t-i)} + \epsilon_t$; test $H_0: \delta = 0$ against $\delta < 0$. Critical values are nonstandard.

SIMPLE MARKET EXAMPLE

A spread rejects the unit-root null with an intercept but not after a trend is added. That disagreement reveals sensitivity to whether deterministic drift belongs in the maintained model.

PYTHON / SYSTEM IMPLEMENTATION

Choose intercept, trend, and lag policy before viewing the result; report test statistic, critical values, sample, and residual diagnostics. Repeat only as a documented robustness analysis.

CONFUSIONS TO AVOID

A small p-value does not prove strict stationarity, fast reversion, or trading profitability. Data-driven selection of trend and lags after seeing the p-value invalidates the advertised evidence.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The validation report stores test specification and links it to the transformation decision. Promotion also requires later residual behavior, half-life uncertainty, and break resistance.

RELATED CONCEPTS AND QUICK RECALL

Related: unit-root null, deterministic term, lag augmentation, nonstandard critical value. Recall task: State the ADF null and explain how choosing an intercept versus a trend changes the economic question.

Phillips-Perron test

INTERMEDIATE

ONE-SENTENCE MEANING

The Phillips-Perron test evaluates a unit-root null while adjusting the test statistic nonparametrically for serial correlation and heteroskedasticity in the errors. It offers a different nuisance correction from ADF.

MEMORY JOGGER

Keep a simple autoregression, then repair its inference for dependent, uneven noise. Agreement with ADF increases confidence; disagreement points to specification sensitivity.

WHY IT MATTERS IN QUANT RESEARCH

Financial increments frequently show heteroskedasticity, making a robust unit-root challenger useful. PP is still a test of one null under finite-sample constraints, not an automatic remedy for market complexity.

FORMULA INTUITION

The regression resembles a Dickey-Fuller form, while long-run variance estimates adjust the statistic. Bandwidth and kernel choices affect the correction and therefore the result.

SIMPLE MARKET EXAMPLE

A rate spread rejects under ADF with several lagged differences but not under PP with a broad bandwidth. The conflict suggests that short-run dynamics and variance estimation materially influence inference.

PYTHON / SYSTEM IMPLEMENTATION

Predeclare bandwidth policy, deterministic terms, and sample treatment; report PP beside ADF and simulated size under plausible fitted processes. Preserve exact library version and options.

CONFUSIONS TO AVOID

Heteroskedasticity-robust does not mean robust to breaks, outliers, nonlinear mean reversion, or incorrect deterministic terms. Treating any one test as decisive hides model uncertainty.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The test suite exposes multiple governed specifications as a sensitivity panel. A human review resolves disagreements using economics, plots, and out-of-sample behavior rather than majority voting.

RELATED CONCEPTS AND QUICK RECALL

Related: long-run variance, bandwidth, nuisance correction, unit-root testing. Recall task: Explain why ADF and PP can disagree and what evidence you would inspect before transforming the series.

ONE-SENTENCE MEANING

The KPSS test takes level or trend stationarity as its null and looks for evidence that a random-walk component is present. It reverses the burden of proof relative to ADF and PP.

MEMORY JOGGER

Use opposing nulls as a decision matrix. One test asks whether persistence is permanent; the other asks whether a stable baseline is still plausible.

WHY IT MATTERS IN QUANT RESEARCH

Combining unit-root and stationarity-null tests clarifies ambiguous cases caused by low power. The result can distinguish strong evidence from a zone where more data or conservative modeling is needed.

FORMULA INTUITION

KPSS decomposes X_t into deterministic terms, a random walk, and stationary error; a statistic based on cumulative residual sums is scaled by a long-run variance estimate.

SIMPLE MARKET EXAMPLE

ADF fails to reject a unit root and KPSS rejects stationarity, supporting differencing. If both fail to reject, evidence is weak rather than contradictory proof of two states.

PYTHON / SYSTEM IMPLEMENTATION

Match level versus trend-stationary specification to the model, publish bandwidth sensitivity, and interpret the four-cell ADF-KPSS matrix. Simulate near-unit processes to calibrate practical uncertainty.

CONFUSIONS TO AVOID

Failing to reject stationarity is not proof that parameters are constant. KPSS can over-reject under breaks, outliers, or poorly estimated long-run variance.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The transformation decision record includes both null directions and a conservative action for inconclusive cells. Production can retain level and difference challengers until later evidence separates them.

RELATED CONCEPTS AND QUICK RECALL

Related: stationarity null, complementary testing, cumulative residual, inconclusive evidence. Recall task: Interpret all four possible ADF-KPSS outcomes without calling failure to reject a proof.

Rolling estimates as visual diagnostics

INTERMEDIATE

ONE-SENTENCE MEANING

Rolling estimates reveal how local means, scales, dependence, and model errors evolve across forecast origins. They are exploratory state summaries whose smoothness and uncertainty depend strongly on overlapping windows.

MEMORY JOGGER

Slide the same lens through history and watch what moves. Remember that adjacent estimates share most observations and therefore are not independent confirmations.

WHY IT MATTERS IN QUANT RESEARCH

Rolling plots can localize instability, connect it to events, and reveal adaptation delay. They are especially useful for comparing fast and slow state estimates before formal break modeling.

FORMULA INTUITION

$\theta_{\hat{t}} = T(X_{(t-L+1)}, \dots, X_t)$. Adjacent L -period windows overlap in $L-1$ observations, creating high serial dependence even if X_t itself is independent.

SIMPLE MARKET EXAMPLE

A 252-day beta appears to decline smoothly after a sector reclassification. The smooth path partly reflects replacing one observation per day rather than a truly gradual economic transition.

PYTHON / SYSTEM IMPLEMENTATION

Show confidence bands, window length, pair counts, and event markers; compare several predeclared L values. Compute live-compatible estimates that stop at the availability cutoff.

CONFUSIONS TO AVOID

Selecting the most dramatic window after inspection is multiple testing. Narrow windows exaggerate noise, broad windows blur breaks, and centered windows leak future observations.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The monitoring layer stores rolling estimates as derived artifacts with lineage and effective sample sizes. Alerts require persistent, material divergence across approved windows.

RELATED CONCEPTS AND QUICK RECALL

Related: moving window, overlap dependence, state estimate, event annotation. Recall task: Explain why a smooth rolling statistic can accompany an abrupt break and how uncertainty should be displayed.

Chow and known-break tests

INTERMEDIATE

ONE-SENTENCE MEANING

A known-break test compares model fit before and after a prespecified candidate date to determine whether parameters differ more than sampling noise would suggest. Its validity is strongest when the date came from external information.

MEMORY JOGGER

If a rule change has a clear effective date, compare one model across the boundary with separate models on each side. The date must not be selected because it produced the largest test statistic.

WHY IT MATTERS IN QUANT RESEARCH

Exchange redesigns, index methodology changes, accounting rules, and strategy launches create natural break hypotheses. Testing those dates connects statistical instability to a plausible mechanism.

FORMULA INTUITION

For linear regression, a Chow-type F statistic compares pooled residual sum of squares with the sum from two segment fits, adjusted for parameter count and segment sample sizes.

SIMPLE MARKET EXAMPLE

Tick-size reduction changes spread behavior on the implementation date. Separate pre- and post-rule regressions materially reduce residual error and reveal a new coefficient structure.

PYTHON / SYSTEM IMPLEMENTATION

Register the event date and affected parameters, require enough observations on both sides, and inspect variance changes and residual dependence. Use robust alternatives when assumptions fail.

CONFUSIONS TO AVOID

Searching all dates and reporting a standard Chow p-value ignores the selection process. A significant break can reflect variance or outliers even when only mean coefficients are tested.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The event catalog provides prespecified dates to validation jobs. Confirmed breaks trigger version boundaries in training data, features, and monitoring references.

RELATED CONCEPTS AND QUICK RECALL

Related: parameter stability, event study, structural break, segmented regression. Recall task: Design a known-break analysis for a tick-size change and state what makes the date exogenous.

Unknown changepoint detection

INTERMEDIATE

ONE-SENTENCE MEANING

Unknown changepoint detection searches for one or more times at which a process parameter changes. Because the location is estimated from the data, thresholds must account for the search and any penalty on complexity.

MEMORY JOGGER

Find where one stable description stops and another starts. More breaks always improve in-sample fit, so the penalty and minimum segment length are essential parts of the method.

WHY IT MATTERS IN QUANT RESEARCH

Changepoints can flag new volatility states, liquidity changes, data defects, or strategy decay. Their value lies in faster diagnosis and safer adaptation, not in retroactively optimizing regimes.

FORMULA INTUITION

A common objective minimizes sum over segments of cost(data segment) plus beta times the number of breaks. Cost can target mean, variance, regression coefficients, or a distribution.

SIMPLE MARKET EXAMPLE

A variance detector finds a sharp jump during a crisis, while a mean detector finds none. The difference matters because the action should resize risk rather than reverse the expected-return signal.

PYTHON / SYSTEM IMPLEMENTATION

Choose target parameter, penalty, minimum length, and online delay before evaluation. Measure false alarms and detection lag on simulations containing realistic dependence, tails, and missingness.

CONFUSIONS TO AVOID

An algorithm will label outliers, seasonality, or feed changes as breaks if the cost model cannot represent them. Retrospective changepoints cannot be used as if known in real time.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The regime service stores detection time separately from estimated event time. Downstream backtests may act only after detection latency and confirmation rules are satisfied.

RELATED CONCEPTS AND QUICK RECALL

Related: segmentation, break penalty, detection delay, minimum segment. Recall task: Explain why retrospective break time and live action time differ, then define both for a volatility jump.

CUSUM monitoring

INTERMEDIATE

ONE-SENTENCE MEANING

CUSUM accumulates signed deviations from a reference level so small persistent shifts become detectable even when individual observations look ordinary. A threshold converts cumulative evidence into an alert.

MEMORY JOGGER

Add each standardized surprise to a running balance. Random noise tends to cancel, while a sustained drift pushes the balance toward a boundary.

WHY IT MATTERS IN QUANT RESEARCH

CUSUM can monitor residual bias, execution slippage, quote latency, or calibration drift with low delay for specified shift sizes. It turns ongoing diagnostics into an explicit sequential policy.

FORMULA INTUITION

One-sided $C_t^+ = \max(0, C_{t-1}^+ + z_t - k)$; signal when $C_t^+ > h$. Reference k targets a shift size and h controls the false-alarm versus delay tradeoff.

SIMPLE MARKET EXAMPLE

Daily residuals average only 0.1 standard deviation above zero, too small to flag individually. Repetition accumulates and triggers a warning before the rolling mean becomes visually obvious.

PYTHON / SYSTEM IMPLEMENTATION

Standardize with a point-in-time scale estimate, calibrate k and h by simulation, and reset only under a documented policy. Track missing observations and state snapshots for replay.

CONFUSIONS TO AVOID

Autocorrelation and changing variance distort nominal false-alarm rates. Tuning thresholds on the incident that motivated the monitor overstates future protection.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The observability service runs CUSUM states beside batch drift metrics and connects warning levels to investigation, risk reduction, and hard-stop procedures.

RELATED CONCEPTS AND QUICK RECALL

Related: sequential detection, control chart, residual drift, average run length. Recall task: Choose k and h for a small positive bias and describe how autocorrelation changes calibration.

Breakpoint uncertainty

INTERMEDIATE

ONE-SENTENCE MEANING

A detected break has uncertainty in location, magnitude, duration, and affected parameters. Treating one estimated date as exact creates false precision in regime labels and subsequent validation.

MEMORY JOGGER

The evidence often supports a transition zone rather than one magic timestamp. Decision rules should remain robust across plausible boundaries.

WHY IT MATTERS IN QUANT RESEARCH

Performance attribution, sample segmentation, and retraining choices can change materially when a boundary moves by weeks. Honest uncertainty prevents results from depending on a convenient cutoff.

FORMULA INTUITION

Profile loss evaluates candidate break dates τ by refitting the model at each τ ; near-minimum dates form a practical confidence region under suitable assumptions.

SIMPLE MARKET EXAMPLE

A correlation break is estimated on March 12, but dates from late February through early April have nearly equal objective values. A strategy profitable only with the central date is fragile.

PYTHON / SYSTEM IMPLEMENTATION

Plot objective versus candidate date, bootstrap boundaries with dependence-aware blocks, and rerun conclusions over the plausible region. Separate abrupt and gradual-transition alternatives.

CONFUSIONS TO AVOID

Standard bootstrap resampling destroys serial structure, while post-selection intervals can be too narrow. A wide date range is informative evidence about limited temporal resolution.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The regime registry stores boundary distributions or ranges, not only point labels. Validation gates require stability to boundary perturbation and action after realistic confirmation delay.

RELATED CONCEPTS AND QUICK RECALL

Related: profile likelihood, transition zone, boundary sensitivity, post-selection inference. Recall task: Move a candidate boundary across its uncertainty range and identify which strategy conclusions must remain stable.

A regime is a decision-relevant state

INTERMEDIATE

ONE-SENTENCE MEANING

A regime is a persistent state in which the distribution, dynamics, liquidity, or payoff mapping differs enough to change a research or risk decision. It should be defined by measurable behavior rather than a memorable narrative alone.

MEMORY JOGGER

Name the state only if it changes what the system estimates, trades, limits, or monitors. Labels such as risk-on are incomplete without observable criteria and transition rules.

WHY IT MATTERS IN QUANT RESEARCH

Regime conditioning can improve sizing, model selection, scenarios, and diagnostics when state differences are stable. It also adds estimation and selection risk because the current state is rarely known with certainty.

FORMULA INTUITION

A latent state S_t indexes parameters $\theta_{(S_t)}$, so $p(X_t|F_{(t-1)})$ differs by state. A useful regime has persistence, separability, and an operational action.

SIMPLE MARKET EXAMPLE

A momentum strategy works in high-dispersion trending months and loses in low-dispersion reversals. A two-state overlay reduces exposure only when the state estimate is available before trading.

PYTHON / SYSTEM IMPLEMENTATION

Write a regime charter containing observables, target parameters, minimum duration, detection latency, uncertainty, allowed actions, and retirement rule. Validate state-conditioned decisions on later transitions.

CONFUSIONS TO AVOID

Regimes are not objective natural kinds, and naming them after outcomes creates look-ahead bias. More labels make historical fit easier while leaving fewer independent episodes per state.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The regime service publishes probabilities and health metadata rather than an unqualified label. Consumers specify how confidence affects exposure and which fallback applies when state evidence is weak.

RELATED CONCEPTS AND QUICK RECALL

Related: latent state, state probability, persistence, regime charter. Recall task: Define one regime whose existence changes a concrete capital decision and list its measurable criteria.

Volatility regimes

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility regimes divide time into states with materially different return scale, tail frequency, correlation, and liquidity conditions. They are often the most observable regimes but still require a causal timing policy.

MEMORY JOGGER

The same expected return has a different risk meaning in calm and stressed markets. Scale first, then decide whether direction or dependence also changed.

WHY IT MATTERS IN QUANT RESEARCH

Position sizing, leverage, option exposure, stop distance, and covariance forecasts depend on volatility state. A slow estimator can remain overleveraged during transitions and underinvested during normalization.

FORMULA INTUITION

A simple state rule may compare point-in-time $\sigma_{\hat{t}}$ with fixed training quantiles; a probabilistic model estimates $P(S_t = \text{high} | F_t)$. Threshold uncertainty affects both turnover and delay.

SIMPLE MARKET EXAMPLE

Realized volatility jumps from 10% to 30% annualized, while forecast alpha stays constant. Risk targeting should cut position size to roughly one third before considering any alpha adjustment.

PYTHON / SYSTEM IMPLEMENTATION

Use lagged robust scale, forecast calibration, and hysteresis to avoid threshold chatter. Test transition loss, detection delay, costs, and crisis coverage rather than only average state classification.

CONFUSIONS TO AVOID

Defining regimes from future realized volatility leaks the outcome into historical decisions. A high-volatility state can also be caused by a bad tick or stale-price correction rather than market risk.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The risk-state artifact feeds sizing and scenario engines with probability, estimate age, and data quality. Conservative fallback uses the higher credible scale during uncertainty.

RELATED CONCEPTS AND QUICK RECALL

Related: risk targeting, hysteresis, realized volatility, scale state. Recall task: Design a live volatility-state rule and explain how its label differs from an ex post realized-volatility label.

Markov-switching models

INTERMEDIATE

ONE-SENTENCE MEANING

A Markov-switching model lets parameters depend on an unobserved state whose transition probabilities depend on the previous state. It jointly estimates state persistence, transition risk, and state-specific distributions.

MEMORY JOGGER

The market moves among hidden operating modes, and yesterday's mode informs today's probability. Observations provide noisy evidence rather than certain labels.

WHY IT MATTERS IN QUANT RESEARCH

Switching means, variances, or regression coefficients can represent recurring market behavior more coherently than arbitrary date splits. The model is useful only if probabilities improve later decisions after latency and costs.

FORMULA INTUITION

$P(S_t=j|S_{(t-1)=i})=p_{ij}$ and $X_t|S_t=j$ follows a distribution with parameters θ_j . Filtered probabilities use data through t ; smoothed probabilities also use future data.

SIMPLE MARKET EXAMPLE

A two-state return model estimates calm volatility at 8% and stressed volatility at 28%, with stressed-state persistence of 0.85. The filtered probability rises gradually after a shock.

PYTHON / SYSTEM IMPLEMENTATION

Fit several initializations, constrain identification, archive filtered probabilities, and score state-conditional density out of sample. Use only filtered or predicted probabilities for trading simulations.

CONFUSIONS TO AVOID

Smoothed state paths are excellent explanations but illegal live features. State labels can swap across fits, and a third transient state may be represented as extreme observations in two states.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The state-model artifact includes transition matrix, emission parameters, label-alignment rule, probability timestamp, and fallback. Monitoring checks occupancy, transition frequency, and likelihood calibration.

RELATED CONCEPTS AND QUICK RECALL

Related: transition matrix, filtered probability, smoothed probability, label switching. Recall task: Explain why a backtest using smoothed probabilities is biased and specify the probability available before an order.

Hidden Markov models

INTERMEDIATE

ONE-SENTENCE MEANING

A hidden Markov model represents observed features as emissions from a latent discrete state with Markov transitions. Unlike a simple switching return model, emissions can combine several market measurements to infer state.

MEMORY JOGGER

Infer an unseen market condition from a bundle of noisy clues, while requiring state changes to have temporal continuity. The feature bundle and emission distribution define what the state means.

WHY IT MATTERS IN QUANT RESEARCH

An HMM can combine volatility, trend strength, spread, and cross-sectional dispersion for a richer regime signal. Added dimensions increase sensitivity to scaling, missingness, and unstable feature relationships.

FORMULA INTUITION

$p(S_{1:T}, X_{1:T}) = p(S_1) \prod_t p(X_t | S_t) \prod_{t=1}^{T-1} p(S_{t+1} | S_t)$. Filtering marginalizes over possible prior paths.

SIMPLE MARKET EXAMPLE

A liquidity-stress state emits wide spreads, high volatility, and low depth. If only volatility is unusual, the model assigns partial rather than certain stress probability.

PYTHON / SYSTEM IMPLEMENTATION

Normalize features within training folds, compare emission families, fit multiple seeds, align labels by stable characteristics, and evaluate probability calibration on later episodes.

CONFUSIONS TO AVOID

Gaussian emissions can turn outliers into extra states, while too many states create attractive but unrepeated narratives. Missing-feature handling can itself become a state detector.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature and regime registries are versioned together because emission semantics depend on transforms. Production logs posterior probability, evidence contribution, and uncertainty for every decision.

RELATED CONCEPTS AND QUICK RECALL

Related: emission distribution, forward algorithm, multivariate state, posterior probability. Recall task: Choose three emissions for liquidity stress and describe how you would verify the state is not merely an outlier bucket.

Threshold autoregression

INTERMEDIATE

ONE-SENTENCE MEANING

A threshold autoregressive model changes dynamics when an observed transition variable crosses a boundary. The state is explicit and nonlinear, so reversion speed or sign can differ across regions.

MEMORY JOGGER

Behavior on one side of the threshold need not match behavior on the other. The threshold must be observable at decision time and stable enough to support separate estimates.

WHY IT MATTERS IN QUANT RESEARCH

Spreads may revert only when deviations are large enough to overcome costs, and volatility feedback may strengthen after losses. Threshold models express those asymmetric mechanisms directly.

FORMULA INTUITION

$X_t = a_j + \phi_j X_{t-1} + \epsilon_t$ when q_{t-d} lies in region j . Threshold q , delay d , and coefficients are estimated or prespecified.

SIMPLE MARKET EXAMPLE

A basis spread is nearly random inside plus or minus ten basis points but mean reverts beyond that band. The no-trade region reflects fees and inventory frictions.

PYTHON / SYSTEM IMPLEMENTATION

Estimate thresholds inside training folds, require enough observations and transitions per region, and compare with smooth-transition challengers. Include threshold uncertainty in trading rules.

CONFUSIONS TO AVOID

Choosing the threshold from final PnL overfits both dynamics and execution. Hard switching can cause excessive turnover when the transition variable oscillates near the boundary.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The signal artifact stores threshold, delay, hysteresis, region coefficients, and eligibility. The execution layer applies the same band and cost model used in research.

RELATED CONCEPTS AND QUICK RECALL

Related: SETAR, smooth transition, no-trade band, asymmetric dynamics. Recall task: Design a threshold spread rule with hysteresis and explain why the threshold cannot be selected on final Sharpe.

Feature-space clustering

INTERMEDIATE

ONE-SENTENCE MEANING

Feature-space clustering groups observations with similar measured states without imposing a temporal transition model. It can discover recurring configurations, but cluster identity depends on scaling, distance, and selected features.

MEMORY JOGGER

Group similar market snapshots, then ask whether the groups persist through time and lead to different decisions. A geometric cluster is not automatically an economic regime.

WHY IT MATTERS IN QUANT RESEARCH

Clustering helps explore combinations of volatility, correlation, dispersion, trend, and liquidity. It is useful for hypothesis generation and scenario libraries when followed by strict out-of-sample validation.

FORMULA INTUITION

K-means minimizes within-cluster squared distance to centroids; other algorithms change shape assumptions. Time persistence is absent unless explicitly added.

SIMPLE MARKET EXAMPLE

Snapshots cluster into calm trend, calm reversal, and stressed high-correlation groups. The stressed cluster is rare but dominates drawdowns and turnover costs.

PYTHON / SYSTEM IMPLEMENTATION

Fit scalars and clusters on training data, map later points to frozen clusters, test bootstrap stability, and report occupancy and transition behavior. Compare with simple rule-based states.

CONFUSIONS TO AVOID

Using future standardized means leaks information, and selecting k from strategy performance is data snooping. Nearby days can be split across clusters even when a persistent regime story is desired.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The exploratory store treats cluster models as provisional artifacts with feature lineage and stability scores. Only governed state definitions can control production capital.

RELATED CONCEPTS AND QUICK RECALL

Related: k-means, mixture model, distance metric, cluster stability. Recall task: Explain how scaling changes a market-state clustering and what evidence upgrades a cluster into a usable regime.

Regime labels and validation

INTERMEDIATE

ONE-SENTENCE MEANING

Regime validation asks whether state definitions, probabilities, and actions generalize to later episodes rather than merely explaining the fitted history. The unit of evidence is often a transition or episode, not a row.

MEMORY JOGGER

A label is useful only when the next state occurrence behaves similarly enough to support the same decision. Count independent episodes and preserve realistic detection delay.

WHY IT MATTERS IN QUANT RESEARCH

State-conditioned backtests can multiply apparent insights while shrinking samples. Validation must show incremental value beyond unconditional models and simple risk controls.

FORMULA INTUITION

Evaluate log score or classification calibration for states and decision loss for actions. Use episode-level resampling because adjacent rows within one regime are highly dependent.

SIMPLE MARKET EXAMPLE

A crisis overlay appears successful across 800 stress days, but 700 belong to one financial crisis. Its evidence base is closer to a few episodes than hundreds of independent observations.

PYTHON / SYSTEM IMPLEMENTATION

Freeze state definition, use filtered labels, test later episodes, and report transition confusion, occupancy, turnover, costs, and state-conditional uncertainty. Perturb boundaries and state count.

CONFUSIONS TO AVOID

Using ex post labels or optimizing actions separately in every state creates look-ahead and severe multiplicity. A rare regime can show spectacular Sharpe from one event.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The validation ledger links every state claim to episodes, independent transitions, and allowed actions. Promotion can limit exposure until new regimes accumulate evidence.

RELATED CONCEPTS AND QUICK RECALL

Related: episode count, filtered label, state calibration, conditional backtest. Recall task: Recalculate the effective evidence for a regime strategy when most state-days come from one episode.

Regime uncertainty and soft decisions

INTERMEDIATE

ONE-SENTENCE MEANING

Regime uncertainty recognizes that current state membership is probabilistic, delayed, and model-dependent. Soft decisions scale actions with evidence instead of pretending each timestamp has a perfectly known label.

MEMORY JOGGER

Treat state probability as confidence, not truth. The action should remain safe when probabilities are wrong or close to the decision boundary.

WHY IT MATTERS IN QUANT RESEARCH

Gradual exposure, mixtures of experts, and conservative risk overlays can reduce threshold chatter and transition losses. They also require calibration so probabilities have operational meaning.

FORMULA INTUITION

A mixture forecast is $\sum_j p_j(t) f_j(X_t)$, while exposure can be a bounded function of $p_j(t)$. Calibration compares predicted probabilities with realized state frequencies under an external definition.

SIMPLE MARKET EXAMPLE

With 60% stress probability, a risk system moves halfway between normal and stressed leverage rather than making a full switch. A hard stop remains available for severe observable breaches.

PYTHON / SYSTEM IMPLEMENTATION

Backtest probability-to-action maps fixed in advance, add hysteresis and caps, and inspect reliability curves. Stress the system under deliberately wrong state probabilities and delayed transitions.

CONFUSIONS TO AVOID

Model posterior probability is not automatically calibrated and can be overly sharp under misspecification. Smooth exposure can still trade excessively if probabilities fluctuate rapidly.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The decision layer receives a probability vector, confidence grade, and model-health state. Governance owns the bounded mapping from those values to capital rather than embedding it in research notebooks.

RELATED CONCEPTS AND QUICK RECALL

Related: mixture of experts, probability calibration, hysteresis, robust control. Recall task: Design a soft exposure map from stress probability and specify its cap, hysteresis, and hard-stop override.

Point-in-time regime features

INTERMEDIATE

ONE-SENTENCE MEANING

A point-in-time regime feature uses only observations, transformations, and fitted parameters available by the decision cutoff. This includes the delay required to finalize bars, volatility, macro releases, and model state.

MEMORY JOGGER

Ex post regime pictures are explanations, not tradable inputs. Move every event through availability, computation, decision, and execution timestamps.

WHY IT MATTERS IN QUANT RESEARCH

Regime leakage is especially dangerous because future data improves both the state label and the state-conditioned model. The resulting backtest can avoid transitions that were unknowable in real time.

FORMULA INTUITION

Require source event time $\leq$ availability $\leq$ feature cutoff $\leq$ decision $\leq$ order $\leq$ fill. Predicted state $P(S_{t|F}(t-1))$ differs from filtered state $P(S_{t|F_t})$ when X_t is not yet tradable.

SIMPLE MARKET EXAMPLE

A monthly regime uses month-end volatility and fills at that same close. Unless an auction order was placed before the cutoff using partial information, the state arrives too late for that price.

PYTHON / SYSTEM IMPLEMENTATION

Persist availability times, transform states, model version, probability timestamp, and earliest legal action. Add future-mutation tests that verify historical features remain unchanged.

CONFUSIONS TO AVOID

Shifting a regime column by one row does not solve time-zone, publication, revision, or bar-finalization delays. Smoothed probabilities use future evidence even when their timestamps look historical.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature platform enforces as-of joins and produces an audit record for every regime update. The simulator consumes earliest-action metadata to choose a legal fill model.

RELATED CONCEPTS AND QUICK RECALL

Related: information set, as-of join, filtered state, execution latency. Recall task: Audit one monthly regime decision from raw observations through earliest executable fill and identify every delay.

Walk-forward validation under drift

INTERMEDIATE

ONE-SENTENCE MEANING

Walk-forward validation repeatedly fits on past data and evaluates on the next unseen block, reproducing adaptation and state uncertainty through time. It measures whether the complete process works as conditions evolve.

MEMORY JOGGER

Rebuild the model at each historical forecast origin using only the history then available. Aggregate results without hiding which regimes and transitions produced them.

WHY IT MATTERS IN QUANT RESEARCH

Stationarity choices, break detectors, scalars, and regime models are all data-dependent preprocessing. Walk-forward testing prevents future structure from defining earlier transformations.

FORMULA INTUITION

For origins t_j , fit A_j on data through t_j , predict block $(t_j, t_{j+1}]$, then concatenate only out-of-sample predictions. Purge overlapping labels where needed.

SIMPLE MARKET EXAMPLE

A 500-day training window adapts quickly after a break but forgets older crises; an expanding window preserves crises but dilutes the new state. Walk-forward results reveal the operating tradeoff.

PYTHON / SYSTEM IMPLEMENTATION

Predeclare refit cadence, window policy, hyperparameter search, purge, and fallback. Report performance by origin, regime, transition, and calendar era with uncertainty.

CONFUSIONS TO AVOID

Tuning the window after viewing all fold results overfits the validation sequence. Random cross-validation destroys chronology and lets both regimes appear in every fold.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The experiment tracker stores one artifact lineage per origin and reconstructs the production state machine. Approval requires replayable forecasts, alerts, and actions, not only summary metrics.

RELATED CONCEPTS AND QUICK RECALL

Related: rolling origin, refit cadence, purging, temporal generalization. Recall task: Compare expanding and rolling windows after a break and state which metrics determine the production choice.

Strategy robustness across regimes

INTERMEDIATE

ONE-SENTENCE MEANING

Regime robustness evaluates whether a strategy's return, risk, cost, and exposure remain acceptable across materially different states and during transitions. It is a loss-control exercise before it is an optimization opportunity.

MEMORY JOGGER

First ask where the strategy breaks; only then ask whether state conditioning can improve it. A strategy requiring perfect regime timing is not robust.

WHY IT MATTERS IN QUANT RESEARCH

Average Sharpe can hide drawdowns concentrated in liquidity stress, trend reversals, or correlation spikes. State analysis guides limits, diversification, and fallback without assuming every regime is forecastable.

FORMULA INTUITION

Decompose total PnL by predeclared state, but weight evidence by independent episodes and include transition costs. Stress loss uses adverse state confusion and detection delay.

SIMPLE MARKET EXAMPLE

A mean-reversion strategy earns steadily in calm markets but loses six months of profits during the first three days of each volatility transition. A lagged overlay may arrive too late to help.

PYTHON / SYSTEM IMPLEMENTATION

Report state occupancy, episode count, conditional return distribution, turnover, slippage, drawdown, and transition path. Test unconditional risk reduction before complex regime switching.

CONFUSIONS TO AVOID

Optimizing a separate model in every historical state multiplies parameters and selection risk. Excluding bad regimes ex post is equivalent to using future information.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The portfolio layer applies approved state limits and diversification rules, while the model remains responsible for unconditional behavior. Incident reviews distinguish detector failure from strategy fragility.

RELATED CONCEPTS AND QUICK RECALL

Related: conditional performance, transition risk, robust sizing, state confusion. Recall task: Identify whether a regime overlay helps before, during, or after transitions and include the correct execution delay.

Regime monitoring and response

INTERMEDIATE

ONE-SENTENCE MEANING

Regime monitoring tracks data quality, state evidence, transition behavior, model calibration, and downstream exposure, then maps persistent deviations to predefined operational responses. Detection without ownership is only a dashboard.

MEMORY JOGGER

Monitor the clock, the inferred state, and the capital consequence. Each alert needs a measurable threshold, an owner, and a safe next action.

WHY IT MATTERS IN QUANT RESEARCH

A state model can fail through stale data, unseen combinations, transition explosion, probability collapse, or changed payoff mapping. Monitoring separates these mechanisms so the system responds proportionally.

FORMULA INTUITION

Health combines freshness, feature support, state occupancy, transition rate, likelihood or residual score, probability calibration, and decision sensitivity. Thresholds account for expected sampling variation.

SIMPLE MARKET EXAMPLE

A feed freeze produces low measured volatility and a calm-state label while quote age rises. A quality-first monitor blocks leverage expansion before the state model can misinterpret the defect.

PYTHON / SYSTEM IMPLEMENTATION

Implement warning, degraded, and stop tiers; retain conservative fallback parameters; snapshot state for replay; and run incident simulations. Test alerts on historical and synthetic failures.

CONFUSIONS TO AVOID

Automatic retraining on every state alert can absorb corrupt data or crisis noise. A stable label is not healthy if probabilities are unsupported or capital sensitivity is extreme.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The control plane binds model health to position caps, refit authorization, fallback, and rollback. Every transition in operational mode is logged with evidence and approver.

RELATED CONCEPTS AND QUICK RECALL

Related: observability, fallback, alert tier, model-health state. Recall task: Draft three response tiers for a regime model and include one data-quality trigger and one decision-sensitivity trigger.

Regime model governance

INTERMEDIATE

ONE-SENTENCE MEANING

Regime model governance defines ownership, versioning, evidence, permitted uses, and change control for state-dependent decisions. It prevents attractive historical labels from becoming untraceable production controls.

MEMORY JOGGER

A regime is a governed model artifact, not a colored background on a chart. Its definition, uncertainty, and authorized action must survive audit and replay.

WHY IT MATTERS IN QUANT RESEARCH

State models touch leverage, risk limits, strategy routing, and incident response, so errors can have broad consequences. Governance separates exploratory labels from models allowed to move capital.

FORMULA INTUITION

An artifact identity hashes data snapshot, feature transforms, state model, hyperparameters, fit cutoff, label alignment, and probability-to-action policy. Any dependency change creates a new version.

SIMPLE MARKET EXAMPLE

A researcher changes volatility thresholds after a recent drawdown. Governance requires a new experiment, later validation, impact review, shadow period, and explicit promotion rather than overwriting history.

PYTHON / SYSTEM IMPLEMENTATION

Maintain a model card with state semantics, episode evidence, assumptions, limitations, monitors, fallback, and retirement criteria. Reproduce representative historical decisions from immutable inputs.

CONFUSIONS TO AVOID

A sophisticated model with unclear state meaning is harder to govern than a simple threshold. Versioning code alone is insufficient when data vintages or action maps can change outcomes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The registry links data, features, fit artifacts, forecasts, decisions, incidents, and approvals. Deployment refuses unapproved combinations and supports independent rollback of state model or action policy.

RELATED CONCEPTS AND QUICK RECALL

Related: model registry, artifact lineage, shadow deployment, retirement rule. Recall task: List the minimum artifacts needed to reproduce one historical regime-driven exposure change and assign an owner to each.


One path, four stationarity views

INTERMEDIATE


READING OBJECTIVE

The synthetic path combines a deterministic slope, a level break, and a volatility shift. Each transform reveals a different failure of stability; no single plot supplies a complete diagnosis.


Level: trend plus break




First difference: location shock exposed



Detrended residual: scale still changes



Trailing scale: delayed state estimate



INTERPRETATION DISCIPLINE

Differencing removes the trend and one-time level displacement from the level path, but it does not stabilize the later variance. Detrending with the correct break still leaves a scale regime, while the rolling scale reacts only after large observations enter its window.

Choosing a stability treatment

INTERMEDIATE

USE THE TREE AS A HYPOTHESIS MAP

Every branch requires point-in-time plots, economic context, and later validation. The aim is to choose a defensible challenger set, not to turn noisy tests into an automatic transformation oracle.



FINAL GATE

Whichever path is chosen, refit transformations inside walk-forward folds, preserve legal timing, compare against simpler alternatives, and measure decision sensitivity to near-equivalent specifications.

ADF, PP, and KPSS evidence matrix

INTERMEDIATE

ADF OR PP REJECTS; KPSS DOES NOT REJECT

Evidence is consistent with level or trend stationarity under the selected deterministic terms. Continue with break checks, effect-size judgment, and later forecast validation before accepting a stable level model.

ADF OR PP DOES NOT REJECT; KPSS REJECTS

Evidence supports a unit-root or otherwise nonstationary interpretation. Compare first differences, stochastic-trend models, and equilibrium combinations; preserve uncertainty about near-unit behavior.

BOTH REJECT THEIR NULLS

The series may contain breaks, nonlinear dynamics, changing variance, or specification errors that make both simple nulls inadequate. Inspect events and run break-aware challengers rather than choosing the preferred p-value.

NEITHER REJECTS ITS NULL

The sample is inconclusive, often because persistence is near one or power is low. Use conservative forecasts, wider uncertainty, longer or pooled evidence where valid, and parallel level-and-difference challengers.

DETERMINISTIC-TERM RULE

Intercept, trend, seasonality, and break dummies change the null model. Select them from economics and data construction before reading results, then disclose sensitivity rather than silently optimizing significance.

DECISION RULE

Tests narrow a model set; they do not certify tradability. The final choice uses walk-forward loss, residual diagnostics, state stability, transaction costs, and the consequence of choosing the wrong transformation.

FINITE-SAMPLE CALIBRATION

Simulate fitted stationary, near-unit, broken, and heteroskedastic processes at the actual sample length. Empirical rejection rates reveal whether nominal p-values are credible for the research setting.

QUICK RECALL

For each cell, state the conservative transformation choice, the key alternative explanation, and the next piece of evidence that would change the decision.

Break event, evidence, and action

INTERMEDIATE

CORE DISTINCTION

A market mechanism can change before the detector has enough evidence to signal. Realistic research charges the strategy for that detection interval and never assigns the final regime label backward in time.



LOSS DURING DETECTION

The strategy remains exposed while evidence accumulates. Measure transition PnL, leverage, slippage, and alert latency explicitly; a regime overlay that helps only with hindsight does not control live risk.

BOUNDARY UNCERTAINTY

After the episode, the estimated event time may have a wide range. Re-run attribution and validation across that range, but retain the original warning, confirmation, and action timestamps for operational truth.

SYSTEM RECORD

Persist detector state, reference version, raw evidence, threshold crossings, approvals, fallback, state probability, decision, order, and fill. This ledger separates market delay, computational delay, and governance delay.

Regime stability across choices

INTERMEDIATE

QUESTION

Do state assignments persist when the lookback, threshold, and data period move within reasonable bounds? A robust regime has recognizable structure beyond one tuned configuration.

Calm-trend occupancy	0.82	0.75	0.88	0.81	0.90	0.86
Stress recall	0.91	0.78	0.86	0.74	0.69	0.60
Transition delay	0.88	0.80	0.72	0.65	0.48	0.40
Label agreement	0.68	0.62	0.81	0.76	0.72	0.69
Strategy risk reduction	0.76	0.71	0.84	0.80	0.73	0.70
Turnover control	0.42	0.58	0.71	0.79	0.88	0.92

INTERPRETATION

Fast settings recall stress quickly but trade excessively; slow settings control turnover but miss the start of stress. Medium settings are not automatically optimal: choose by the specified loss function, then stress near-equivalent settings and later episodes.

DENSITY RULE

Every cell must be computed only from walk-forward state estimates. If a row uses retrospective labels or the final threshold, apparent stability is a property of hindsight rather than the deployable process.

Method selection by failure mode

INTERMEDIATE

OBSERVED PROBLEM	FIRST CHALLENGER	KEY ASSUMPTION	PRIMARY RISK
Known linear trend	Trend regression	Stable slope	Extrapolation after break
Accumulated shocks	First difference	I(1)-like path	Over-differencing noise
Stable level combination	Cointegration	Persistent equilibrium	Hedge-ratio break
Known event date	Segmented test	Exogenous boundary	Variance confounding
Unknown abrupt break	Changepoint model	Adequate segment length	Search false alarms
Slow parameter drift	Local model	Smooth evolution	Window lag and noise
Recurring hidden states	Markov or HMM	Repeated episodes	Label and model risk
Observed nonlinear state	Threshold model	Causal transition variable	Boundary chatter

ESCALATION RULE

Prefer the simplest challenger that represents the observed failure and supports the decision. Add state complexity only after plots, event context, residuals, and walk-forward loss show that a simpler transform or robust limit is inadequate.

Failure modes and diagnostic actions

INTERMEDIATE

SYMPTOM	LIKELY CAUSE	CHECK	ACTION
High level R-squared	Shared stochastic trend	Differences and residual tests	Change target or test equilibrium
ADF and KPSS conflict	Break or low power	Event plot and simulations	Keep competing specifications
Rolling mean drifts smoothly	Overlap masks abrupt break	Nonoverlap blocks; event dates	Estimate boundary range
Many rapid state switches	Threshold chatter or noise	Probability path and dwell time	Add hysteresis or simplify
Perfect historical state timing	Smoothed-label leakage	Replay filtered probabilities	Charge detection delay
Calm state during feed freeze	Data defect	Quote age and event counts	Block leverage; quarantine
Regime Sharpe from one crisis	Episode dependence	Episode-level resampling	Limit claim and exposure
Live state differs from batch	State or cutoff skew	Batch-stream replay	Unify artifact and clock

TRIAGE ORDER

Verify identity, clock, calendar, revision, gaps, and staleness before changing a model. Then inspect transformation and deterministic terms, then break or state structure, and only afterward increase model complexity.

Formula sheet and interpretation map

INTERMEDIATE

WEAK STATIONARITY

$E[X_t] = \mu$ and $\text{Cov}(X_t, X_{t-k}) = \gamma_k$. The center and lag relation are invariant, but tails and higher moments may still change.

FIRST DIFFERENCE

$\Delta X_t = (1-B)X_t = X_t - X_{t-1}$. One difference removes one common stochastic trend but changes the economic target.

AR PERSISTENCE

$X_t = c + \phi X_{t-1} + \epsilon_t$. A unit root has $\phi = 1$; for $|\phi| < 1$, half-life is $\log(0.5)/\log(|\phi|)$.

TREND-STATIONARY MODEL

$X_t = a + b t + u_t$ with stationary u_t . Shocks to u_t decay while the deterministic path continues.

COINTEGRATION

$z_t = Y_t - \beta X_t$ is $I(0)$ although X_t and Y_t are $I(1)$. Estimate β inside the training window and test later residual behavior.

CUSUM

$C_t^+ = \max(0, C_{t-1}^+ + z_{t-k})$; alert when $C_t^+ > h$. Reference k and threshold h define detectable shift and false-alarm behavior.

MARKOV TRANSITION

$P(S_t = j | S_{t-1} = i) = p_{ij}$. Filtered state probabilities use only current and past evidence; smoothed probabilities use future observations.

MIXTURE DECISION

$\text{Forecast} = \sum_j p_{t,j} f_j(X_t)$. Bound and calibrate the probability-to-action map so wrong-state decisions remain survivable.

Worked example: break to legal action

INTERMEDIATE

1. INSPECT THE PATH

A spread has monthly levels 2.0, 2.3, 2.5, 2.8, 5.9, 6.1, 6.0, and 6.2. The jump near the midpoint suggests a level break rather than a single stable center.

2. COMPARE SEGMENT MEANS

The first four observations average 2.40; the last four average 6.05. The 3.65-unit difference is large relative to within-segment movement, so one full-sample mean of 4.23 is not representative of either state.

3. DIFFERENCE THE LEVEL

Changes are +0.3, +0.2, +0.3, +3.1, +0.2, -0.1, and +0.2. Differencing localizes the break as one extreme innovation but does not explain whether it is permanent.

4. PRESERVE TIMING

Suppose the detector needs two post-break observations to confirm. The event occurs at observation five, warning arrives at six, confirmation at seven, and the earliest new-state trade occurs afterward.

5. CHARGE TRANSITION LOSS

Any position maintained from observations five through seven belongs to the old live policy even if retrospective segmentation labels those dates new regime. Attribute its PnL to detection delay.

6. COMPARE MODELS

Evaluate a stable-mean model, first-difference model, and two-level break model at rolling origins. Fit centers and detector thresholds only with history available at each origin.

7. STRESS THE BOUNDARY

Move the candidate event between observations four and six and repeat the conclusion. A strategy whose profitability requires the exact best boundary has weak regime evidence.

8. REGISTER THE DECISION

Store raw snapshot, observation and availability times, transform, detector state, threshold, event estimate, confirmation time, model version, exposure action, and earliest legal fill.

Final active recall and system index

INTERMEDIATE

INVARIANT CONTRACT

State exactly which mean, scale, dependence, tail, or payoff property must remain stable, for which horizon, grid, universe, and consuming decision.

TREND MECHANISM

Separate deterministic trend, stochastic trend, seasonality, equilibrium combinations, fractional persistence, local drift, and abrupt breaks before choosing a transform.

EVIDENCE PANEL

Use plots, event context, opposing null tests, residuals, break models, boundary ranges, simulations, and later forecast loss; no single p-value certifies stationarity.

REGIME DESIGN

Define observables or emissions, persistence, state meaning, probability, detection latency, minimum episode evidence, allowed action, uncertainty, and fallback.

CAUSAL TIMING

Require event $\leq$ availability $\leq$ feature cutoff $\leq$ decision $\leq$ order $\leq$ fill. Use filtered or predicted state probabilities, never retrospective smoothed labels for trading.

VALIDATION

Walk forward the entire preprocessing and state process, purge overlapping labels, count independent episodes, charge transition costs, and stress wrong-state decisions.

OPERATIONAL CONTROLS

Monitor freshness, support, state occupancy, transition rate, probability calibration, residuals, decision sensitivity, alert tiers, fallback, rollback, and retirement.

SYSTEM LINEAGE

Version snapshots, calendars, transforms, fitted parameters, state alignment, detector state, probabilities, actions, forecasts, fills, metrics, incidents, and approvals.

FINAL RECALL

Explain why a process can be mean-stable but risk-unstable, why a break can look like long memory, and why a perfect retrospective regime chart can be unusable live.