

Volatility Modeling Foundations

Volatility targets, conditional variance, realized variation, forecast validation, and capital controls.

Library of the Quant

Educational reference manual

Manual Profile

LIBRARY ID	LOTQ-0030
CATEGORY	Time Series
DIFFICULTY	Intermediate
FORMAT	Single-column field-guide manual
USE	Study, lookup, and research-system reference

Educational Use Only

This manual is educational material for quant research study. It is not financial advice, not a trading recommendation, and not a guarantee of correctness or live trading usefulness.

Table of contents

INTERMEDIATE

ORIENTATION

Title	1
How to use this manual	5
Notation and volatility contract	6
Volatility research and control pipeline	7

VOLATILITY OBJECTS AND MEASUREMENT

Volatility is scale, not direction	8
Price levels, returns, and scale targets	9
Variance, standard deviation, and units	10
Squared and absolute returns as proxies	11
Volatility clustering	12
Conditional and unconditional volatility	13
Asymmetry and the leverage effect	14
Heavy tails, jumps, and data errors	15

Table of contents

INTERMEDIATE

CONDITIONAL VARIANCE MODELS

Rolling-window volatility	16
Exponentially weighted variance	17
ARCH models	18
GARCH(1,1) recursion	19
Persistence and volatility half-life	20
Long-run variance and variance targeting	21
Multi-step GARCH forecasts	22
GJR-GARCH and threshold response	23
EGARCH and log variance	24
Innovation distributions and Student-t tails	25

REALIZED AND MARKET-BASED VOLATILITY

Realized variance from intraday returns	26
Sampling frequency and market noise	27
Range-based volatility estimators	28
Jumps and bipower variation	29
HAR-RV and heterogeneous horizons	30
Intraday seasonality	31
Implied versus expected realized volatility	32

Table of contents

INTERMEDIATE

VALIDATION, RISK, AND SYSTEM DESIGN

Forecast target and horizon alignment	33
Loss functions for variance forecasts	34
Walk-forward volatility validation	35
Calibration and standardized residual diagnostics	36
Structural breaks and volatility regimes	37
Parameter and model uncertainty	38
Volatility forecasts in sizing and limits	39
Production monitoring and artifact lineage	40

VISUAL LABS AND REVIEW

Returns, clustering, and conditional scale	41
Choosing a volatility model family	42
Persistence, half-life, and horizon risk	43
From market event to scored forecast	44
Variance forecast path after a shock	45
Method selection by measurement problem	46
Failure modes and diagnostic actions	47
Formula sheet and interpretation map	48
Worked calculation: shock to governed risk	49
Validation scorecard for a volatility model	50
Final active recall and system index	51

How to use this manual

INTERMEDIATE

1. DEFINE THE RISK OBJECT

Name price source, return convention, horizon, session, target type, and earliest legal forecast origin before selecting a model.

2. START WITH TRANSPARENT BASELINES

Compare rolling variance and EWMA before adding likelihood, asymmetry, tails, or regimes. Complexity must earn later decision value.

3. SEPARATE MEASUREMENT FROM FORECASTING

Squared returns, realized variance, ranges, and implied variance are different measurements. A model forecasts one declared target.

4. CARRY THE FULL FORECAST PATH

After a shock, risk changes across horizons. Store stepwise variance, cumulative variance, long-run anchor, and uncertainty.

5. VALIDATE DISTRIBUTION, NOT ONE NUMBER

Inspect loss, standardized residuals, coverage, tails, regimes, and capital consequences on walk-forward forecasts.

6. CONVERT FORECASTS INTO BOUNDED ACTION

Use scale floors, caps, turnover limits, liquidity constraints, health discounts, and stress overrides.

DECISION STANDARD

Promote only when target timing, data quality, model diagnostics, later calibration, model risk, fallback, and lineage are reproducible.

Notation and volatility contract

INTERMEDIATE

R_T AND EPSILON_T

r_t is the return; ϵ_t is the mean-model innovation. State whether scale models raw, demeaned, or standardized values.

H_T AND SIGMA_T

h_t is conditional variance in squared-return units; $\sigma_t = \sqrt{h_t}$ is conditional volatility in return units.

OMEGA, ALPHA, BETA

GARCH baseline, shock reaction, and variance carryover. Total symmetric persistence is commonly $\alpha + \beta$.

Z_T

Standardized innovation ϵ_t / σ_t . Calibration expects unit conditional variance and the declared tail distribution.

RV_T AND BV_T

Realized variance sums intraday squares; bipower variation estimates continuous variation under stated high-frequency conditions.

LAMBDA AND P

Lambda controls EWMA decay; p is model persistence. Both require a declared observation interval.

F_T AND FORECAST ORIGIN

F_t is the legal information set. Record event time, availability, computation, decision, order, and fill.

MINIMUM ARTIFACT CONTRACT

Universe, prices, calendar, returns, target, horizon, units, cutoff, model equation, parameters, loss, costs, controls, owner, and fallback.

Volatility research and control pipeline

INTERMEDIATE

1. FRAME DECISION

Specify risk action, horizon, execution clock, loss asymmetry, and required distribution output.

2. BUILD LEGAL RETURNS

Align prices, adjustments, sessions, gaps, quality flags, and availability timestamps.

3. CHOOSE MEASUREMENT

Select squared return, realized variance, range, jump component, or implied measure with matching units.

4. ESTABLISH BASELINES

Run historical, rolling, EWMA, and seasonal benchmarks before expanding the model family.

5. FIT DYNAMICS

Estimate symmetric scale, asymmetry, tails, horizons, and long-run anchor inside each origin.

6. DIAGNOSE RESIDUALS

Test z and z squared, quantiles, exceedance clusters, response curves, and parameter boundaries.

7. WALK FORWARD

Issue frozen forecasts, wait for legal targets, purge overlap, and score state-conditioned loss.

8. MAP TO CAPITAL

Apply confidence buffers, floors, caps, covariance, costs, liquidity, stress, and health.

9. GOVERN PRODUCTION

Monitor clocks, data, calibration, sensitivity, fallback, incidents, and complete decision lineage.

Volatility is scale, not direction

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility describes the conditional size and dispersion of future price changes, not whether the next return will be positive or negative. A volatility forecast can be excellent while having no directional information.

MEMORY JOGGER

Direction asks which way; volatility asks how far. Keep the mean forecast and the scale forecast as separate artifacts even when one model produces both.

WHY IT MATTERS IN QUANT RESEARCH

Risk limits, option values, execution tactics, leverage, and forecast intervals all depend on scale. Treating volatility as a buy-or-sell signal confuses uncertainty with expected return.

FORMULA INTUITION

A common return model is $r_t = \mu_t + \sigma_t z_t$, where $E[z_t | F_{t-1}] = 0$ and $\text{Var}(z_t | F_{t-1}) = 1$. σ_t controls conditional scale.

SIMPLE MARKET EXAMPLE

An earnings announcement may leave expected return near zero but raise the predicted daily standard deviation from 1% to 4%. The trade implication is less leverage and wider scenarios, not an automatic long or short.

PYTHON / SYSTEM IMPLEMENTATION

Publish μ , σ , horizon, units, cutoff, and earliest action separately. Test that downstream services never infer a sign from a positive volatility forecast.

CONFUSIONS TO AVOID

High volatility does not imply falling prices, and low volatility does not imply safety. Quiet markets can accumulate leverage and gap sharply when the state changes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The forecast layer emits a conditional distribution; portfolio and execution services consume its scale. Directional alpha remains an independently validated input.

RELATED CONCEPTS AND QUICK RECALL

Related: conditional distribution, return forecast, risk forecast, standardized innovation. Recall task: Describe a market event that changes volatility without changing expected direction. Name two systems that should react.

Price levels, returns, and scale targets

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility is usually modeled on returns or price changes because raw price levels carry units, trends, and splits that obscure comparable variation. The target must match the economic decision and sampling grid.

MEMORY JOGGER

First transform the price into the change the strategy actually bears. Then define whether the model forecasts variance, standard deviation, absolute movement, or a quantile.

WHY IT MATTERS IN QUANT RESEARCH

Incorrect targets contaminate annualization, model comparison, and risk budgets. A one-minute midquote return, a daily close return, and a ten-day overlapping return are different statistical objects.

FORMULA INTUITION

Log return $r_t = \log(P_t/P_{(t-1)})$; simple return $R_t = P_t/P_{(t-1)} - 1$. For small moves they are close, but aggregation and extreme losses differ.

SIMPLE MARKET EXAMPLE

A stock split creates a huge raw price change in an unadjusted series but no investor loss. A corrected return series removes that false volatility shock while retaining genuine gap risk.

PYTHON / SYSTEM IMPLEMENTATION

Version price source, adjustment policy, timezone, calendar, sampling, gap rule, and return convention. Reconcile computed returns against a small hand-audited sample.

CONFUSIONS TO AVOID

Volatility of price levels is not a stable substitute for return volatility. Forward-filled prices manufacture zeros followed by delayed jumps and distort both clustering and tails.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The data contract produces point-in-time returns with quality flags. Every model registry entry names the exact target field and observation clock.

RELATED CONCEPTS AND QUICK RECALL

Related: simple return, log return, corporate action, sampling grid. Recall task: Explain why an unadjusted split corrupts volatility. State which metadata makes the corrected target reproducible.

Variance, standard deviation, and units

INTERMEDIATE

ONE-SENTENCE MEANING

Variance measures average squared dispersion, while standard deviation is its square root and returns to the units of the modeled change. Models often evolve variance because squared shocks add naturally.

MEMORY JOGGER

Variance is the engine; standard deviation is the human-readable speedometer. Never compare or annualize them with the same exponent.

WHY IT MATTERS IN QUANT RESEARCH

Portfolio variance, likelihood functions, and GARCH recursions operate in squared-return units. Position sizing and scenario communication often use standard deviation or percent moves.

FORMULA INTUITION

If conditional variance is h_t , conditional volatility is $\sigma_t = \sqrt{h_t}$. Under independent equal-variance periods, $h_H = H \cdot h_1$ and $\sigma_H = \sqrt{H} \cdot \sigma_1$.

SIMPLE MARKET EXAMPLE

Daily variance 0.0004 corresponds to 2% daily volatility. Multiplying the standard deviation by 252 instead of $\sqrt{252}$ would report an impossible 504% annualized scale.

PYTHON / SYSTEM IMPLEMENTATION

Attach type and unit metadata to every field: `variance_decimal2`, `volatility_decimal`, or `volatility_percent`. Add unit tests for squaring, square roots, and annualization.

CONFUSIONS TO AVOID

Variance forecasts are nonnegative by definition, but unconstrained regressions can output negatives. Percentage and decimal conventions can differ by factors of 100 or 10,000.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The semantic layer exposes both h_t and σ_t from one artifact ID so risk, reporting, and optimization cannot silently combine incompatible scales.

RELATED CONCEPTS AND QUICK RECALL

Related: second moment, conditional variance, annualization, dimensional analysis. Recall task: Convert daily variance 0.0004 into daily volatility and an IID annualized volatility. Identify the assumptions.

Squared and absolute returns as proxies

INTERMEDIATE

ONE-SENTENCE MEANING

The latent conditional variance is not directly observed, so squared or absolute returns are common noisy proxies for realized scale. Each proxy emphasizes large moves differently.

MEMORY JOGGER

One daily return is a very noisy volatility measurement. The model must separate a persistent scale state from the randomness of a single shock.

WHY IT MATTERS IN QUANT RESEARCH

Proxy choice affects estimation loss, diagnostics, and apparent forecast quality. Absolute returns are less dominated by extremes, while squared returns align directly with variance equations.

FORMULA INTUITION

If $r_t = \sigma_t z_t$ with $E[z_t^2] = 1$, then $E[r_t^2 | F_{t-1}] = \sigma_t^2$ when the conditional mean is negligible. Absolute returns require a distribution-specific scale factor.

SIMPLE MARKET EXAMPLE

A 5% return contributes 0.0025 to squared-return loss, twenty-five times the contribution of a 1% return. In absolute-return space it contributes only five times as much.

PYTHON / SYSTEM IMPLEMENTATION

Retain raw return, demeaned residual, square, absolute value, and robust flags. Evaluate conclusions with realized measures when reliable intraday data exist.

CONFUSIONS TO AVOID

Squaring converts bad ticks into dominant observations. Replacing all large moves as outliers can erase the tail events a risk model is supposed to learn.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature layer publishes multiple scale proxies with provenance. The evaluation service maps each forecast to a target and loss designed for that target.

RELATED CONCEPTS AND QUICK RECALL

Related: latent variance, realized variance, absolute deviation, robust loss. Recall task: Compare the influence of a 5% and 1% return under squared and absolute proxies. Which proxy is more tail-sensitive?

Volatility clustering

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility clustering is the tendency for large moves to be followed by large moves and quiet periods by quiet periods, even when return signs remain hard to predict. Dependence appears in magnitude rather than direction.

MEMORY JOGGER

Returns may look uncorrelated while their squares remember the recent past. Diagnose both the mean and the scale channels.

WHY IT MATTERS IN QUANT RESEARCH

Clustering makes constant-volatility risk limits stale after shocks and too conservative after long calm periods. It motivates conditional variance models and horizon-specific forecasts.

FORMULA INTUITION

Typical evidence is $\text{Corr}(r_t, r_{t-k})$ near zero but $\text{Corr}(r_t^2, r_{t-k}^2) > 0$ for several lags. A model represents this with a persistent h_t process.

SIMPLE MARKET EXAMPLE

After a 4% index move, the next day's direction may be unclear, but a 2% move is more plausible than during an ordinary 0.7% regime.

PYTHON / SYSTEM IMPLEMENTATION

Plot returns, absolute returns, squared returns, rolling scale, and ACFs on the same legal sample. Mark events and missing-data repairs before telling an economic story.

CONFUSIONS TO AVOID

Clusters can be exaggerated by stale prices, changing market hours, or mixed intraday seasonality. A moving window also creates persistence mechanically.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Exploration establishes the fingerprint; forecasting estimates tomorrow's scale; monitoring checks whether realized shock frequency matches the issued distribution.

RELATED CONCEPTS AND QUICK RECALL

Related: ARCH effect, magnitude dependence, conditional heteroskedasticity, persistence. Recall task: Explain how returns can be serially uncorrelated while squared returns are predictable. What model component changes?

Conditional and unconditional volatility

INTERMEDIATE

ONE-SENTENCE MEANING

Conditional volatility is the scale forecast given current information, while unconditional volatility averages across the long-run distribution under a stable model. They answer different questions.

MEMORY JOGGER

Conditional is today's weather forecast; unconditional is the climate. A risk engine needs the former, while calibration and stress context use both.

WHY IT MATTERS IN QUANT RESEARCH

Confusing the two can freeze leverage during calm periods or ignore recent shocks. Long-run variance also anchors mean reversion in persistent models.

FORMULA INTUITION

For stationary GARCH(1,1), $h_{\text{bar}} = \omega / (1 - \alpha - \beta)$. The conditional h_t moves around h_{bar} as shocks arrive, provided $\alpha + \beta < 1$.

SIMPLE MARKET EXAMPLE

A model may have 15% long-run annualized volatility but forecast 32% tomorrow after a shock. Using 15% for overnight risk understates immediate uncertainty.

PYTHON / SYSTEM IMPLEMENTATION

Store conditional forecasts at each horizon plus the model-implied long-run level. Compare rolling realized outcomes with both and flag persistent divergence.

CONFUSIONS TO AVOID

An unconditional value may not exist or be operationally meaningful when persistence is near one or regimes change. A sample standard deviation is not automatically a stable long-run parameter.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The forecast API returns current path, long-run anchor, convergence rate, and uncertainty. Capital limits can use the larger of current and stressed scale.

RELATED CONCEPTS AND QUICK RECALL

Related: ergodic variance, long-run mean, forecast path, variance reversion. Recall task: Give a case where conditional volatility is twice the unconditional level. Which value should size tomorrow's trade?

Asymmetry and the leverage effect

INTERMEDIATE

ONE-SENTENCE MEANING

Negative returns often raise future equity volatility more than positive returns of the same magnitude. This asymmetry can reflect financial leverage, risk aversion, feedback trading, and market mechanics.

MEMORY JOGGER

Equal-sized shocks need not have equal scale consequences. Test sign and magnitude jointly rather than assuming a symmetric square is sufficient.

WHY IT MATTERS IN QUANT RESEARCH

Asymmetric forecasts improve downside scenarios, option risk, and post-shock leverage controls. They are especially relevant when losses and liquidity stress arrive together.

FORMULA INTUITION

A GJR term adds $\gamma I(\epsilon_{t-1} < 0) \epsilon_{t-1}^2$ to h_t . Positive γ gives negative shocks an extra variance impact.

SIMPLE MARKET EXAMPLE

A -3% index return may increase the next variance forecast far more than a +3% return. A symmetric GARCH model assigns them identical immediate squared-shock effects.

PYTHON / SYSTEM IMPLEMENTATION

Estimate sign effects with event and regime controls, then validate later forecasts separately after positive and negative shocks. Preserve the residual sign convention.

CONFUSIONS TO AVOID

Calling every asymmetry leverage is too strong; volatility feedback and data construction may contribute. Inverting a return series also changes the sign interpretation.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model registry records shock sign, indicator rule, parameter restrictions, and side-specific diagnostics. Risk reports show forecast response curves by shock direction.

RELATED CONCEPTS AND QUICK RECALL

Related: GJR-GARCH, EGARCH, downside risk, news impact curve. Recall task: Compare the one-step variance response to +3% and -3% shocks in a symmetric and asymmetric model.

Heavy tails, jumps, and data errors

INTERMEDIATE

ONE-SENTENCE MEANING

Financial returns contain more extreme observations than a Gaussian model predicts, but an extreme can be genuine market information, a discrete jump, or a data defect. The model and pipeline must distinguish these cases.

MEMORY JOGGER

Do not clean away risk and do not train on bad ticks. Preserve the raw observation, the repair decision, and an event-aware alternative.

WHY IT MATTERS IN QUANT RESEARCH

Tail assumptions determine forecast intervals, likelihood influence, VaR, and stress losses. One corrupted tick can also dominate squared-return estimation for months.

FORMULA INTUITION

Excess kurtosis above zero indicates heavier tails than Gaussian under a finite fourth moment. Student-t innovations add a degrees-of-freedom parameter controlling tail thickness.

SIMPLE MARKET EXAMPLE

A genuine central-bank gap and an isolated crossed quote can both produce a six-sigma return. Only the former belongs in economic tail evidence; both belong in the audit log.

PYTHON / SYSTEM IMPLEMENTATION

Use cross-venue checks, quote conditions, neighboring observations, event calendars, and immutable raw data. Run models on approved raw, repaired, and robust variants.

CONFUSIONS TO AVOID

Winsorization can improve numerical stability while understating capital risk. A heavy-tailed likelihood does not excuse unresolved bad data.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The quality service classifies observations; the model consumes flags; the stress library retains verified jumps even when robust estimation downweights them.

RELATED CONCEPTS AND QUICK RECALL

Related: kurtosis, jump process, Student-t, robust estimation. Recall task: Design a rule that separates a market-wide gap from one bad tick. Which version belongs in stress testing?

Rolling-window volatility

INTERMEDIATE

ONE-SENTENCE MEANING

Rolling volatility estimates scale from a fixed recent window and discards older observations abruptly. It is transparent but produces a stepwise memory rule chosen by the researcher.

MEMORY JOGGER

Every observation has equal weight until it falls off a cliff. Window length controls responsiveness, noise, and mechanical overlap.

WHY IT MATTERS IN QUANT RESEARCH

Rolling estimates are strong baselines for dashboards and risk limits. Their simplicity makes it easier to detect when a complex model adds no decision value.

FORMULA INTUITION

$s_t^2 = \sum_{i=1}^L (r_{t-i} - \bar{r}_t)^2 / (L-1)$, using only returns available before t . Forecast convention must state whether the current return is included.

SIMPLE MARKET EXAMPLE

A 20-day window reacts quickly to a shock but remains elevated until that shock exits on one date. A 252-day window adapts slowly and mixes older regimes.

PYTHON / SYSTEM IMPLEMENTATION

Version L, demeaning, minimum history, missing policy, annualization, and cutoff. Compare windows at identical origins and include a nonoverlapping diagnostic.

CONFUSIONS TO AVOID

Overlapping windows make adjacent estimates appear smoother and more stable than independent evidence warrants. Optimizing L on final PnL is data snooping.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The baseline service publishes rolling estimates and their observation coverage. More complex forecasts must beat them on approved losses and decisions.

RELATED CONCEPTS AND QUICK RECALL

Related: moving window, sample variance, lookback, effective sample size. Recall task: Contrast a 20-day and 252-day estimator after one shock. Explain the step that occurs when the shock leaves.

Exponentially weighted variance

INTERMEDIATE

ONE-SENTENCE MEANING

EWMA updates variance by combining yesterday's forecast with yesterday's squared shock, giving recent observations more weight without a hard window boundary. One decay parameter controls memory.

MEMORY JOGGER

Old information fades geometrically. A larger lambda is smoother and slower; a smaller lambda is faster and noisier.

WHY IT MATTERS IN QUANT RESEARCH

EWMA is a practical baseline for daily risk and covariance updates. It adapts faster than long rolling windows while remaining easy to reproduce.

FORMULA INTUITION

$h_t = \lambda h_{t-1} + (1-\lambda)\epsilon_{t-1}^2$, usually with $0 < \lambda < 1$. Weight at age k is $(1-\lambda)\lambda^{k-1}$.

SIMPLE MARKET EXAMPLE

With $\lambda = 0.94$, yesterday's squared shock receives 6% weight and the prior variance 94%. A large shock lifts forecasts and then decays gradually.

PYTHON / SYSTEM IMPLEMENTATION

Define initialization, mean treatment, missing updates, calendar gaps, lambda, and warm-up. Reconcile recursive output against direct geometric weights.

CONFUSIONS TO AVOID

Standard EWMA has no separate long-run anchor, so a shock decays toward zero absent new shocks. Lambda chosen by convention is not automatically optimal for every asset or horizon.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The risk service uses EWMA as an interpretable fallback. Monitoring checks update freshness and sensitivity to lambda and initialization.

RELATED CONCEPTS AND QUICK RECALL

Related: decay factor, recursive filter, RiskMetrics, geometric weights. Recall task: For $\lambda = 0.94$, identify yesterday's shock weight and explain how a larger lambda changes half-life.

ARCH models

INTERMEDIATE

ONE-SENTENCE MEANING

An ARCH model expresses current conditional variance as a positive baseline plus weighted past squared shocks. It formalizes the idea that recent surprises predict near-term scale.

MEMORY JOGGER

Variance reacts to a finite history of news. The model explains magnitude dependence even when mean returns show little autocorrelation.

WHY IT MATTERS IN QUANT RESEARCH

ARCH provides a likelihood-based foundation for conditional risk, standardized residual diagnostics, and later GARCH extensions.

FORMULA INTUITION

ARCH(q): $h_t = \omega + \sum_{i=1}^q \alpha_i \epsilon_{t-i}^2$, with $\omega > 0$ and $\alpha_i \geq 0$ for a simple nonnegative specification.

SIMPLE MARKET EXAMPLE

An ARCH(2) can give yesterday's shock twice the weight of the prior day's shock. Forecasts fall back rapidly once both observations leave the lag set.

PYTHON / SYSTEM IMPLEMENTATION

Fit the mean and variance equations together or document the sequence. Inspect standardized residuals, squared residual ACF, constraints, and optimizer convergence.

CONFUSIONS TO AVOID

Large q creates many correlated parameters and abrupt memory. Significant ARCH effects do not prove that the selected order forecasts better out of sample.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The experiment tracker stores mean model, variance order, likelihood, constraints, initialization, diagnostics, and fit cutoff as one artifact.

RELATED CONCEPTS AND QUICK RECALL

Related: conditional heteroskedasticity, LM test, squared residual, likelihood. Recall task: Explain why ARCH can model volatility clustering with uncorrelated raw returns. What residual diagnostic should become flat?

GARCH(1,1) recursion

INTERMEDIATE

ONE-SENTENCE MEANING

GARCH(1,1) combines the last squared shock with the last variance forecast, producing a parsimonious persistent conditional variance process. It is the standard baseline for daily volatility dynamics.

MEMORY JOGGER

Alpha measures reaction to new shock magnitude; beta carries the existing variance state forward. Omega anchors the process.

WHY IT MATTERS IN QUANT RESEARCH

The model supplies one-step and multi-step variance forecasts, standardized residuals, and an interpretable persistence measure with only a few parameters.

FORMULA INTUITION

$h_t = \omega + \alpha \epsilon_{t-1}^2 + \beta h_{t-1}$. A common covariance-stationary condition is $\alpha + \beta < 1$ with nonnegative coefficients.

SIMPLE MARKET EXAMPLE

If $\alpha = 0.08$ and $\beta = 0.90$, a shock moves variance immediately while 90% of yesterday's forecast persists. The sum 0.98 implies slow expected decay.

PYTHON / SYSTEM IMPLEMENTATION

Scale returns sensibly, constrain parameters, use multiple starting values, and verify the likelihood optimum. Save filtered h_t and standardized residuals.

CONFUSIONS TO AVOID

High beta alone is not total persistence; $\alpha + \beta$ governs expected variance decay. A converged optimizer can still sit on a boundary or local optimum.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model registry publishes coefficients, covariance, long-run variance, forecast path, diagnostics, and fallback. Consumers never rebuild recursion ad hoc.

RELATED CONCEPTS AND QUICK RECALL

Related: ARCH, variance state, maximum likelihood, standardized residual. Recall task: Interpret $\alpha = 0.08$, $\beta = 0.90$, and their sum. Which parameter carries yesterday's forecast?

Persistence and volatility half-life

INTERMEDIATE

ONE-SENTENCE MEANING

GARCH persistence measures how slowly expected variance deviations return toward the long-run level. Near-unit persistence creates long forecast memory and large horizon uncertainty.

MEMORY JOGGER

Use $\alpha + \beta$ for symmetric GARCH variance persistence, then translate it into periods. Small coefficient changes near one create large duration changes.

WHY IT MATTERS IN QUANT RESEARCH

Persistence controls risk normalization, stress decay, forecast convergence, and how long a shock affects capital. It is an operational parameter, not just a fit statistic.

FORMULA INTUITION

Let $p = \alpha + \beta$. Expected variance deviation decays approximately as p^h , with half-life $\log(0.5)/\log(p)$ when $0 < p < 1$.

SIMPLE MARKET EXAMPLE

Persistence 0.95 implies about 13.5 periods; 0.99 implies about 69. A confidence interval spanning both values describes radically different capital duration.

PYTHON / SYSTEM IMPLEMENTATION

Propagate parameter draws into half-life and forecast paths. Report asymmetric duration intervals and stress any upper bound close to one.

CONFUSIONS TO AVOID

Clipping an estimated sum below one hides uncertainty. Structural breaks can masquerade as highly persistent but stationary volatility.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Capital policy uses a conservative persistence quantile and caps exposure when the upper duration violates limits. Monitoring tracks realized decay after shocks.

RELATED CONCEPTS AND QUICK RECALL

Related: integrated GARCH, local-to-unity, shock decay, forecast convergence. Recall task: Compute half-life for persistence 0.95 and 0.99. Why is the upper confidence bound a capital concern?

Long-run variance and variance targeting

INTERMEDIATE

ONE-SENTENCE MEANING

Long-run variance is the model-implied equilibrium level toward which conditional variance forecasts converge. Variance targeting estimates that anchor separately and expresses omega through persistence.

MEMORY JOGGER

Separate the climate from the current storm. Targeting reduces one free parameter but transfers uncertainty into the chosen long-run estimate.

WHY IT MATTERS IN QUANT RESEARCH

A stable anchor improves multi-horizon forecasts and parameter estimation. It also exposes whether a model implies an implausible long-run risk level.

FORMULA INTUITION

For GARCH(1,1), $h_bar = \omega / (1 - \alpha - \beta)$, so variance targeting sets $\omega = (1 - \alpha - \beta)h_bar$.

SIMPLE MARKET EXAMPLE

If annualized long-run volatility is 20%, convert it to the model's daily variance before setting omega. A unit mistake can shift the entire forecast path.

PYTHON / SYSTEM IMPLEMENTATION

Estimate h_bar on a declared, point-in-time sample and propagate its uncertainty. Compare fixed, rolling, and regime-aware anchors in later forecasts.

CONFUSIONS TO AVOID

Targeting a full-history sample leaks future regimes into early forecasts. When $\alpha + \beta$ is nearly one, tiny anchor or coefficient errors materially change omega.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model artifact binds anchor source, units, cutoff, persistence, and implied omega. Governance treats an anchor revision as a model change.

RELATED CONCEPTS AND QUICK RECALL

Related: unconditional variance, shrinkage, parameter reduction, long-run anchor. Recall task: Derive omega from h_bar and persistence. Explain why a full-sample target is illegal in a historical backtest.

Multi-step GARCH forecasts

INTERMEDIATE

ONE-SENTENCE MEANING

A multi-step variance forecast iterates the conditional model into the future, gradually moving from current risk toward the long-run level. Forecast variance is not obtained by merely multiplying today's one-step value.

MEMORY JOGGER

Today's storm fades according to persistence. Horizon variance aggregates the full path, not only the endpoint.

WHY IT MATTERS IN QUANT RESEARCH

Holding-period VaR, options, leverage, and scenario bands require horizon-consistent forecasts. A ten-day decision needs ten conditional variance steps and any covariance assumptions.

FORMULA INTUITION

For GARCH(1,1), $E_t[h_{t+k}] = h_{\text{bar}} + p^{k-1}(h_t - h_{\text{bar}})$, where $p = \alpha + \beta$. Cumulative variance sums forecast steps under zero conditional covariance.

SIMPLE MARKET EXAMPLE

After a shock, day-one variance may be four times normal while day-ten is closer to normal. Ten-day risk lies between using the elevated value for all days and using the long-run value immediately.

PYTHON / SYSTEM IMPLEMENTATION

Generate and store the entire forecast term structure. Reconcile one-step recursion, long-horizon convergence, and cumulative variance with simulation.

CONFUSIONS TO AVOID

Square-root-of-time scaling of today's volatility assumes constant independent variance. Summing variances also fails when return autocovariance or parameter uncertainty matters.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The forecast service returns stepwise and cumulative paths by horizon. Risk consumers specify whether they need endpoint scale, cumulative return variance, or scenario quantiles.

RELATED CONCEPTS AND QUICK RECALL

Related: term structure, cumulative variance, horizon scaling, forecast recursion. Recall task: Explain why ten-day risk is not ten times day-one variance after a shock. Which path must be summed?

GJR-GARCH and threshold response

INTERMEDIATE

ONE-SENTENCE MEANING

GJR-GARCH adds a sign-dependent shock term so negative and positive residuals of equal magnitude can have different variance effects. It is a direct extension of symmetric GARCH.

MEMORY JOGGER

The indicator opens an extra variance channel after shocks on the designated side. Always preserve the residual sign convention.

WHY IT MATTERS IN QUANT RESEARCH

Threshold response improves downside risk forecasts when bad news has stronger scale impact. It also creates an interpretable news-impact curve.

FORMULA INTUITION

$h_t = \omega + \alpha \epsilon_t^2 + \gamma I(\epsilon_t < 0) \epsilon_t^2 + \beta h_{t-1}$. Under symmetric innovations, a common persistence expression is $\alpha + \beta + \gamma/2$.

SIMPLE MARKET EXAMPLE

With $\alpha=0.05$ and $\gamma=0.10$, a negative shock receives 0.15 immediate squared-shock weight versus 0.05 for a positive shock.

PYTHON / SYSTEM IMPLEMENTATION

Fit with constraints compatible with nonnegative variance, inspect side-specific sample counts, and validate forecast loss after both shock signs.

CONFUSIONS TO AVOID

A significant γ can reflect a few crashes or misspecified mean. The usual persistence formula changes if innovation sign probabilities are asymmetric.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The registry stores indicator side, persistence convention, response curves, and side-specific diagnostics. Stress tests use the negative-shock path.

RELATED CONCEPTS AND QUICK RECALL

Related: TARCH, leverage effect, indicator function, news-impact curve. Recall task: Compute immediate shock weights when $\alpha=0.05$ and $\gamma=0.10$. What changes for a negative residual?

EGARCH and log variance

INTERMEDIATE

ONE-SENTENCE MEANING

EGARCH models the logarithm of conditional variance, which keeps variance positive without requiring every coefficient to be nonnegative. Signed standardized shocks represent asymmetry.

MEMORY JOGGER

Work in log scale: additive effects become multiplicative changes in variance. Magnitude and sign enter through separate terms.

WHY IT MATTERS IN QUANT RESEARCH

EGARCH can capture asymmetric reactions and scale effects more flexibly than a linear variance recursion. It is useful when symmetric GARCH leaves signed residual patterns.

FORMULA INTUITION

$\log h_t = \omega + \beta \log h_{t-1} + \alpha(|z_{t-1}| - E|z|) + \gamma z_{t-1}$. The sign of γ depends on the model convention.

SIMPLE MARKET EXAMPLE

A negative standardized shock can raise log variance more than a positive shock when γ has the appropriate sign. The same coefficient is not directly comparable with GJR γ .

PYTHON / SYSTEM IMPLEMENTATION

Document the exact equation and software convention, initialize log variance, and plot the implied news-impact curve. Validate finite forecasts under extreme z .

CONFUSIONS TO AVOID

Parameter names and asymmetry signs vary across libraries. Exponentiation can create very large forecasts even though positivity is guaranteed.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The model card stores the literal equation, innovation scaling, sign convention, constraints, and response plots. Deployment packages the fitted implementation, not coefficients alone.

RELATED CONCEPTS AND QUICK RECALL

Related: log link, asymmetric volatility, standardized shock, multiplicative effect. Recall task: Why does EGARCH not need nonnegative coefficients for positivity? Explain why γ signs cannot be compared blindly.

Innovation distributions and Student-t tails

INTERMEDIATE

ONE-SENTENCE MEANING

The innovation distribution describes standardized shocks after the variance process removes predictable scale. Gaussian, Student-t, skewed, and empirical choices allocate different probability to tails.

MEMORY JOGGER

First standardize by forecast volatility; then ask whether the remaining z scores look like the assumed distribution. Variance dynamics and tail shape solve different problems.

WHY IT MATTERS IN QUANT RESEARCH

Tail assumptions affect likelihood estimates, prediction intervals, VaR, and option-risk scenarios. A better variance recursion can still understate tail probability.

FORMULA INTUITION

Student-t z with degrees of freedom ν has polynomial tails; when standardized to unit variance it requires $\nu > 2$. Its fourth moment is finite only for $\nu > 4$.

SIMPLE MARKET EXAMPLE

A Gaussian model may call a standardized 4-sigma shock extraordinarily rare, while a fitted t model expects such events more often. Neither validates a bad tick.

PYTHON / SYSTEM IMPLEMENTATION

Estimate or predeclare ν , confirm unit-variance standardization, and inspect PIT, QQ, exceedance, and tail calibration out of sample.

CONFUSIONS TO AVOID

A heavy-tailed likelihood can absorb missing variance dynamics and make coefficients look stable. Very low ν may imply unstable moments and fragile risk summaries.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The distribution artifact travels with the variance forecast. Risk services consume full quantiles rather than reconstructing them from sigma and a default Gaussian.

RELATED CONCEPTS AND QUICK RECALL

Related: degrees of freedom, standardized residual, skewed-t, prediction interval. Recall task: State the moment conditions for Student-t variance and kurtosis. Why can a t likelihood not replace data cleaning?

Realized variance from intraday returns

INTERMEDIATE

ONE-SENTENCE MEANING

Realized variance sums squared intraday returns to estimate the quadratic variation accumulated during a day. With suitable sampling and market conditions, it is a more informative daily target than one squared close-to-close return.

MEMORY JOGGER

Many within-day increments reveal the day's scale. The estimator improves target signal but inherits intraday data and microstructure problems.

WHY IT MATTERS IN QUANT RESEARCH

Realized measures support forecast evaluation, HAR models, event studies, and risk monitoring. They let researchers distinguish forecast failure from noisy daily proxies.

FORMULA INTUITION

$RV_t = \sum_{i=1}^M r_{t,i}^2$. It includes continuous variation and jumps under ideal high-frequency asymptotics.

SIMPLE MARKET EXAMPLE

If four intraday returns are 0.4%, -0.3%, 0.2%, and 0.5%, realized variance is the sum of their squared decimal returns, not the square of their sum.

PYTHON / SYSTEM IMPLEMENTATION

Define sessions, overnight treatment, price type, interval, timestamp alignment, cleaning, and minimum coverage. Reconcile intraday sums to daily returns.

CONFUSIONS TO AVOID

Sampling ever faster does not guarantee accuracy because bid-ask bounce and timestamp errors can dominate. Omitting overnight movement understates full-day risk.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The realized-measure pipeline publishes RV, coverage, overnight component, jump flags, and quality tier. Forecast evaluation filters by approved tier.

RELATED CONCEPTS AND QUICK RECALL

Related: quadratic variation, intraday return, integrated variance, overnight gap. Recall task: Why is the sum of squared intraday returns different from the squared daily return? What cross terms disappear?

Sampling frequency and market noise

INTERMEDIATE

ONE-SENTENCE MEANING

Finer sampling captures more price variation but also more bid-ask bounce, discreteness, latency, and asynchronous quotes. The best grid balances information loss against contamination.

MEMORY JOGGER

More rows are not automatically more truth. Build a volatility signature plot across intervals and look for a stable region.

WHY IT MATTERS IN QUANT RESEARCH

Grid choice materially changes realized variance, jump tests, and model labels. A biased target can reward the wrong forecasting model.

FORMULA INTUITION

Observed log price can be written $p_t^{\text{obs}} = p_t^{\text{efficient}} + \eta_t$. At very fine grids, differencing η_t creates extra return variance and negative autocorrelation.

SIMPLE MARKET EXAMPLE

One-second transaction returns may produce twice the realized variance of five-minute midpoint returns because quotes alternate between bid and ask.

PYTHON / SYSTEM IMPLEMENTATION

Compare transaction, midpoint, previous-tick, and synchronized prices across grids. Record coverage and use subsampling or noise-robust estimators when justified.

CONFUSIONS TO AVOID

Choosing the smoothest grid after observing forecast performance leaks the test set. A single grid may not suit both liquid futures and thin equities.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The data registry stores instrument-specific sampling policy and signature diagnostics. A quality change triggers target revalidation before model promotion.

RELATED CONCEPTS AND QUICK RECALL

Related: bid-ask bounce, volatility signature plot, synchronization, market microstructure. Recall task: Explain why one-second RV can exceed five-minute RV even when efficient volatility is unchanged. Which price series helps diagnose it?

Range-based volatility estimators

INTERMEDIATE

ONE-SENTENCE MEANING

Range estimators use high and low prices, sometimes with open and close, to extract more scale information than a close-to-close return. Their efficiency depends on assumptions about drift, jumps, and trading continuity.

MEMORY JOGGER

The day's path leaves information in its range. Use it when OHLC data are reliable and the estimator's assumptions match the session.

WHY IT MATTERS IN QUANT RESEARCH

Range measures are useful when intraday trades are unavailable but daily OHLC is trustworthy. They provide robust benchmark targets and quality checks.

FORMULA INTUITION

Parkinson variance per period is $[\ln(H/L)]^2 / (4 \ln 2)$ under a driftless continuous diffusion without opening jumps.

SIMPLE MARKET EXAMPLE

A stock can close near its open after a wide intraday swing; close-to-close variance sees little, while the high-low range records substantial activity.

PYTHON / SYSTEM IMPLEMENTATION

Validate OHLC ordering, split adjustment, session definition, and zero ranges. Compare Parkinson, Garman-Klass, Rogers-Satchell, and close variance under events.

CONFUSIONS TO AVOID

Opening gaps and price limits violate simple range assumptions. Bad high or low ticks are especially damaging because the estimator uses extremes.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The measurement layer publishes several range estimators with assumption metadata. Model selection uses one predeclared target and others as diagnostics.

RELATED CONCEPTS AND QUICK RECALL

Related: Parkinson estimator, OHLC, Garman-Klass, Rogers-Satchell. Recall task: Give a day when close-to-close movement is small but range variance is large. Which assumption fails after an overnight gap?

Jumps and bipower variation

INTERMEDIATE

ONE-SENTENCE MEANING

Realized variance combines continuous variation and discrete jumps, while bipower variation estimates the continuous component under ideal conditions. Their difference provides evidence about jump variation.

MEMORY JOGGER

Separate ordinary accumulated turbulence from discontinuous repricing. Both are risk, but they may forecast and decay differently.

WHY IT MATTERS IN QUANT RESEARCH

Jump-aware targets improve event risk, tail scenarios, and volatility decomposition. They can prevent a one-time gap from being treated as the same process as persistent intraday noise.

FORMULA INTUITION

$BV_t = (\pi/2) \sum_{i=2}^M |r_t(t,i)| |r_t(t,i-1)|$ for a common high-frequency form. Positive $RV_t - BV_t$ is a simple jump-variation proxy, subject to testing and noise correction.

SIMPLE MARKET EXAMPLE

A macro announcement produces one large five-minute return and otherwise ordinary movement. RV jumps sharply while adjacent-product BV rises less.

PYTHON / SYSTEM IMPLEMENTATION

Use synchronized cleaned returns, enough intraday observations, robust jump tests, and event labels. Retain signed jump size and overnight jumps separately.

CONFUSIONS TO AVOID

Microstructure noise and clustered continuous volatility can create false jump signals. Truncating extreme returns before measurement removes the object being estimated.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The target service publishes continuous and jump components with confidence and quality flags. Forecast models may use separate dynamics and scenario rules.

RELATED CONCEPTS AND QUICK RECALL

Related: bipower variation, quadratic variation, jump test, event risk. Recall task: Explain why one isolated return raises RV more than BV. Name two sources of false jump detection.

HAR-RV and heterogeneous horizons

INTERMEDIATE

ONE-SENTENCE MEANING

The HAR-RV model forecasts realized volatility using daily, weekly, and monthly realized components. It approximates long memory through a small set of economically interpretable horizons.

MEMORY JOGGER

Different market participants operate on different clocks. Blend recent, weekly, and monthly scale instead of fitting dozens of lags.

WHY IT MATTERS IN QUANT RESEARCH

HAR is a strong realized-volatility baseline and often captures persistent multi-scale behavior with simple linear estimation.

FORMULA INTUITION

A common log specification is $\log RV_{t+1} = c + b_d \log RV_t + b_w \log RV_t^{(week)} + b_m \log RV_t^{(month)} + u_{t+1}$, with causal trailing averages.

SIMPLE MARKET EXAMPLE

A one-day shock raises the daily component immediately, while persistent turbulence gradually lifts weekly and monthly components, extending the forecast.

PYTHON / SYSTEM IMPLEMENTATION

Define window lengths, overlapping averages, log transform, zero handling, and target horizon. Fit every component inside each walk-forward origin.

CONFUSIONS TO AVOID

Overlapping weekly and monthly averages create collinearity and dependence. HAR describes horizon persistence but does not prove a specific trader-group mechanism.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The feature registry emits daily, weekly, and monthly RV with cutoffs. Forecast monitoring attributes changes to each horizon component.

RELATED CONCEPTS AND QUICK RECALL

Related: long memory, realized volatility, multi-scale model, overlapping window. Recall task: Describe how a one-day shock enters the daily, weekly, and monthly HAR components over time.

Intraday seasonality

INTERMEDIATE

ONE-SENTENCE MEANING

Intraday volatility follows recurring calendar patterns such as the open, close, auctions, lunch periods, and scheduled announcements. Modeling this seasonality prevents time-of-day effects from being mistaken for shocks.

MEMORY JOGGER

A 20-basis-point move means something different at the open than at midday. Normalize by the expected scale for that clock location.

WHY IT MATTERS IN QUANT RESEARCH

Seasonal baselines improve anomaly detection, execution schedules, intraday risk, and realized-measure comparability across days.

FORMULA INTUITION

One approach is $r_{(d,i)} = s_i u_{(d,i)}$, where s_i is a causal time-of-day scale and u is deseasonalized. Normalize s_i so the daily variance total remains interpretable.

SIMPLE MARKET EXAMPLE

The first fifteen minutes may routinely contribute far more variance than a midday interval. A raw threshold over-alerts at the open and misses a midday event.

PYTHON / SYSTEM IMPLEMENTATION

Estimate profiles by venue, weekday, event class, and daylight-saving regime using only prior dates. Monitor coverage and profile drift.

CONFUSIONS TO AVOID

A profile estimated from all dates leaks future market structure. Averaging scheduled-event days with ordinary days can blur both patterns.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The calendar service supplies phase labels; the feature layer produces seasonal and deseasonalized scale; execution consumes phase-specific risk.

RELATED CONCEPTS AND QUICK RECALL

Related: U-shaped volatility, deseasonalization, auction, calendar effect. Recall task: Explain why one raw return threshold is inappropriate across the trading day. How should s_i be normalized?

Implied versus expected realized volatility

INTERMEDIATE

ONE-SENTENCE MEANING

Implied volatility is the volatility input consistent with option prices under a pricing convention, while expected realized volatility is a statistical forecast of future path variation. Their difference includes risk premia and market frictions.

MEMORY JOGGER

Implied is a priced measure, not a pure physical forecast. Compare matched horizon, strike convention, calendar, and variance units.

WHY IT MATTERS IN QUANT RESEARCH

The relationship informs option valuation, hedging, variance risk premium research, and cross-checks on time-series forecasts.

FORMULA INTUITION

A simple variance risk premium convention is implied variance minus forecast expected realized variance over the same horizon. Sign and annualization conventions must be stated.

SIMPLE MARKET EXAMPLE

Thirty-day at-the-money implied volatility may be 24% while a time-series model forecasts 19% realized. The gap is not free profit; it compensates for tail and volatility risk plus trading costs.

PYTHON / SYSTEM IMPLEMENTATION

Align option timestamp, expiration, forward, day count, overnight treatment, and realized window. Preserve bid-ask and surface quality.

CONFUSIONS TO AVOID

Comparing implied standard deviation with realized variance is a unit error. Using a final settlement-based realized target can leak information into an earlier option snapshot.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The derivatives layer publishes implied surfaces and quality; the forecast layer publishes physical distributions; strategy logic computes matched, versioned gaps.

RELATED CONCEPTS AND QUICK RECALL

Related: risk-neutral measure, physical measure, variance risk premium, volatility surface. Recall task: List four alignment conditions for comparing implied and realized volatility. Why is a positive gap not automatically arbitrage?

Forecast target and horizon alignment

INTERMEDIATE

ONE-SENTENCE MEANING

A volatility model is valid only relative to a declared target, horizon, information cutoff, and decision. One-day close variance, next-session realized variance, and ten-day cumulative variance are not interchangeable labels.

MEMORY JOGGER

Define the answer before selecting the model. Every forecast row needs a legal origin and a later target window that does not overlap hidden information.

WHY IT MATTERS IN QUANT RESEARCH

Target alignment determines training labels, purge gaps, loss functions, and whether the result can support the intended risk action.

FORMULA INTUITION

For origin t and horizon H , a cumulative realized-variance target can be $y_{-}(t,H)=\sum_{(j=1)}^H r_{-}(t+j)^2$ or a high-frequency analogue, with exact endpoints documented.

SIMPLE MARKET EXAMPLE

A model issued after today's close cannot forecast a target that begins at today's open. Calling it one-day-ahead does not repair the timing mismatch.

PYTHON / SYSTEM IMPLEMENTATION

Store origin, cutoff, availability, target start, target end, horizon unit, overlap, and earliest action. Generate labels from a versioned calendar.

CONFUSIONS TO AVOID

Overlapping H -day labels inflate row counts and require purging or dependence-aware inference. A target chosen after viewing model rankings contaminates comparison.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The experiment schema binds forecast and target contracts. Evaluation rejects rows whose timestamp inequalities or units fail.

RELATED CONCEPTS AND QUICK RECALL

Related: label window, forecast origin, purging, cumulative variance. Recall task: Define a legal ten-day variance target for a forecast issued after the close. Which adjacent labels overlap?

Loss functions for variance forecasts

INTERMEDIATE

ONE-SENTENCE MEANING

A loss function translates forecast error into the decision's penalty. MSE emphasizes large absolute variance misses, while QLIKE is scale-sensitive and often robust to noisy variance proxies under stated conditions.

MEMORY JOGGER

Models are ranked by what the loss values, not by an abstract universal accuracy. Use at least one statistical and one decision-aware view.

WHY IT MATTERS IN QUANT RESEARCH

Loss choice can reverse model rankings and determine whether crisis underprediction receives enough weight. It must be locked before final comparison.

FORMULA INTUITION

For target $y > 0$ and forecast $h > 0$, QLIKE can be written $y/h + \log h$ up to a target-only constant. Variance MSE is $(y-h)^2$.

SIMPLE MARKET EXAMPLE

Forecasts 1 and 4 against target 4 have different errors under MSE and QLIKE; QLIKE strongly penalizes a forecast that is too close to zero relative to realized variance.

PYTHON / SYSTEM IMPLEMENTATION

Clip only by a predeclared numerical floor, report loss by state and horizon, and use paired time-series tests or bootstrap on loss differences.

CONFUSIONS TO AVOID

R-squared on squared returns is not a sufficient volatility score. Optimizing the floor or crisis weights after results is another model-selection degree of freedom.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The evaluation service publishes QLIKE, MSE, interval calibration, and decision loss from the same rows. Promotion uses a predeclared scorecard.

RELATED CONCEPTS AND QUICK RECALL

Related: QLIKE, proper scoring rule, forecast comparison, loss differential. Recall task: Explain why QLIKE dislikes a near-zero forecast when realized variance is high. Why can model rankings differ from MSE?

Walk-forward volatility validation

INTERMEDIATE

ONE-SENTENCE MEANING

Walk-forward validation refits every learned component at historical origins and scores only later targets. It reproduces parameter drift, warm-up, data availability, and forecast issuance.

MEMORY JOGGER

Test the procedure, not a full-sample fitted path. Long-run variance, seasonal profiles, distribution tails, and thresholds all belong inside each fold.

WHY IT MATTERS IN QUANT RESEARCH

Volatility models are vulnerable to regime leakage and revised data. Rolling origins reveal whether forecast gains persist through calm, crisis, and recovery states.

FORMULA INTUITION

At each origin t_j , fit on an approved history, issue $h(t_j, H)$, wait for $y(t_j, H)$, and concatenate only out-of-sample loss. Purge overlap where required.

SIMPLE MARKET EXAMPLE

A Student-t GARCH fit on the entire sample uses a future crisis to set tail thickness for earlier forecasts, making historical intervals unrealistically informed.

PYTHON / SYSTEM IMPLEMENTATION

Predeclare training window, cadence, initialization, search family, target, purge, and fallback. Save one immutable artifact per origin.

CONFUSIONS TO AVOID

Random cross-validation destroys chronology. Reusing a holdout to choose GARCH variants converts it into training data.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The tracker links each forecast to data snapshot, code, parameters, target, loss, and decision. Final metrics aggregate only untouched origins.

RELATED CONCEPTS AND QUICK RECALL

Related: rolling origin, temporal validation, refit cadence, holdout contamination. Recall task: List every quantity that must be refit for an intraday seasonal GARCH model. Which rows can enter final loss?

Calibration and standardized residual diagnostics

INTERMEDIATE

ONE-SENTENCE MEANING

A calibrated variance model produces standardized residuals whose scale and tail frequencies match the assumed distribution over time. Diagnostics test remaining mean, variance, asymmetry, and distribution errors.

MEMORY JOGGER

Divide realized shocks by issued sigma, then inspect whether z behaves as promised. A good average loss can hide systematic undercoverage.

WHY IT MATTERS IN QUANT RESEARCH

Calibration determines whether forecast intervals, VaR, and leverage controls have the advertised frequency. Residual patterns identify the missing mechanism.

FORMULA INTUITION

$z_t = \epsilon_t / \sqrt{h_t}$. Desired properties include $E[z_t]$ near zero, $E[z_t^2]$ near one, little ACF in z and z^2 , and calibrated quantiles.

SIMPLE MARKET EXAMPLE

If z^2 remains autocorrelated, scale dynamics remain. If negative z has heavier tails, the innovation distribution or asymmetry is incomplete.

PYTHON / SYSTEM IMPLEMENTATION

Plot z , z^2 ACF, QQ, PIT, coverage, exceedance clusters, and state-conditioned summaries. Use out-of-sample forecasts for decisive diagnostics.

CONFUSIONS TO AVOID

Passing one Ljung-Box test does not prove calibration. Repeated lag and distribution searches require multiplicity discipline.

WHERE THIS FITS IN THE RESEARCH SYSTEM

Monitoring computes the same diagnostics on rolling live cohorts and maps failures to data, variance recursion, tail, or regime response.

RELATED CONCEPTS AND QUICK RECALL

Related: PIT, QQ plot, exceedance test, residual whiteness. Recall task: What does autocorrelation in z squared imply? Name two diagnostics for tail calibration.

Structural breaks and volatility regimes

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility dynamics can change after policy shifts, market redesigns, crises, or liquidity events. A single stationary recursion may average incompatible centers and persistence levels.

MEMORY JOGGER

High persistence may be a slow model or an unmodeled break. Compare both stories before increasing the memory parameter.

WHY IT MATTERS IN QUANT RESEARCH

Breaks affect long-run variance, shock response, tail frequency, and execution costs simultaneously. Regime-aware controls can be safer than automatic refitting during stress.

FORMULA INTUITION

A regime model allows state-specific parameters $h_t(s)$, while a break model changes parameters at detected times. Causal filtered states differ from retrospective smoothed labels.

SIMPLE MARKET EXAMPLE

A market moves from 12% to 35% volatility after a structural event. One GARCH fit may report near-unit persistence because it tries to bridge both eras.

PYTHON / SYSTEM IMPLEMENTATION

Use break-aware challengers, rolling parameter paths, state-conditioned loss, and delayed detection simulations. Define warning, degraded, and stop responses.

CONFUSIONS TO AVOID

Retrospective regime labels use future data and overstate trading usefulness. Refitting immediately on defective crisis data can absorb a feed problem.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The control plane separates data incidents, temporary volatility states, and model invalidation. Each state has an owner and fallback.

RELATED CONCEPTS AND QUICK RECALL

Related: change point, Markov switching, filtered probability, local stationarity. Recall task: Explain how one level break can mimic high GARCH persistence. Why are smoothed regime labels illegal online?

Parameter and model uncertainty

INTERMEDIATE

ONE-SENTENCE MEANING

Volatility forecasts are uncertain because parameters are estimated, models are approximations, and specification choices are selected from data. A point forecast suppresses all three sources.

MEMORY JOGGER

Ask how the forecast changes across plausible coefficients, windows, distributions, and model families. Capital should survive that range.

WHY IT MATTERS IN QUANT RESEARCH

Near-boundary persistence, tail degrees of freedom, and long-run variance can create wide risk intervals. Model disagreement often rises exactly during stress.

FORMULA INTUITION

Predictive variance decomposes conceptually into process uncertainty plus parameter and model uncertainty. Bootstrap, Bayesian draws, and ensembles approximate different pieces.

SIMPLE MARKET EXAMPLE

Two models forecast 20% and 34% annualized volatility after a shock. Averaging blindly may hide a meaningful disagreement that should reduce leverage.

PYTHON / SYSTEM IMPLEMENTATION

Propagate parameter draws, run predeclared challengers, and report forecast quantiles and dispersion. Maintain a search ledger for every tested specification.

CONFUSIONS TO AVOID

Selecting the best of dozens of variants biases its loss downward. An ensemble is not robust if all members share the same target defect.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The forecast API includes point, interval, model dispersion, and health. Policy uses conservative quantiles or explicit disagreement discounts.

RELATED CONCEPTS AND QUICK RECALL

Related: bootstrap, ensemble, model risk, specification search. Recall task: Separate process, parameter, and model uncertainty for a GARCH forecast. How should large model disagreement affect leverage?

Volatility forecasts in sizing and limits

INTERMEDIATE

ONE-SENTENCE MEANING

A volatility forecast becomes useful when translated into bounded positions, risk limits, and scenarios. The mapping must specify horizon, covariance, costs, forecast uncertainty, and model health.

MEMORY JOGGER

Inverse-volatility sizing stabilizes predicted risk only if sigma is aligned and trustworthy. Always cap the position and the rate of change.

WHY IT MATTERS IN QUANT RESEARCH

Scale forecasts drive volatility targeting, portfolio risk, stop distances, liquidity reserves, and order sizing. Bad low-vol forecasts can create dangerous leverage.

FORMULA INTUITION

A simple target is $w_t = \text{clip}(k/\sigma_{\hat{t}}, -w_{\max}, w_{\max})$, possibly using a conservative sigma quantile and turnover constraint. Portfolio risk needs the full covariance matrix.

SIMPLE MARKET EXAMPLE

If forecast volatility falls from 20% to 10%, naive inverse-vol sizing doubles exposure. A floor and model-health discount prevent quiet-period leverage from exploding.

PYTHON / SYSTEM IMPLEMENTATION

Use horizon-matched forecasts, uncertainty buffers, caps, turnover penalties, liquidity limits, and stress overrides. Backtest the whole target-to-fill path.

CONFUSIONS TO AVOID

Inverse volatility ignores expected return, correlation, and costs. Procyclical deleveraging can force sales after a volatility spike and amplify impact.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The portfolio engine consumes scale plus confidence, covariance, cost, and health. It logs raw target, each constraint, and final order.

RELATED CONCEPTS AND QUICK RECALL

Related: volatility targeting, leverage cap, risk budget, procyclicality. Recall task: Design an inverse-volatility rule with a floor, cap, and health discount. What happens when sigma halves?

Production monitoring and artifact lineage

INTERMEDIATE

ONE-SENTENCE MEANING

Production monitoring tracks data freshness, realized calibration, parameter drift, forecast sensitivity, and capital consequences. Lineage makes every live forecast and position reproducible.

MEMORY JOGGER

Monitor the chain from tick to target, not only the final sigma. Every alert needs a material threshold, owner, and safe action.

WHY IT MATTERS IN QUANT RESEARCH

Stale prices can suppress measured volatility, bad ticks can explode it, and optimizer failures can freeze yesterday's state. These faults produce opposite but equally dangerous positions.

FORMULA INTUITION

A health score may combine coverage, freshness, z^2 mean, exceedance rate, parameter boundary distance, model dispersion, and target sensitivity. Thresholds require persistence rules.

SIMPLE MARKET EXAMPLE

A frozen feed creates zero returns and shrinking EWMA volatility, which raises target leverage. Freshness and unchanged-price checks must block the position before calibration notices.

PYTHON / SYSTEM IMPLEMENTATION

Persist raw snapshot, cleaning, returns, features, fit, parameters, forecast, target, decision, order, and fill. Test replay and fallback in shadow mode.

CONFUSIONS TO AVOID

Automatic refitting can absorb feed defects. Monitoring only PnL may miss a dangerous low-vol forecast when directional returns happen to be favorable.

WHERE THIS FITS IN THE RESEARCH SYSTEM

The control plane supports warning, reduced-risk, and hard-stop tiers with versioned fallback and rollback. Incidents classify data, model, policy, or execution failure.

RELATED CONCEPTS AND QUICK RECALL

Related: observability, model registry, fallback, decision lineage. Recall task: Draft three monitoring tiers and a safe action for each. Which artifacts reproduce one live risk decision?


Returns, clustering, and conditional scale

INTERMEDIATE

READING RULE

Return signs alternate unpredictably, while magnitude forms persistent episodes. A conditional model follows the scale envelope rather than forecasting each sign.

RETURN SHOCKS



INTERPRETATION


The first cluster lifts the scale state and then decays. A later cluster arrives before the process fully returns to calm. Evaluate whether the issued envelope makes standardized shocks stable; do not judge the model by sign accuracy.

Choosing a volatility model family

INTERMEDIATE

START WITH THE TARGET

Select the simplest model that addresses a validated residual pattern. Model family follows measurement, horizon, asymmetry, tails, and data frequency.



Persistence, half-life, and horizon risk

INTERMEDIATE

QUESTION

How long does a variance shock remain operationally important? The nonlinear rise near one makes persistence uncertainty central to capital planning.

5 min	0.2h	0.3h	0.5h	1.1h	1.9h	0.1d	0.2d
30 min	1.0h	1.6h	0.1d	0.3d	0.5d	0.7d	1.4d
1 hour	1.9h	0.1d	0.3d	0.6d	0.9d	1.4d	2.9d
1 day	1.9d	3.1d	6.6d	13.5d	22.8d	34.3d	69.0d
1 week	9.7d	15.5d	32.9d	67.6d	113.8d	171.5d	344.8d

INTERPRETATION

Persistence is tied to the observation interval. A daily coefficient cannot be copied into an hourly model, and aggregation does not preserve the same number.

CAPITAL CONSEQUENCE

If the upper parameter interval approaches one, a shock can affect risk for far longer than the point estimate suggests. Use the slow plausible path for limits.

FORMULA


Half-life in periods is $\log(0.5)/\log(p)$. For symmetric GARCH(1,1), p is commonly $\alpha + \beta$.

From market event to scored forecast

INTERMEDIATE

TIMING STANDARD

A forecast can use only finalized observations and fitted state available at its origin. Its target begins later and becomes scoreable only after the entire window completes.



COMMON LEAKAGE

Using the close to update h_t and filling at that same close assumes simultaneity. Estimating a seasonal profile, tail parameter, or long-run variance from the full sample imports future states.

OVERLAPPING TARGETS

A ten-day target issued daily shares nine days with its neighbor. Purging, effective-sample reasoning, and block inference must follow the label geometry.

AUDIT RECORD


Persist event, availability, state cutoff, model version, forecast origin, target bounds, decision, order, fill, realized measure, and score.

Variance forecast path after a shock

INTERMEDIATE

SCENARIO

Current variance is four times the long-run level. Forecasts revert toward the anchor, while parameter and process uncertainty widen the range of plausible outcomes.



DECISION IMPLICATION

Cumulative holding-period risk sums the stepwise variance path. A conservative limit also accounts for the upper band, covariance, and nonlinear costs rather than using the blue point path alone.

Method selection by measurement problem

INTERMEDIATE

QUESTION	FIRST METHOD	KEY ASSUMPTION	PRIMARY RISK
Transparent recent scale	Rolling variance	Stable local window	Hard cutoff and overlap
Smooth daily update	EWMA	Geometric memory	No separate long-run anchor
Daily conditional variance	GARCH(1,1)	Stable recursion	Boundary persistence
Signed shock response	GJR or EGARCH	Repeated side evidence	Convention and tail risk
Heavy standardized tails	Student-t likelihood	Finite required moments	Bad data absorption
High-frequency target	Realized variance	Clean synchronized prices	Microstructure noise
OHLC-only measurement	Range estimator	Reliable session extremes	Gaps and bad highs/lows
Multi-horizon RV forecast	HAR-RV	Stable horizon components	Overlap and regime shifts

SELECTION DISCIPLINE

Choose by target, timing, residual evidence, later loss, and capital behavior. A more elaborate likelihood is not a substitute for a legal target or clean measurement.

Failure modes and diagnostic actions

INTERMEDIATE

SYMPTOM	LIKELY CAUSE	FIRST ACTION
z squared remains autocorrelated	Missing scale dynamics	Add lags only after data checks
Negative z has heavier tail	Asymmetry or skew	Compare signed calibration
RV rises at finer grids	Microstructure noise	Use signature plot and midquotes
Persistence near one	Slow decay or regime break	Run break-aware challengers
Low live sigma with stale prices	Frozen feed	Block targets and use fallback
Good average loss, poor crisis coverage	State underprediction	Report state loss and tail tests
Frequent optimizer boundary hits	Weak identification	Rescale, restart, simplify
PnL falls while QLIKE improves	Decision mismatch	Audit costs, covariance, and mapping

ESCALATION RULE

Quarantine data defects before refitting. When the economic relation changes, reduce risk and compare challengers before promoting a new state model.

Formula sheet and interpretation map

INTERMEDIATE

RETURN AND SCALE

$r_t = \log(P_t/P_{t-1})$; $h_t = \text{Var}(r_t | F_{t-1})$; $\sigma_t = \sqrt{h_t}$. Keep variance and volatility units explicit.

ROLLING VARIANCE

$s_t^2 = \text{sum of trailing centered squared returns divided by } L-1$. State window, cutoff, mean, gaps, and annualization.

EWMA

$h_t = \lambda h_{t-1} + (1-\lambda)\epsilon_{t-1}^2$. Shock weights decay geometrically.

GARCH(1,1)

$h_t = \omega + \alpha \epsilon_{t-1}^2 + \beta h_{t-1}$. Long-run $h_{\text{bar}} = \omega / (1 - \alpha - \beta)$ when stationary.

PERSISTENCE

$\rho = \alpha + \beta$; half-life = $\log(0.5) / \log(\rho)$. Carry parameter intervals through the nonlinear transform.

ASYMMETRY

GJR adds $\gamma I(\epsilon_{t-1} < 0) \epsilon_{t-1}^2$. Under symmetric signs, common persistence is $\alpha + \beta + \gamma/2$.

REALIZED MEASURES

$RV = \text{sum intraday } r_i^2$; $BV = (\pi/2) \sum |r_i| |r_{i-1}|$. Sampling and quality assumptions matter.

RANGE VARIANCE

Parkinson = $[\ln(H/L)]^2 / (4 \ln 2)$ per period under driftless continuous trading without opening jumps.

QLIKE

For $y > 0$ and $h > 0$, loss = $y/h + \log h$ up to a target-only constant. Never feed negative forecasts.

Worked calculation: shock to governed risk

INTERMEDIATE

1. DEFINE THE MODEL

Daily GARCH uses $\omega=0.000004$, $\alpha=0.08$, $\beta=0.90$. Persistence $p=0.98$, so the stationary condition holds but decay is slow.

2. LONG-RUN VARIANCE

$h_{\text{bar}}=0.000004/(1-0.98)=0.0002$. Long-run daily volatility is $\sqrt{0.0002}=1.414\%$.

3. UPDATE AFTER A SHOCK

Yesterday $\epsilon=-3\%$ and prior $h=0.000225$. New $h=0.000004+0.08(0.03^2)+0.90(0.000225)=0.0002785$.

4. CONVERT THE FORECAST

One-day $\sigma=\sqrt{0.0002785}=1.669\%$. IID-style annualized display is $1.669\%*\sqrt{252}=26.50\%$, with assumptions disclosed.

5. COMPUTE HALF-LIFE

$\log(0.5)/\log(0.98)=34.3$ daily periods. This is expected variance-deviation decay, not a guaranteed return to calm.

6. FORECAST DAY TWO

$E[h_{(t+2)}]=h_{\text{bar}}+0.98(h_{(t+1)}-h_{\text{bar}})=0.00027693$. The path converges gradually toward 0.0002.

7. SIZE CONSERVATIVELY

A volatility-target rule uses a confidence-buffered sigma, a floor, cap, covariance, and liquidity constraint, not the point estimate alone.

8. VALIDATE LATER

Score against a legal realized-variance target with QLIKE and MSE; inspect z , z squared, coverage, and crisis cohorts.

9. REGISTER EVIDENCE

Store data snapshot, equation, parameters, covariance, forecast path, target bounds, losses, decision, controls, and fallback.

Validation scorecard for a volatility model

INTERMEDIATE

TARGET EVIDENCE

Price, return, session, horizon, units, target window, overnight treatment, quality tier, and earliest legal origin are fixed.

BASELINE EVIDENCE

Rolling, EWMA, seasonal, and naive long-run forecasts are scored on the same rows and losses.

DYNAMIC EVIDENCE

Variance recursion, asymmetry, tails, persistence, optimizer behavior, residuals, and forecast path pass diagnostics.

MEASUREMENT EVIDENCE

Realized or range target survives grid, microstructure, gap, missing-data, and bad-tick challenges.

VALIDATION EVIDENCE

Every learned quantity is refit by origin; overlap is purged; losses and coverage hold across states.

MODEL-RISK EVIDENCE

Parameter intervals, challenger dispersion, search ledger, boundary cases, and break scenarios are documented.

DECISION EVIDENCE

Sizing includes confidence buffer, floors, caps, covariance, liquidity, turnover, costs, and stress response.

OPERATIONAL EVIDENCE

Freshness, calibration, sensitivity, owner, fallback, rollback, shadow replay, and lineage are tested.

PROMOTION GATE

No statistical gain can compensate for illegal timing, defective targets, missing fallback, or uncontrolled capital sensitivity.

Final active recall and system index

INTERMEDIATE

OBJECT AND UNITS

Distinguish return direction, variance, standard deviation, realized measure, implied measure, and their horizons.

MEASUREMENT MAP

Choose squared returns, intraday RV, ranges, continuous variation, jumps, or option-implied inputs with quality controls.

MODEL MAP

Use rolling and EWMA baselines, GARCH recursion, asymmetry, tail distributions, HAR horizons, and break challengers only as needed.

FORECAST MAP

Carry one-step, multi-step, cumulative, long-run, and uncertainty outputs; avoid blind square-root-of-time scaling.

DIAGNOSTIC MAP

Inspect z , z squared, ACF, QQ, PIT, coverage, exceedance clusters, parameter boundaries, and state loss.

VALIDATION MAP

Freeze every artifact by origin, align targets, purge overlap, count the search family, and preserve untouched evidence.

CAPITAL MAP

Translate scale into bounded targets using covariance, confidence, floors, caps, liquidity, costs, turnover, and stress.

PRODUCTION MAP

Monitor data clocks, unchanged prices, target coverage, calibration, drift, sensitivity, model disagreement, and fallback.

FINAL RECALL

Explain why volatility is not direction, why RV depends on sampling, why persistence differs from beta, and why implied volatility is not a pure forecast.